Supports from IA-32

- >AT&T Assembly language
- >Function calls
- >Recap IA-32 memory management
- >I/O & Peripherals
- >IA-32 Interrupt Processing

LabO is completely out!

- 打字游戏 (deadline: 2014/03/09 23:59:59)
- 自己的游戏 (deadline: 2014/03/23 23:59:59)
- 完整的工程和实验报告 (deadline: 2014/03/25 23:59:59)
- Have fun!

Accessing Google

- For some reasons, Google maybe not easy to access.
- Modify the "hosts" file
 - /etc/hosts
- Search "google hosts" for more details.

AT&T assembly language

Assembly language

- a "middle-level" programming language
- the symbolization of machine language
- can operate registers and hardware directly
 - To know the state of hardware, we simply read from its control register(s)
 - To let the hardware do something, we simply write the specific "command word" to its control register(s)

Assembly language (cont.)

inb \$0x60, %al

- read one byte from the port 0x60, and store it to the register "al"
 - obtain the scancode of the key pressed down

outb %al, \$0x20

- write the byte storing in the register "al" to the port 0x20
 - issue a command word to the I8259A processor

Addressing

• syntax:

mem_location[(base, [index, [scale]])]

example:

gcc -S -o test.S test.c

```
8 main:
                %ebp
       pushl
       movl
                %esp, %ebp
10
       andl
                $-16, %esp
11
       subl
12
                $432, %esp
       movl
                $1, 428(%esp)
13
       movl
14
                428(%esp), %eax
       addl
15
                $3, %eax
       movl
                $97, 28(%esp,%eax,4)
16
```

```
1 #include "stdio.h"
2
3 int main(){
4    int n = 1;
5    int a[100];
6    a[n + 3] = 97;
7    printf("a[n + 3] = %d\n", a[n + 3]);
8    return 0;
9 }
```

```
#esp &= Oxfffffff0
#esp -= 432
#n = 1
#eax = n
#eax += 3
#a[eax] = 97
```

Example - "Hello, world!"

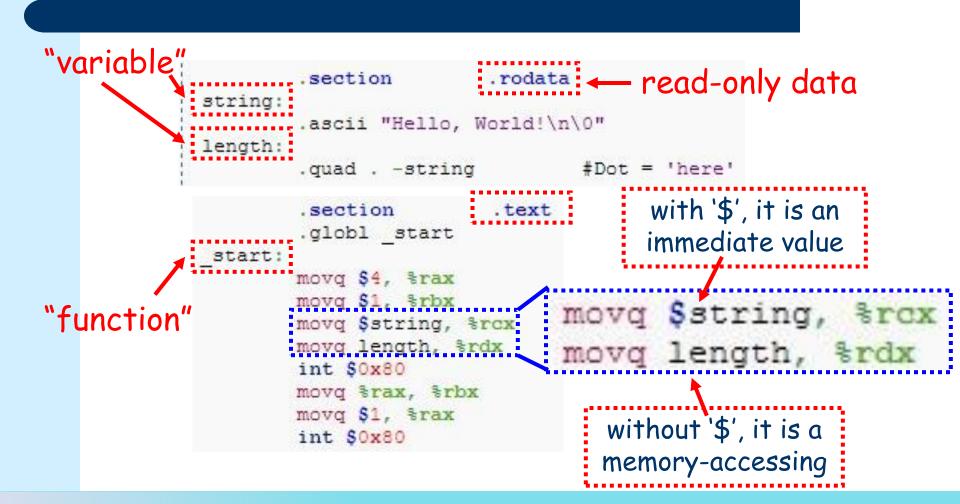
-http://en.wikipedia.org/wiki/List_of_Hello_world_program_examples

Assembly language – x86-64 Linux, AT&T syntax

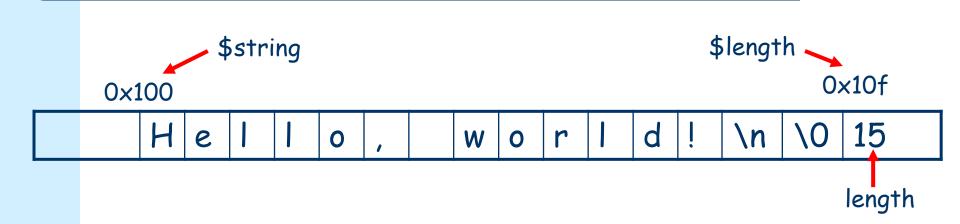
Main article: Assembly language

```
section
                     .rodata
                                           in IA-32 assembly language:
string:
                                                   movl $4, %eax
        .ascii "Hello, World!\n\0"
length:
                               #Dot = 'here'
        .guad . -string
                   .text
        .section
                               #Make entry point visible to linker
        .globl start
start:
       movq $4, %rax
                               #4=write
                               #1=stdout
       movq $1, %rbx
       movg Sstring, %rcx
       movg length, %rdx
       int $0x80
                               #Call Operating System
                               #Make program return syscall exit status
       movq %rax, %rbx
       movq $1, %rax
                               #1=exit
       int $0x80
                               #Call System Again
```

Example - "Hello, world!" (cont.)



Example - "Hello, world!" (cont.)



movl \$1, %eax movl 1, %eax

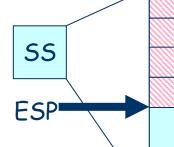


http://sock-raw.org/netsec/stack

Oxfffffff

Stack

- in data structure
 - stack = array + top pointer
- in hardware
 - stack = SS + ESP
 - 55 specifies a segment of memory
 - ESP indicates the top of stack
 - grow downward (to the direction of 0)
 - 4 bytes per elements



Stack (cont.)

```
operations - push, pop
push! %eax
```

- It is identical to

```
subl $4, %esp
movl %eax, (%esp)
```

popl %ebx

- It is identical to

```
movl (%esp), %ebx
addl $4, %esp
```

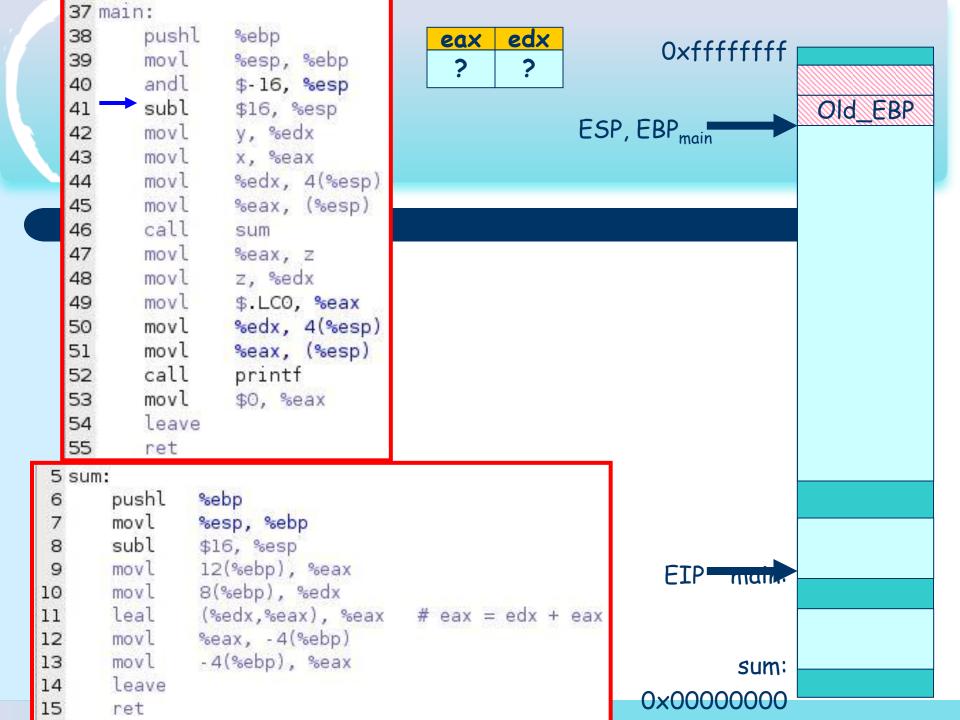
Function calls

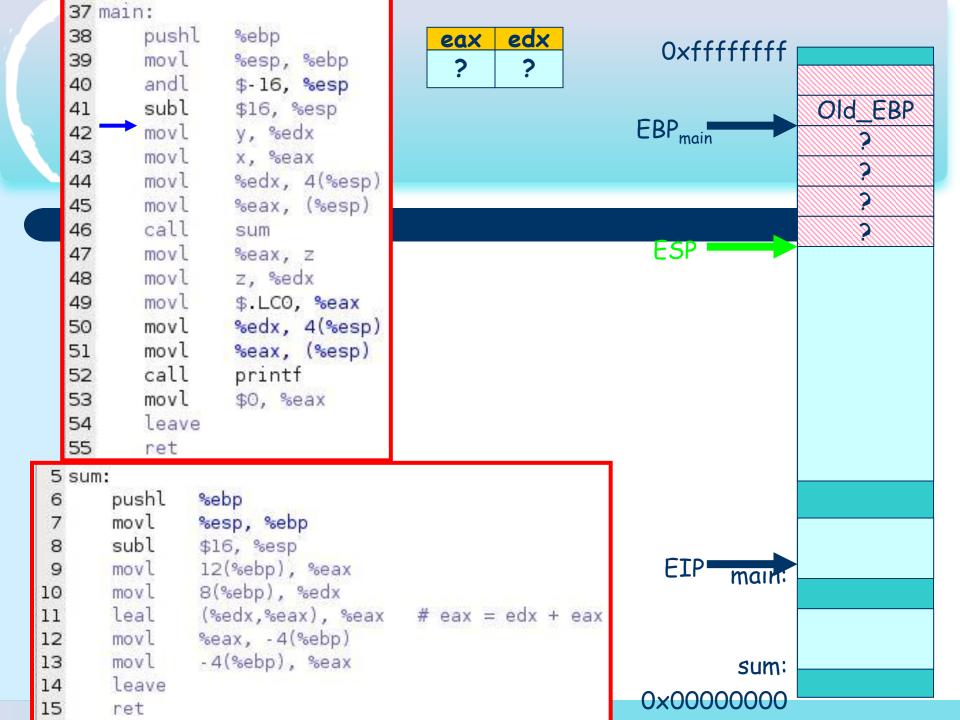
- Calling a function can be easy jmp printf
- But how to assure it can return correctly?
 - save the return address before calling
- Where to save?
 - Stack !!!
- Memories for local variables are usually allocated in stack.
 - Why?

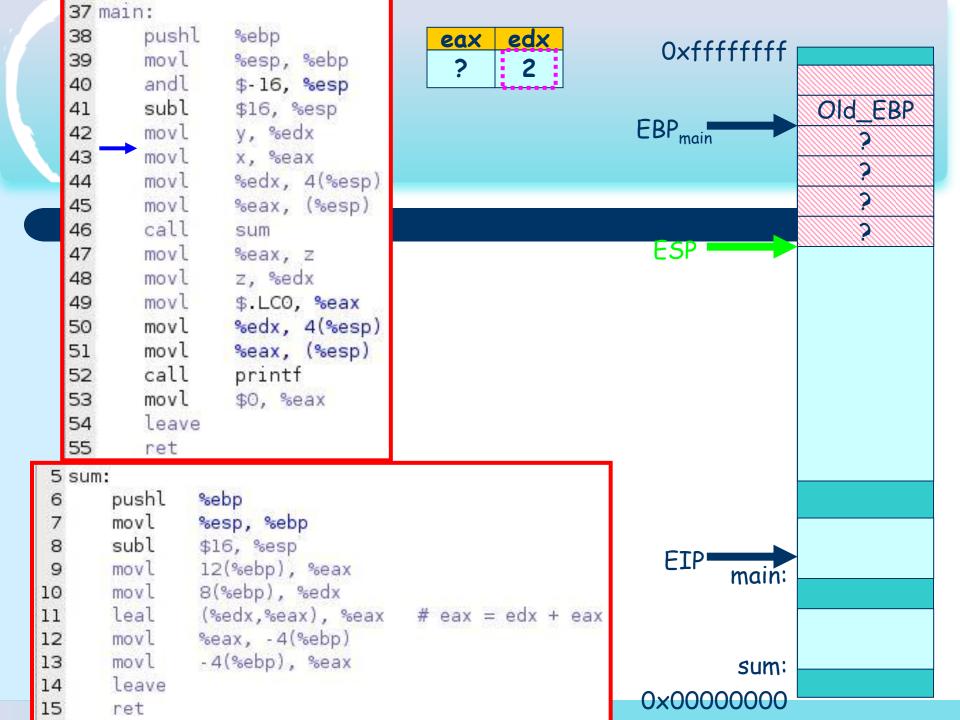
An example

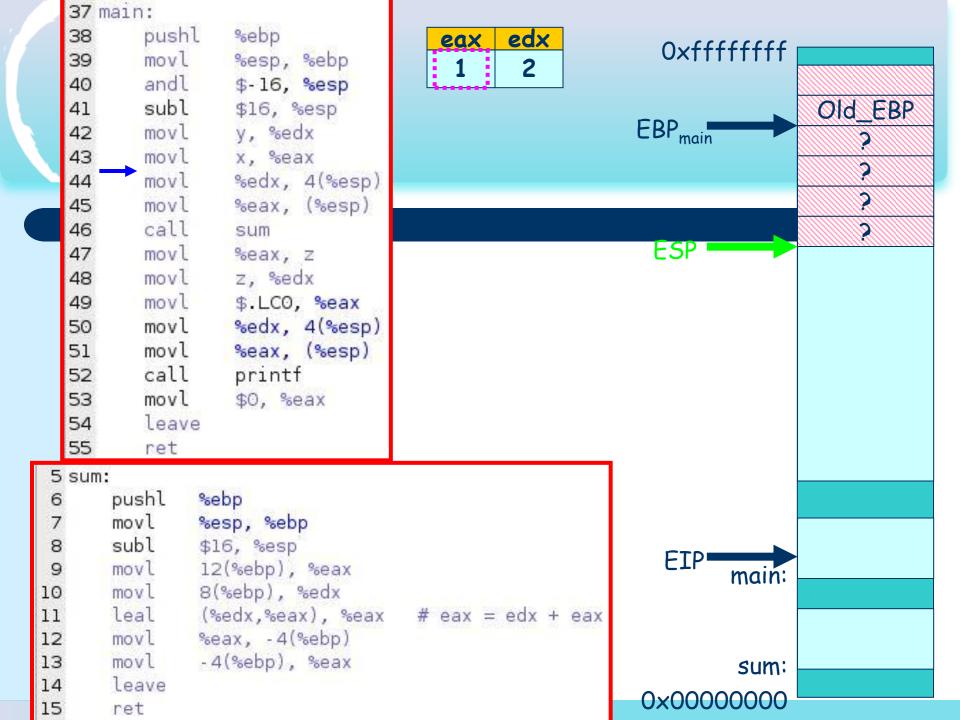
```
1 #include "stdio.h"
 3 int sum(int a, int b){
      int c = a + b;
    return c;
8 int x = 1, y = 2, z;
10 int main(){
  z = sum(x, y);
11
12 printf("%d\n", z);
   return 0;
13
14 }
```

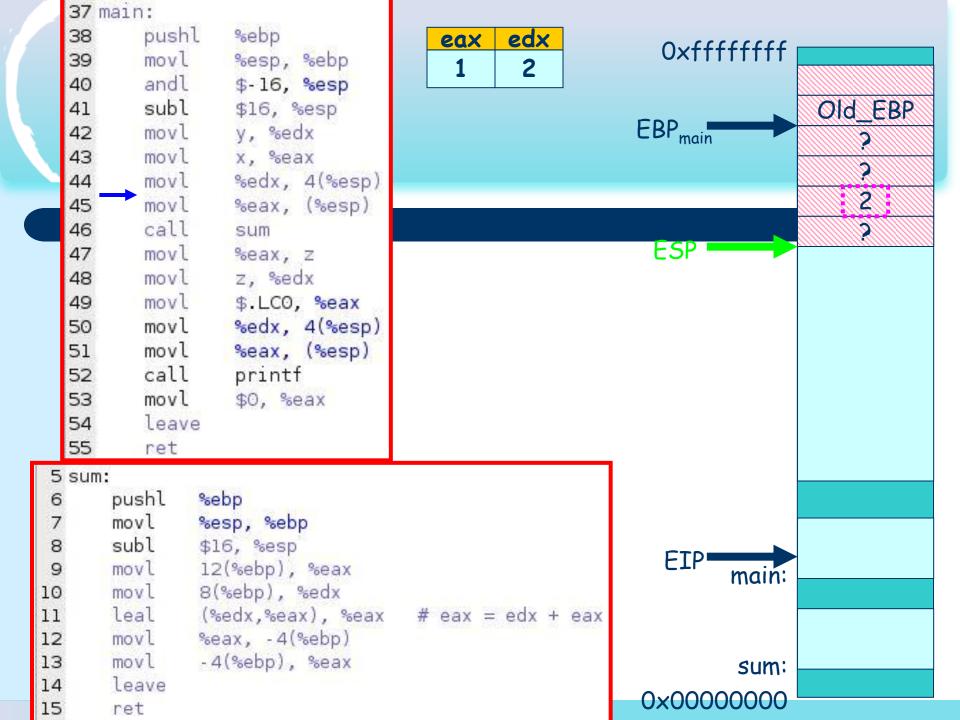
```
17 .globl x
18 .data
19 .align 4
20 .type x, @object
21
     .size x, 4
22 x:
23 .long 1
24 .globl y
25 .align 4
26 .type y, @object
27
     .size y, 4
28 y:
     .long 2
29
30
     .comm z,4,4
     .section .rodata
31
32 .LCO:
     .string "%d\n"
33
```

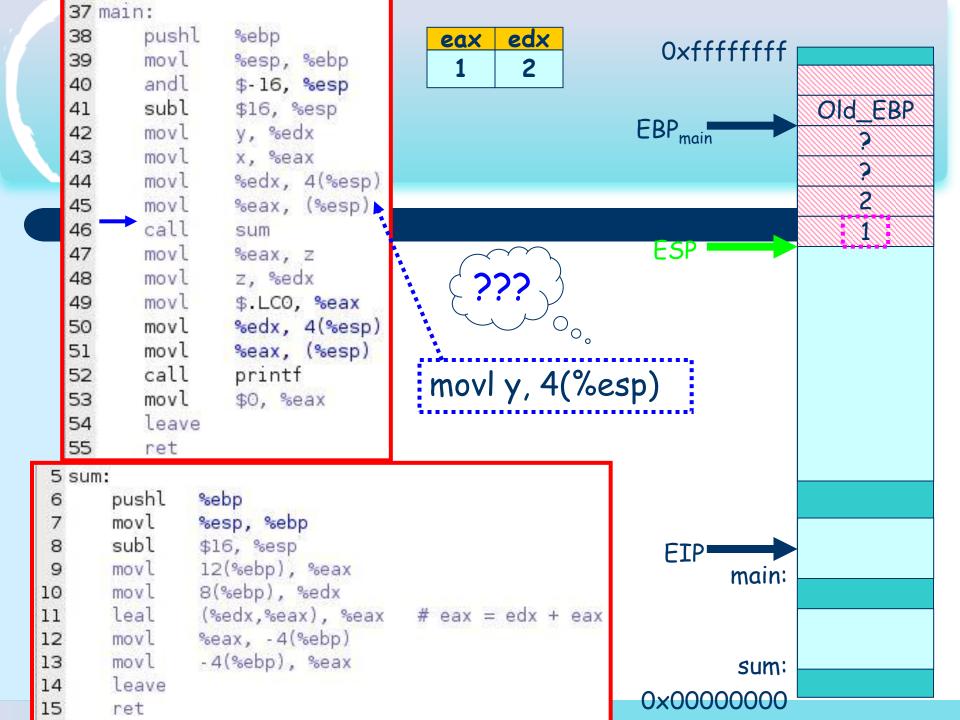


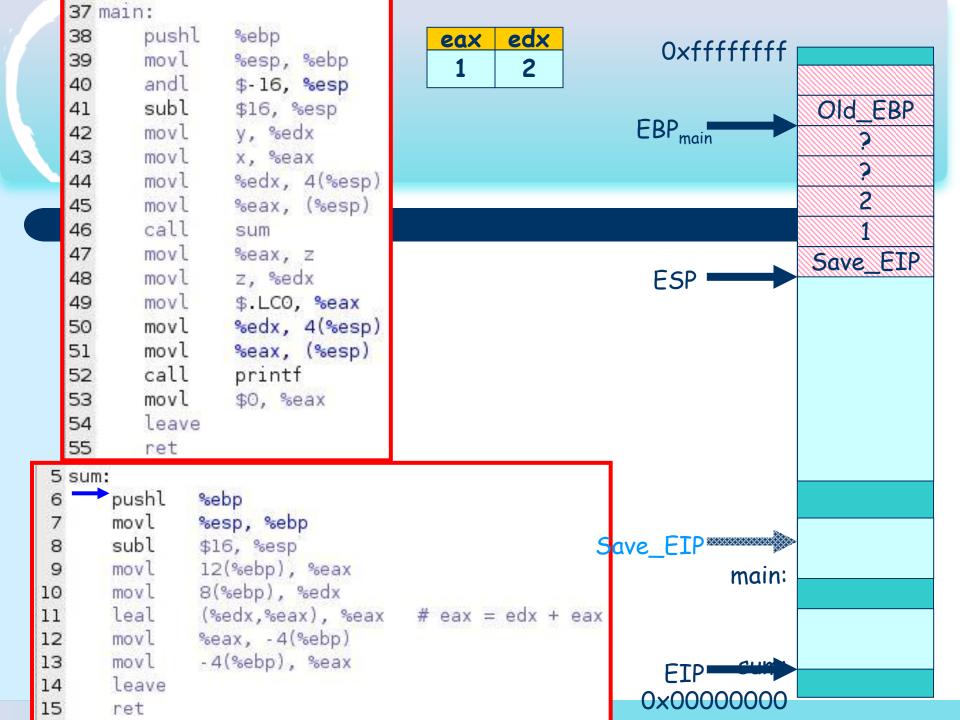


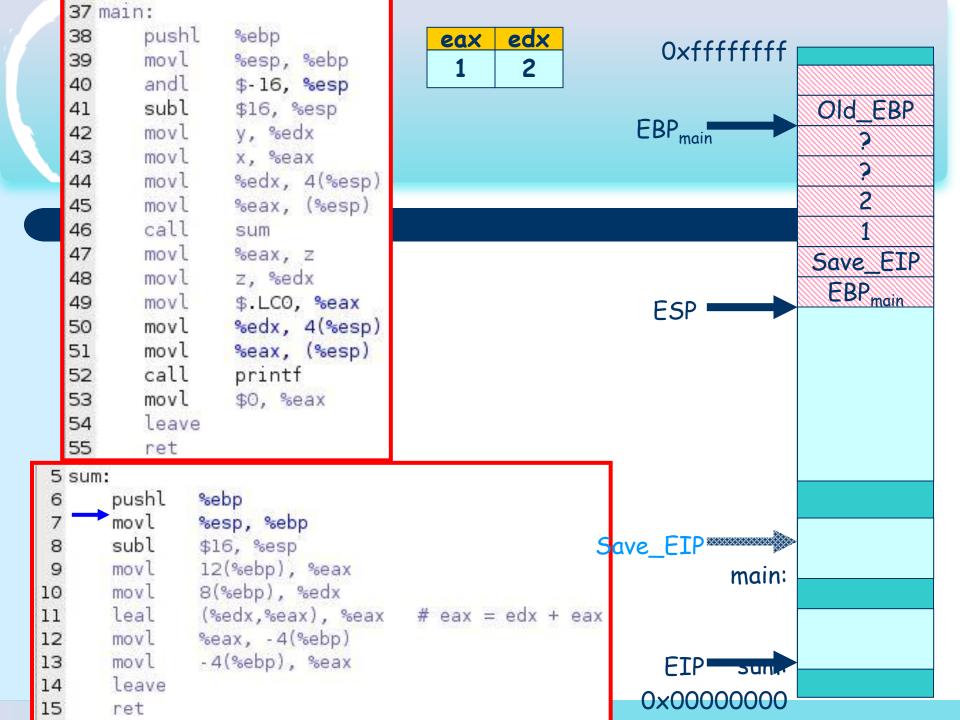


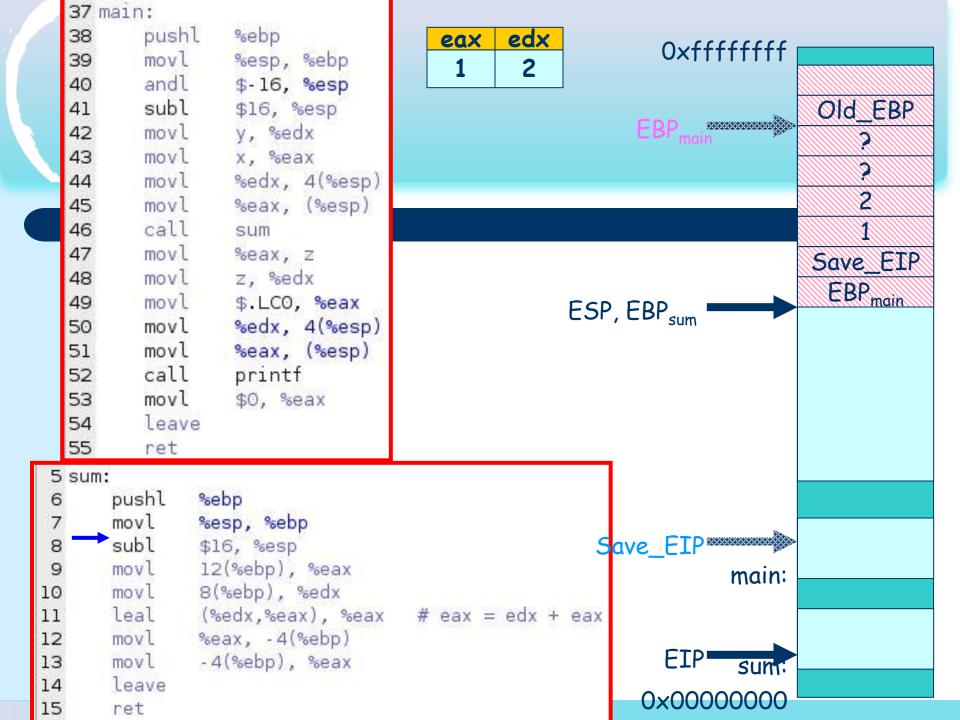


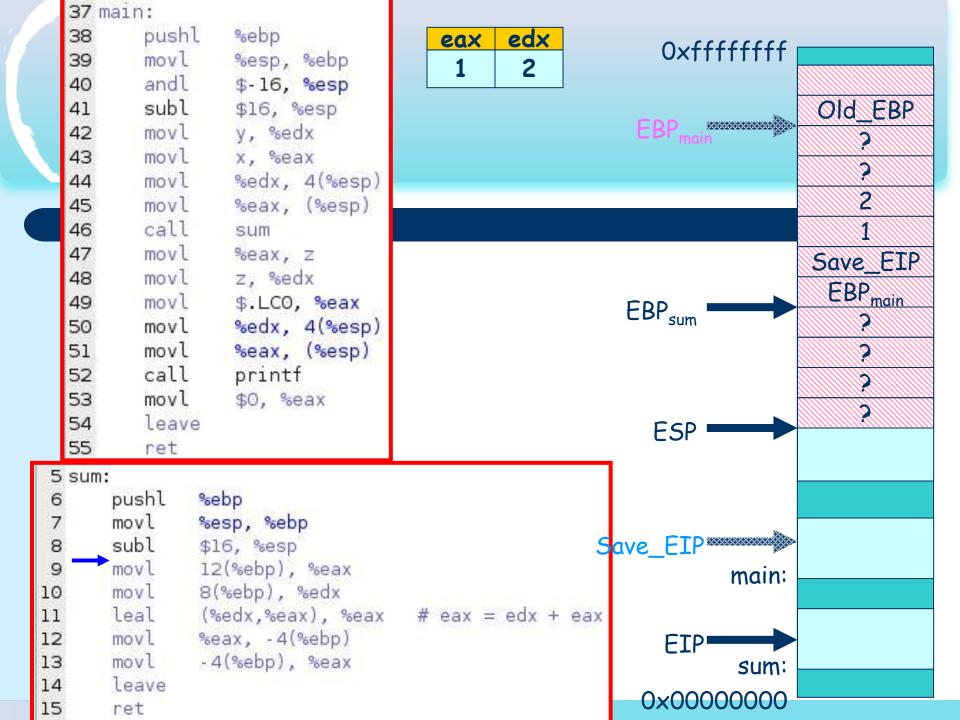


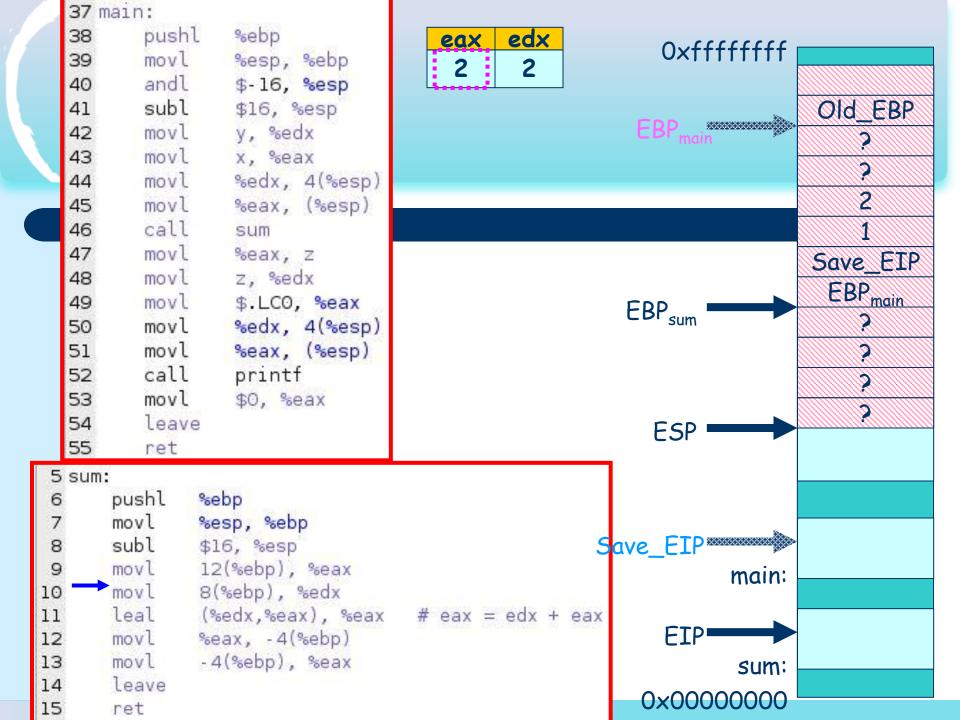






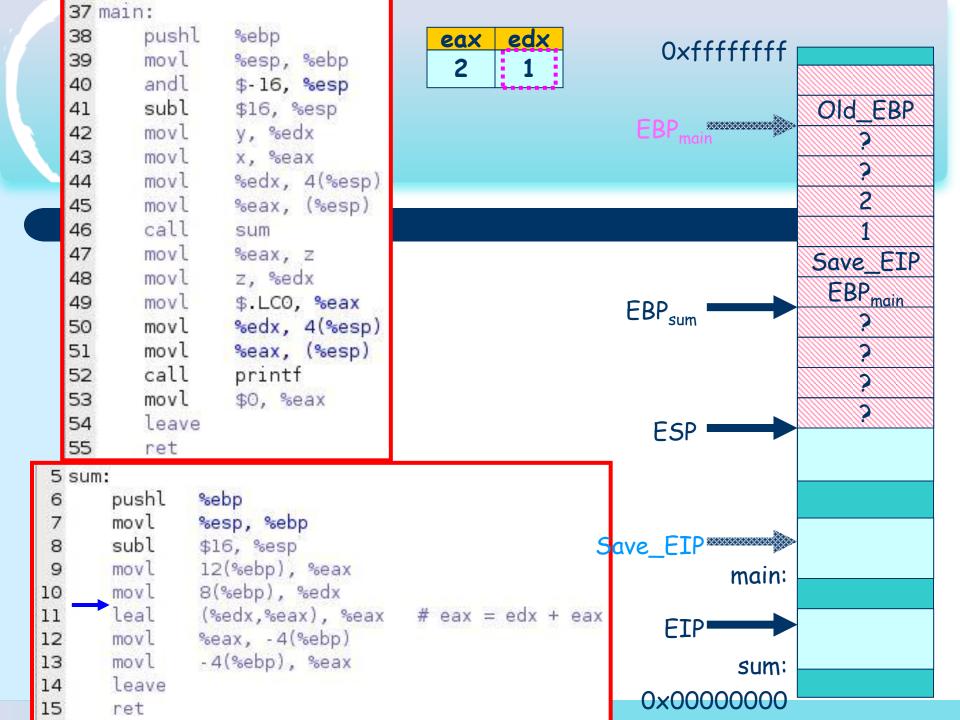






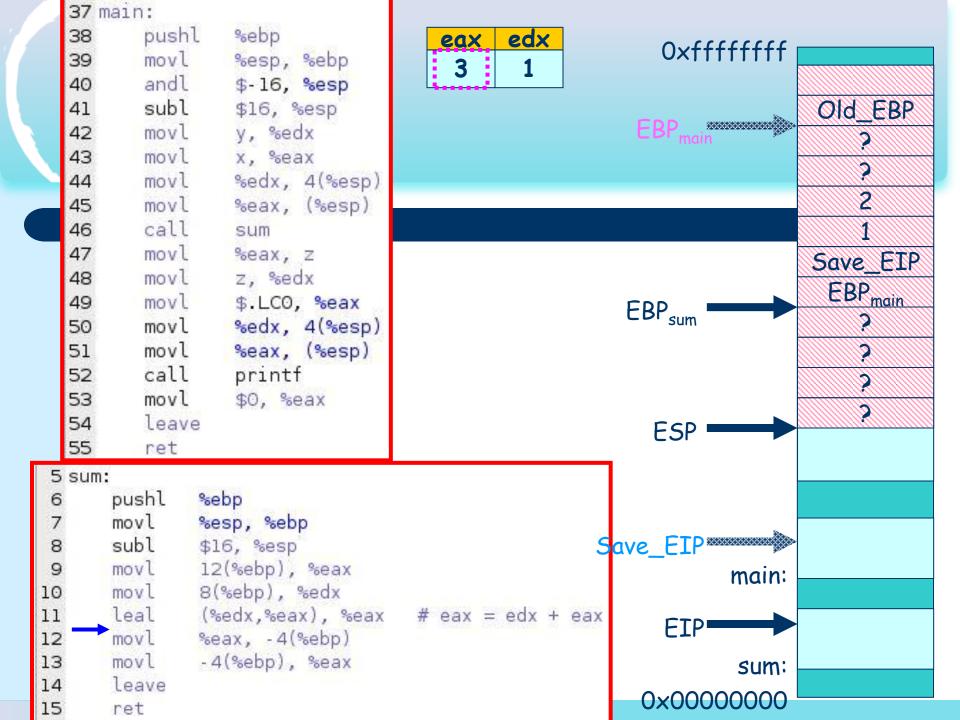
Is EBP necessary?

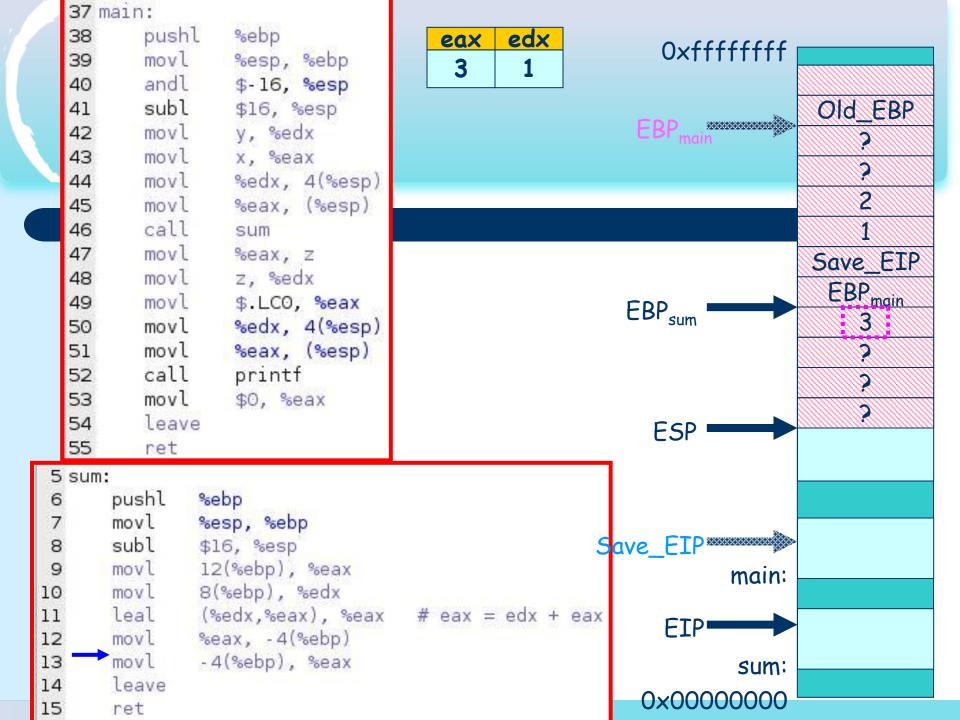
- Why use EBP?
 - accessing parameters
 - all EBP values saved in the stack form a trace of function calls
- Without EBP,
 - do function calls still work?
 - can you somehow obtain the trace of function calls?

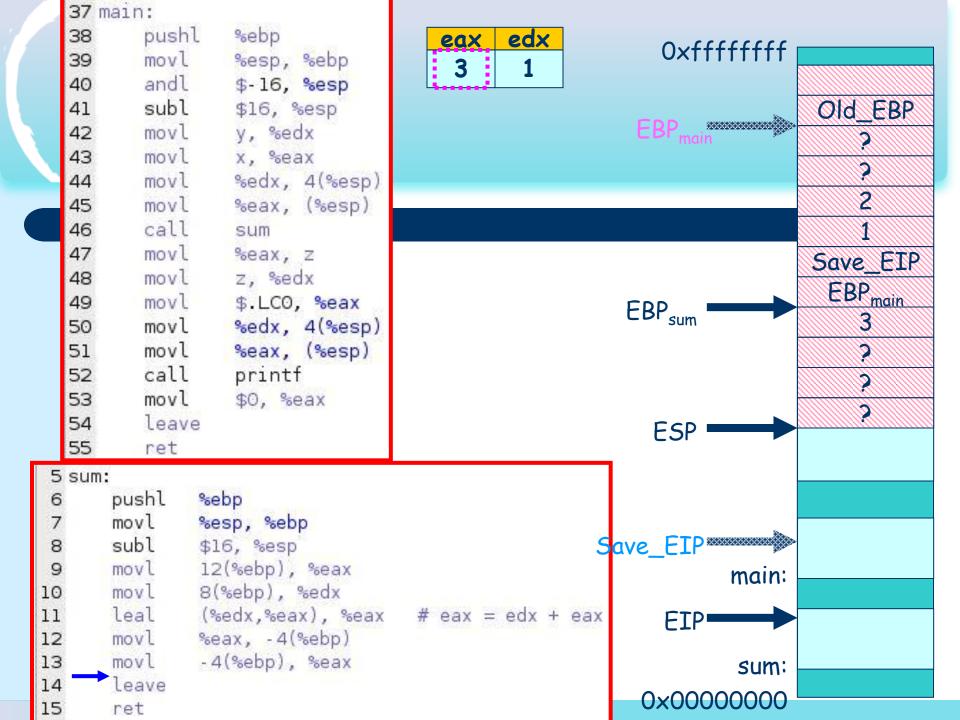


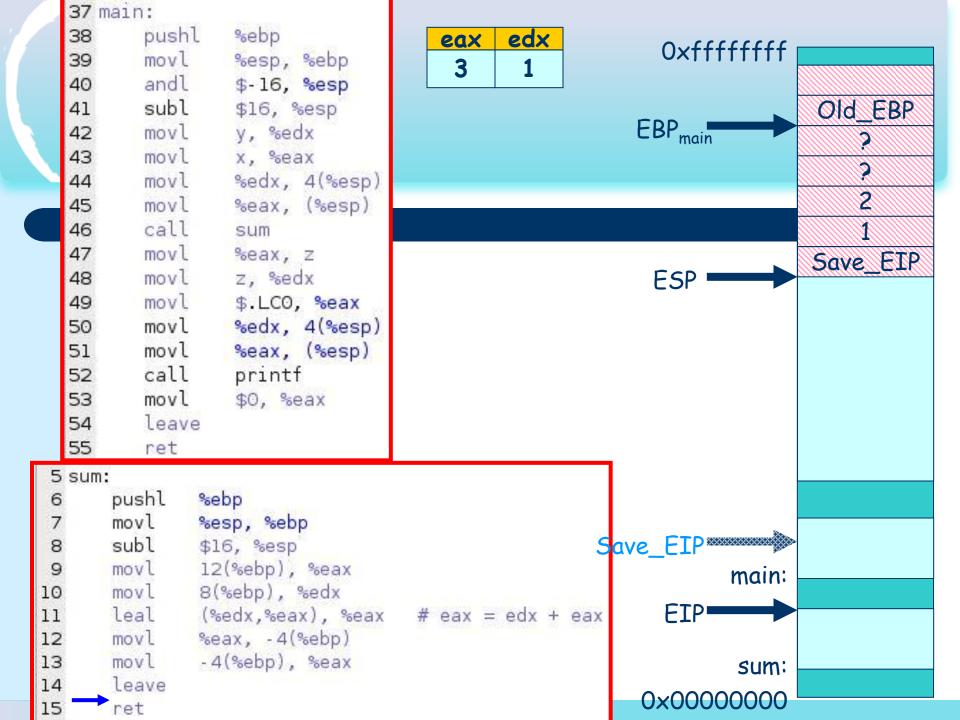
Why "lea", but not "add"?

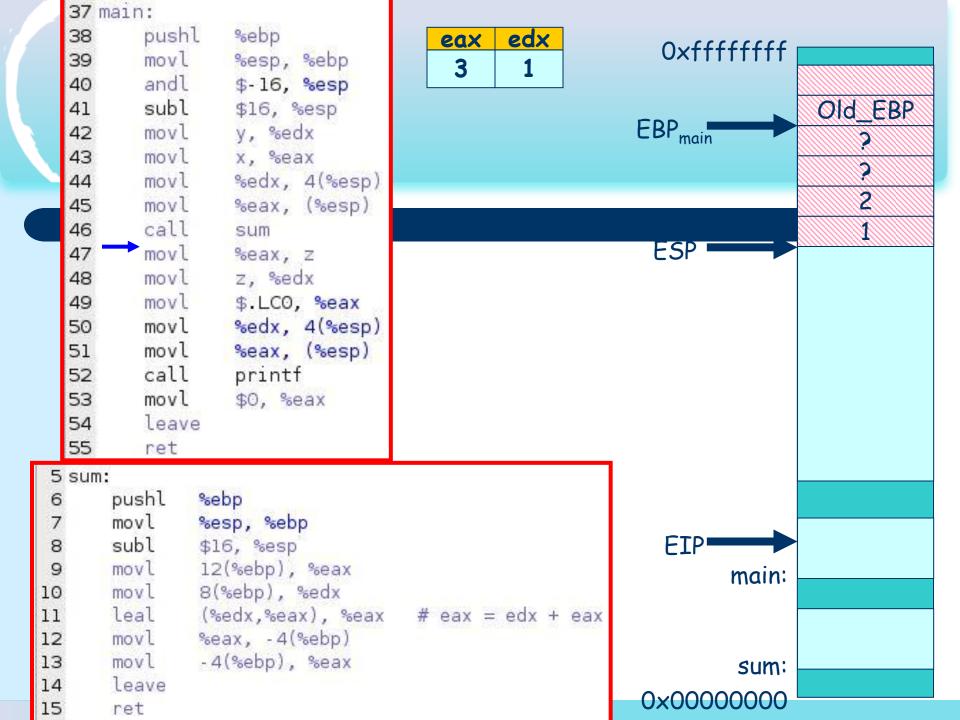
- leal (%edx, %eax), %eax
 - lea = load effective address
 - R[eax] = R[edx] + R[eax]
- Which is faster? Why?
- Design of machine-dependent optimization in a compiler requires deep insight into computer architecture.
 - This is systematic view.













Save_EIP EBP_{main} 1234

fun:

pushl %ebp

movl %esp, %ebp

....

strcpy(%esp, string)

Save_EIP main:

EIP fun:

0xffffffff



EBP_{fun} ???

333

0xffffffff

fun:

push! %ebp

movl %esp, %ebp

....

strcpy(%esp, string)

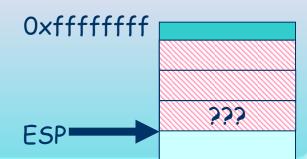
leave

Save_EIP main:

fun:

0x00000000

Example - buffer overflow attack (cont.)



fun:

push! %ebp

movl %esp, %ebp

....

strcpy(%esp, string)

leave

ret

EBP_{main}

Save_EIP main:

fun:

0x00000000

Example - buffer overflow attack (cont.)

Oxfffffff
ESP

fun:

pushl %ebp

movl %esp, %ebp

••••

strcpy(%esp, string)

leave

ret





fun:





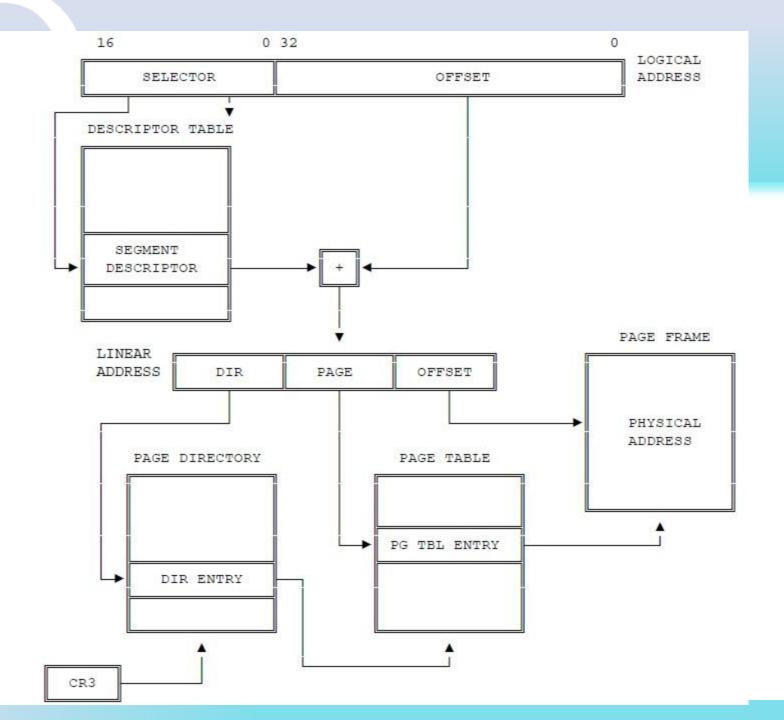
0x0000000

Try to attack this program!

```
#include "stdio.h"
 2 #include "string.h"
 4 void hacker(void) {
       printf("being hacked\n");
 5
 6
   void outputs(char *str) {
       char buffer[16];
       strcpy(buffer,str);
10
       printf("%s\n", buffer);
11
12 }
13
14 int main(int argc, char *argv[]) {
       outputs(argv[1]);
15
       return 0;
16
```

Hint: x86 is little-endian

 Recap - IA-32 memory management



• I/O & Peripherals

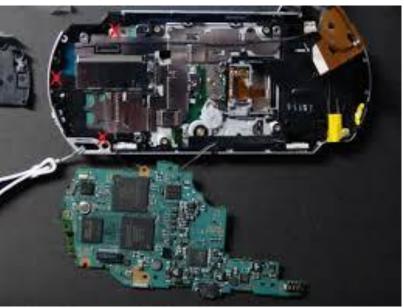
Peripherals

- Why peripherals?
 - CPU can only execute instructions.
 - Peripherals make computer more powerful.
- But how do they communicate with each other to finish a task?

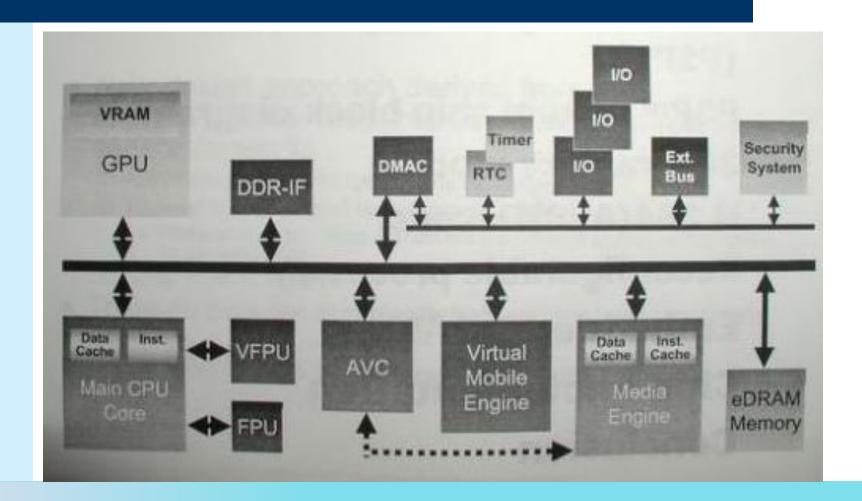
That is the topic of I/O system.

Example - SONY PSP





SONY PSP block diagram



Some processing units

Game processing unit

- CPU Core
 - MIPS R4000 32bit Core:1-333MHz
 - FPU, VFPU (Vector Unit)
- 3D graphics
 - 'Rendering Engine' + 'Surface Engine'
 - 2MByte eDRAM(VRAM):512bit/166MHz bus I/F
- eDRAM
 - Total capacity of 4Mbytes
 - One half for the game processing unit
 - The other half for the media processing unit
- Security block
 - Protect game/ audio-video contents.
 - Protect the PSP™ as a system.
- 1/0
 - Mobile DDR I/F
 - USB 2.0(Device)
 - Memory Stick[™]

Media processing unit

- Media Engine
 - MIPS R4000 32bit Core with FPU
 - 1-333MHz
- H.264 codec engine
 - H.264 hardware accelerator
- VME(Virtual Mobile Engine)
 - A reconfigurable processor to decode audio/video codec

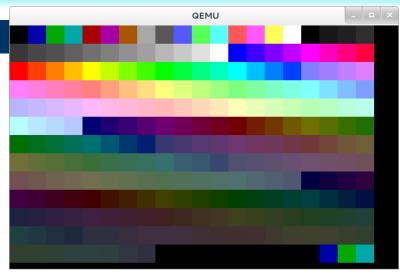
I/O addressing

- port-mapped I/O
 - use special instructions
 - separate address space with memory
 - more control signals for CPU and address bus
- memory-mapped I/O
 - use memory-accessing instructions
 - the same address space with memory
 - no extra control signals for CPU
 - some control signals for address bus

Port-mapped I/O example

```
3 #define PORT PIC MASTER 0x20
 4 #define PORT PIC SLAVE 0xA0
 5 #define IRQ SLAVE
 6
  /* 初始化8259中断控制器:
   * 硬件中断IRQ从32号开始,自动发送E0I */
 9 void
10 init intr(void) {
11
       out byte(PORT PIC MASTER + 1, 0xFF);
12
       out byte(PORT PIC SLAVE + 1 , 0xFF);
       out byte(PORT PIC MASTER, 0x11);
13
       out byte(PORT PIC MASTER + 1, 32);
14
       out byte(PORT_PIC_MASTER + 1, 1 << 2);
15
16
       out byte(PORT PIC MASTER + 1, 0x3);
17
       out byte(PORT PIC SLAVE, 0x11);
18
       out byte(PORT PIC SLAVE + 1, 32 + 8);
       out byte(PORT PIC SLAVE + 1, 2);
19
       out byte(PORT PIC SLAVE + 1, 0x3);
20
21
22
       out byte(PORT PIC MASTER, 0x68);
       out byte(PORT PIC MASTER, 0x0A);
23
       out byte(PORT PIC SLAVE, 0x68);
24
       out byte(PORT PIC SLAVE, 0x0A);
25
26 }
```

Memory-mapped I/O example



```
11 #define VMEM_ADDR ((uint8_t*)0xA0000)
12
13 extern uint8_t *vmem;
14
15 static inline void
16 draw_pixel(int x, int y, int color) {
17    assert(x >= 0 && y >= 0 && x < SCR_HEIGHT && SCR_WIDTH);
18    vmem[(x << 8) + (x << 6) + y] = color;
19 }</pre>
```

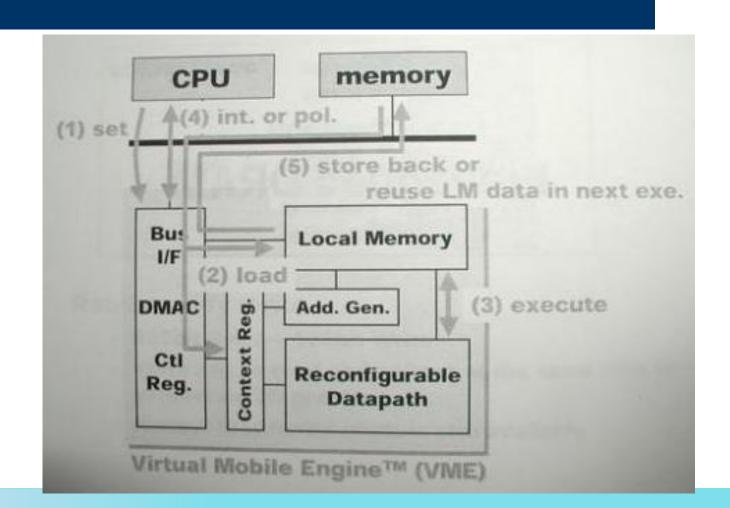
I/O communication

- Polling
 - CPU check the I/O status register repeatedly
 - until the peripheral is ready
- Interrupt
 - when peripheral is ready, or something happen to it, it sends a signal to CPU
 - it is a notification mechanism
- Is interrupt always more efficient then polling?

Polling the serial port

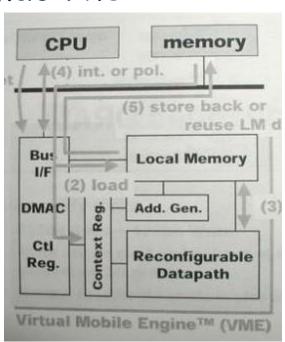
```
16 static inline int
17 serial_idle(void) {
18     return (in_byte(SERIAL_PORT + 5) & 0x20) != 0;
19 }
20
21 void
22 serial_printc(char ch) {
23     while (serial_idle() != TRUE);
24     out_byte(SERIAL_PORT, ch);
25 }
```

Typical VME operation in SONY PSP



Hardware is not mysterious!

- CPU = registers + controller + exectuion unit
- CPU controller = finite state machine
- So do peripherals!
 - have their own instructions
 - command words from CPU's view
 - generate their own control signals
 - use their own private registers



Hardware is not mysterious! (cont.)

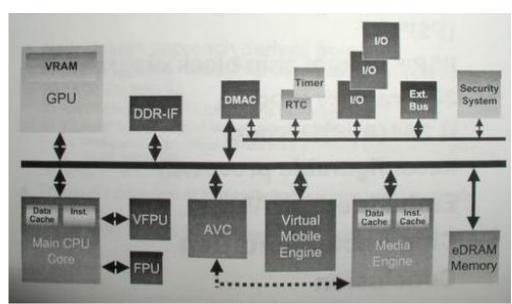
Use bus to connect them together

- bus = arbiter + MUX + DEMUX

- arbiter generates the control signals for MUX and

DEMUX

 This is the world of hardware!



• IA-32 Interrupt Processing

Interrupt

 A signal to inform the processor about something that must be handled immediately



Types

- hardware(external) interrupt
 - sent by a device
 - pressing a key
 - timer
- software(internal) interrupt
 - exception
 - divided by 0
 - page fault
 - trap instruction
 - int \$0x80

What to do?

- CPU must give response to the interrupt source as soon as possible
 - Why?
- What should CPU do?
 - save the current state
 - call the interrupt handler
 - restore the "current state"

Save the current state

- What should be saved?
 - EFLAGS
 - the IF bit in EFLAGS will be cleared if it is an hardware interrupt
 - other flags should be saved, too
 - CS
 - the handler may be in another code segment
 - EIP
 - return address
- They are done by hardware sequentially

Call the interrupt handler

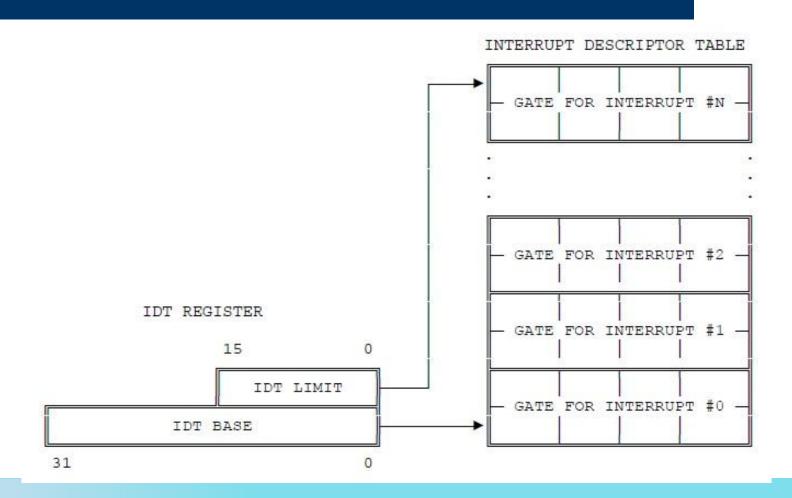
How to find the entry point of handler?

- IA-32
 - uses gate descriptors to indicate the entry point of an interrupt handler.
 - uses IDT to manage gate descriptors.

IDT

- interrupt descriptor table
 - an array of gate descriptors
 - pointed by IDTR
- There are 3 types of gate descriptor
 - interrupt gate
 - IF in EFLAGS will be cleared disable interrupt
 - trap gate
 - IF in EFLAGS remains unchanged
 - system gate
 - not used in Nanos

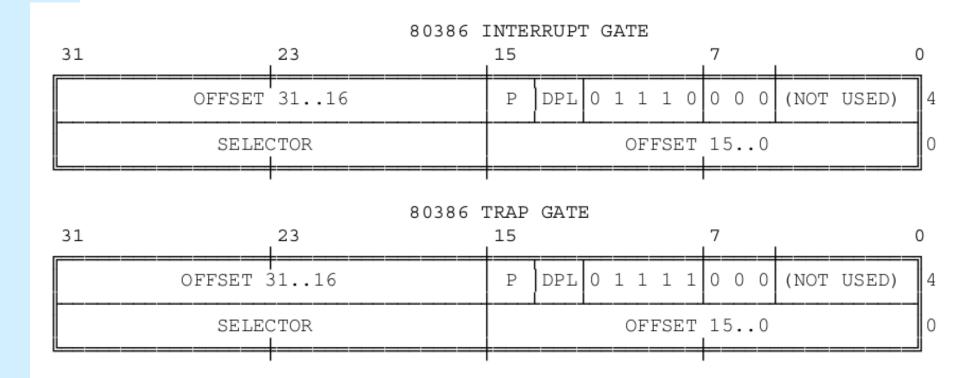
IDT (cont.)



IDT initialization

```
54 void init idt() {
      int i;
55
56
      /* 为了防止系统异常终止,所有irq都有处理函数(irq_empty)。 */
      for (i = 0; i < NR IRQ; i \leftrightarrow) {
57
          set trap(idt + i, SEG KERNEL CODE, (uint32 t)irg empty, DPL KERNEL);
58
59
60
61
      /* 设置异常的中断处理 */
      set trap(idt + 0, SEG KERNEL CODE, (uint32 t)vec0, DPL KERNEL);
62
      set trap(idt + 1, SEG KERNEL CODE, (uint32 t)vec1, DPL KERNEL);
63
      set trap(idt + 2, SEG KERNEL CODE, (uint32 t)vec2, DPL KERNEL);
64
      set trap(idt + 3, SEG_KERNEL_CODE, (uint32_t)vec3, DPL_KERNEL);
65
      set_trap(idt + 4, SEG_KERNEL_CODE, (uint32_t)vec4, DPL_KERNEL);
66
      set trap(idt + 5, SEG KERNEL CODE, (uint32 t)vec5, DPL KERNEL);
67
      set trap(idt + 6, SEG KERNEL CODE, (uint32 t)vec6, DPL KERNEL);
68
69
      set trap(idt + 7, SEG KERNEL CODE, (uint32 t)vec7, DPL KERNEL);
      set trap(idt + 8, SEG KERNEL CODE, (uint32 t)vec8, DPL_KERNEL);
70
      set_trap(idt + 9, SEG_KERNEL_CODE, (uint32_t)vec9, DPL_KERNEL);
71
      set trap(idt + 10, SEG KERNEL CODE, (uint32 t)vec10, DPL KERNEL);
72
73
      set trap(idt + 11, SEG KERNEL CODE, (uint32 t)vec11, DPL KERNEL);
      set trap(idt + 12, SEG KERNEL CODE, (uint32 t)vec12, DPL KERNEL);
74
      set trap(idt + 13, SEG KERNEL CODE, (uint32 t)vec13, DPL KERNEL);
75
76
      /* 设置外部中断的处理 */
77
      set intr(idt + 32, SEG KERNEL CODE, (uint32 t)irq0, DPL KERNEL);
78
79
80
      /* 写入IDT */
      save idt(idt, sizeof(idt));
81
82 }
```

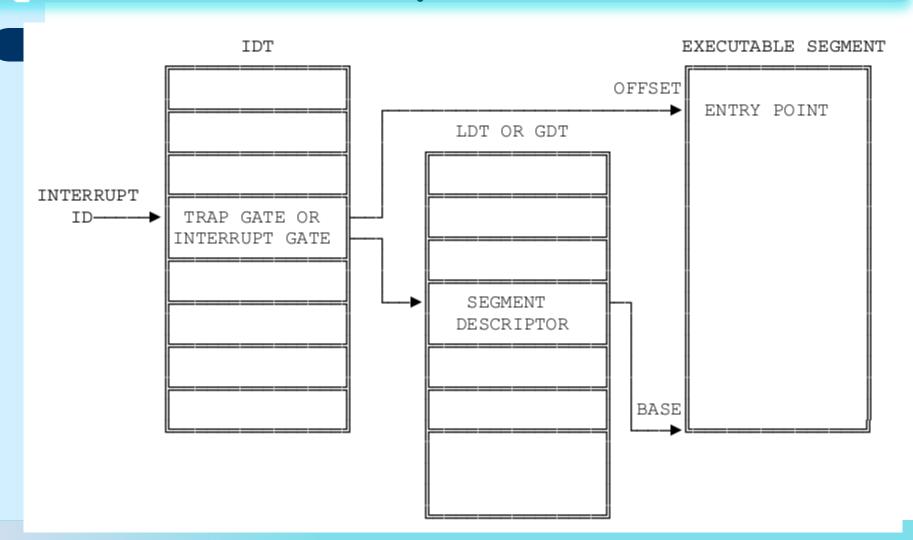
Gate descriptors



Call the interrupt handler (cont.)

- How to find the desired gate descriptor in IDT?
 - use an #id for each interrupt to index in IDT
- Where does #id come from?
 - hardware interrupt: I8259A PIC
 - software interrupt:
 - exception: CPU
 - trap instruction
 - int \$0x80

Call the interrupt handler (cont.)



Restore the "current state"

- Use the "iret" instruction to return from interrupt handling
 - restore EIP
 - pop CS
 - restore EFLAGS

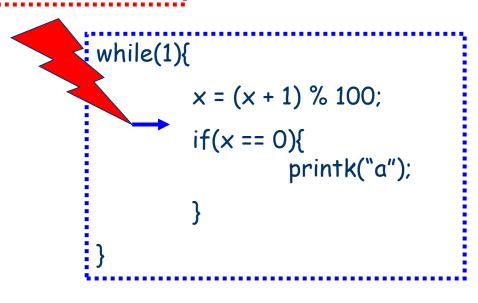
More information is needed

- The "current state" saved by hardware is not sufficient
- GPRs should be saved
 - interrupt is unpredictable for CPU
 - why not save them in a function call?
- #irq should be saved
 - to implement interrupt distribution
 - in LabO, external(hardware) interrupts have#irq >= 1000
 - easy to distinguish

ESP 12

Example

time expires



12 EFLAGS CS EIP

ESP

12 EFLAGS CS EIP

ESP!

```
中断处理函数会在IDT中为相应的中断/异常设置处理程序
 3 # 中断/异常的行为参见i386手册
  .globl vec0; vec0: pushl $0; jmp asm_do_irq
  .globl vec1; vec1: pushl $1; jmp asm do irq
  .globl vec2; vec2: pushl $2; jmp asm do irq
  .globl vec3; vec3: pushl $3; jmp asm_do_irq
  .globl vec4; vec4: pushl $4; jmp asm_do_irq
  .globl vec5; vec5: pushl $5; jmp asm do irq
  .globl vec6; vec6: pushl $6; jmp asm do irq
11 .globl vec7; vec7: pushl $7; jmp asm_do_irq
12 .globl vec8; vec8: pushl $8; jmp asm_do_irq
  .globl vec9; vec9: pushl $9; jmp asm do irq
  .globl vec10; vec10: pushl $10; jmp asm do irq
  .globl vec11; vec11: pushl $11; jmp asm do irq
16 .globl vec12; vec12: pushl $12; jmp asm_do_irq
17 .globl vec13; vec13: pushl $13; jmp asm do irq
18
  .globl irq0; irq0: pushl $1000; jmp asm do irq
20
21 .globl irq empty; irq empty: pushl $-1; jmp asm do irq
```

```
12
EFLAGS
CS
EIP
1000
```

```
ESP
24 .globl asm_do_irq
25 .extern irq handle push all GPRs
26 asm do irq:
       pushal
27
28
       pushl %esp
29
                            # ???
       call irq_handle
30
       addl $4, %esp
31
32
33
       popal
       addl $4, %esp
34
       iret
```

```
24 .globl asm do irq
25 .extern irq_handle
26 asm do irq:
                  555
       pushal
27
28
       pushl %esp
29
                              ???
       call irq_handle
30
       addl $4, %esp
31
32
       popal
33
       addl $4, %esp
34
```

12
EFLAGS
C5
EIP
1000
EAX
EBX
ECX
EDX
old_ESP
EBP
ESI
EDI

```
24 .globl asm do irq
25 .extern irq handle
26 asm do irq:
       pushal
27
28
       pushl %esp
29
                             # ???
       call irq_handle
30
       addl $4, %esp
31
                             ESP
32
       popal
33
       addl $4, %esp
34
```

12	
EFLAGS	
CS	
EIP	
1000	
EAX	
EBX	
ECX	
EDX	
old_ESF	
EBP	
ESI	
EDI	
ESP???	

```
18 void
19 irg handle(struct TrapFrame *tf) {
       if(tf->irq < 1000) {
21
           if(tf->irq == -1) {
               printk("%s, %d: Unhandled exception!\n",
23
                          FUNCTION , LINE );
24
25
           else {
               ፡ {
printk("%s, %d: Unexpected exception #%d!\n",
26
27
                          FUNCTION , LINE , tf->irq);
28
29
           assert(0);
30
31
32
       if (tf->irg == 1000) {
33
           do timer():
       } else if (tf->irq == 1001) {
34
           uint32 t code = in byte(0x60);
35
           uint32 t val = in \overline{b}yte(0x61);
36
           out byte(0x61, val | 0x80);
37
           out byte(0x61, val);
38
39
           printk("%s, %d: key code = %x\n",
40
41
                      FUNCTION , LINE , code);
           do keyboard(code);
42
43
       } else {
44
           assert(0);
45
46 }
```

12 **EFLAGS** CS. EIP 1000 EAX EBX ECX EDX old_ESP EBP ESI EDI ESP??? save EIP

```
24 .globl asm do irq
25 .extern irq handle
26 asm do irq:
       pushal
27
28
       pushl %esp
29
                             # ???
       call irq_handle
30
       addl $4, %esp
31
32
                       remove "ESP?
       popal
33
       addl $4, %esp
34
```

12
EFLAGS
C5
EIP
1000
EAX
EBX
ECX
EDX
old_ESP
EBP
ESI
EDI
ESP???

```
24 .globl asm do irq
25 .extern irq handle
26 asm do irq:
       pushal
27
28
       pushl %esp
29
                             # ???
       call irq_handle
30
       addl $4, %esp
31
32
                       restore GPRs
       popal
33
       addl $4, %esp
34
```

12
EFLAGS
C5
EIP
1000
EAX
EBX
ECX
eDX old ESP
EBP
ESI
EDI

```
12
EFLAGS
CS
EIP
1000
```

```
ESP
24 .globl asm do irq
25 .extern irq handle
26 asm do irq:
       pushal
27
28
       pushl %esp
29
                            # ???
       call irq_handle
30
       addl $4, %esp
31
32
                       remove #irq
33
       popal
34
       addl $4, %esp
       iret
```

12 EFLAGS CS EIP

```
ESPI
```

```
24 .globl asm do irq
25 .extern irq handle
26 asm do irq:
        pushal
27
28
        pushl %esp
29
                             # ???
        call irq_handle
30
       addl $4, %esp
31
32
                       restore the
33
        popal
                      "current state"
       addl $4, %esp stored by hardware
34
        iret
```



we are back!!!

Exceptions

They are triggered ONLY by your wrong code

```
.globl vec0; vec0: pushl:$0; jmp asm_do_irq
 5 .globl vec1; vec1: pushl:$1; jmp asm do irq
  .globl vec2; vec2: pushl $2; jmp asm do irq
  .globl vec3; vec3: pushl $3; jmp asm do irq
  .globl vec4; vec4: pushl $4; jmp asm do irq
  .globl vec5; vec5: pushl:$5; jmp asm do irq
10 .globl vec6; vec6: pushl $6; jmp asm_do_irq
  .globl vec7; vec7: pushl;$7; jmp asm do irq
12 .globl vec8; vec8: pushl:$8; jmp asm_do_irq
  .globl vec9; vec9: pushl $9; jmp asm_do_irq
  .globl vec10; vec10: pushl $10; jmp asm_do_irq
15 .globl vec11; vec11: pushl $11; jmp asm_do_irq
16 .globl vec12; vec12: pushl $12; jmp asm_do_irq
17 .globl vec13; vec13: pushl $13; jmp asm do irq
```



The magic of Interrupt

- DMA
- for the game in LabO
 - frame-based rendering
 - drive the logical time
 - get inputs from player
- for OS
 - context switching (very significant)
- refer to INTEL 80386 PROGRAMMER'S REFERENCE MANUAL for more details about interrupts