

Critical region

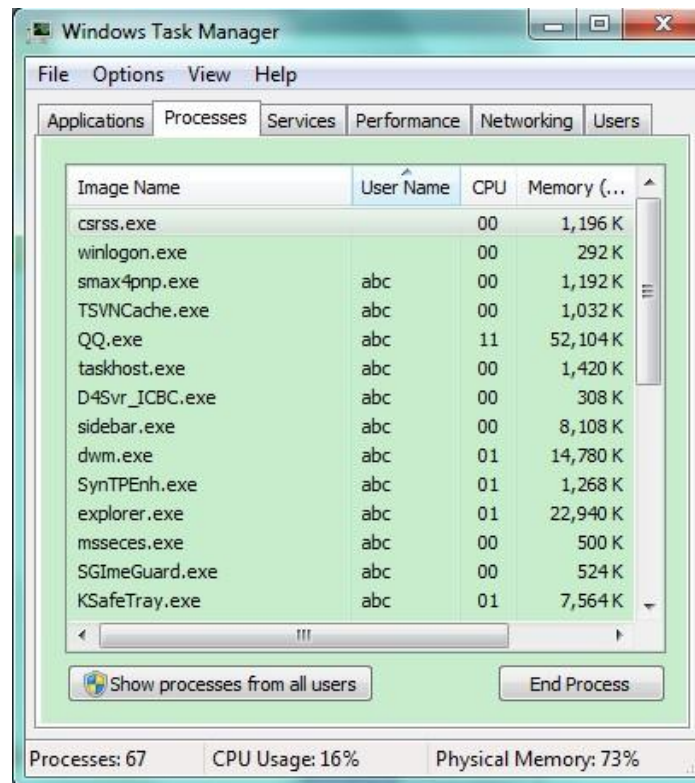
- Critical region
- Semaphore
- Locking



-
- Critical region

With context switch

- different processes can coexist





IPC

- Rather than running separately, processes may cooperate with each other.
 - Xunlei, chrome
- Inter-Process Communication
 - extension: network, database transaction
- Some resources may be shared
 - memory, files

New problem

- summing from 1 to 100

sum	0
-----	---

temp	???
------	-----

```
void sum_odd(){  
    int i = 1, temp;  
    for (; i < 100; i += 2){  
        temp = sum;  
        temp += i;  
        sum = temp;  
    }  
}
```

temp	???
------	-----

```
void sum_even(){  
    int i = 2, temp;  
    for (; i < 100; i += 2){  
        temp = sum;  
        temp += i;  
        sum = temp;  
    }  
}
```


New problem

- summing from 1 to 100

sum	0
-----	---

temp	0
------	---

```
void sum_odd(){  
    int i = 1, temp;  
    for (; i < 100; i += 2){  
        temp = sum;  
        temp += i;  
        sum = temp;  
    }  
}
```



temp	???
------	-----

```
void sum_even(){  
    int i = 2, temp;  
    for (; i < 100; i += 2){  
        temp = sum;  
        temp += i;  
        sum = temp;  
    }  
}
```


New problem

- summing from 1 to 100

sum	0
-----	---

temp	1
------	---

```
void sum_odd(){  
    int i = 1, temp;  
    for (; i < 100; i += 2){  
        temp = sum;  
        temp += i;  
        sum = temp;  
    }  
}
```



temp	???
------	-----

```
void sum_even(){  
    int i = 2, temp;  
    for (; i < 100; i += 2){  
        temp = sum;  
        temp += i;  
        sum = temp;  
    }  
}
```

New problem

- summing from 1 to 100

sum	0
-----	---

temp	1
------	---

temp	???
------	-----

```
void sum_odd(){  
    int i = 1, temp;  
    for (; i < 100; i += 2){  
        temp = sum;  
        temp += i;  
        sum = temp;  
    }  
}
```

```
void sum_even(){  
    int i = 2, temp;  
    for (; i < 100; i += 2){  
        temp = sum;  
        temp += i;  
        sum = temp;  
    }  
}
```


New problem

- summing from 1 to 100

sum	0
-----	---

temp	1
------	---

```
void sum_odd(){  
    int i = 1, temp;  
    for (; i < 100; i += 2){  
        temp = sum;  
        temp += i;  
        sum = temp;  
    }  
}
```

temp	???
------	-----

```
void sum_even(){  
    int i = 2, temp;  
    → for (; i < 100; i += 2){  
        temp = sum;  
        temp += i;  
        sum = temp;  
    }  
}
```

New problem

- summing from 1 to 100

sum	0
-----	---

temp	1
------	---

```
void sum_odd(){  
    int i = 1, temp;  
    for (; i < 100; i += 2){  
        temp = sum;  
        temp += i;  
        sum = temp;  
    }  
}
```

temp	0
------	---

```
void sum_even(){  
    int i = 2, temp;  
    for (; i < 100; i += 2){  
        temp = sum;  
        temp += i;  
        sum = temp;  
    }  
}
```

New problem

- summing from 1 to 100

sum	0
-----	---

temp	1
------	---

```
void sum_odd(){  
    int i = 1, temp;  
    for (; i < 100; i += 2){  
        temp = sum;  
        temp += i;  
        sum = temp;  
    }  
}
```

temp	2
------	---

```
void sum_even(){  
    int i = 2, temp;  
    for (; i < 100; i += 2){  
        temp = sum;  
        temp += i;  
        sum = temp;  
    }  
}
```

New problem

- summing from 1 to 100

temp	1
------	---

sum	2
-----	---

temp	2
------	---

```
void sum_odd(){  
    int i = 1, temp;  
    for (; i < 100; i += 2){  
        temp = sum;  
        temp += i;  
        sum = temp;  
    }  
}
```

```
void sum_even(){  
    → int i = 2, temp;  
    for (; i < 100; i += 2){  
        temp = sum;  
        temp += i;  
        sum = temp;  
    }  
}
```

New problem

- summing from 1 to 100

temp	1
------	---

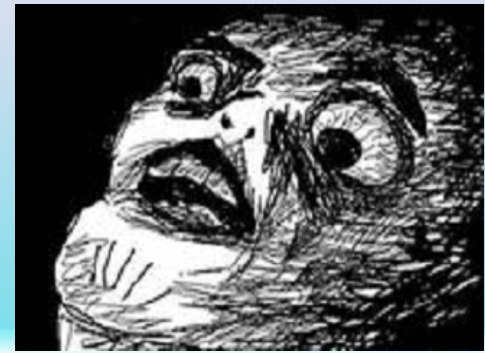
sum	2
-----	---

temp	2
------	---

```
void sum_odd(){  
    int i = 1, temp;  
    for (; i < 100; i += 2){  
        temp = sum;  
        temp += i;  
        sum = temp;  
    }  
}
```

```
void sum_even(){  
    int i = 2, temp;  
    for (; i < 100; i += 2){  
        temp = sum;  
        temp += i;  
        sum = temp;  
    }  
}
```

New problem



- summing from 1 to 100

sum	1
-----	---

temp	1
------	---

```
void sum_odd(){  
    → int i = 1, temp;  
    for (; i < 100; i += 2){  
        temp = sum;  
        temp += i;  
        sum = temp;  
    }  
}
```

lost update

temp	2
------	---

```
void sum_even(){  
    int i = 2, temp;  
    for (; i < 100; i += 2){  
        temp = sum;  
        temp += i;  
        sum = temp;  
    }  
}
```



Why?

- allow concurrent access
 - shared by two or more processes
- allow update
 - allow write operation
- not atomic
 - remain inconsistent while interrupted
- without concurrency control
 - access critical regions casually



Some mechanisms are needed

- allow concurrent access?
 - back to uni-tasking? ❌
- allow update?
 - read only? ❌
- not atomic?
 - make the process uninterruptable
- without concurrency control?
 - control the actions of accessing critical regions



Atomicity & Consistency

- atomicity
 - “all or nothing”
 - when seen by others, it performs either all operations or no operations in the critical region
- consistency
 - this is actually what correctness means
 - induction on the consistency of states
 - consistency → inconsistency → consistency
 - inconsistent states should not be exposed to others
- no atomicity → no consistency

Consistency & Assertion

- Consistency can be used as the condition of assertion.

`assert(temp == sum);`

```
void sum_odd(){  
    int i = 1, temp;  
    for (; i < 100; i += 2){  
        temp = sum;  
        temp += i;  
        sum = temp;  
    }  
}
```

```
void sum_even(){  
    int i = 2, temp;  
    for (; i < 100; i += 2){  
        temp = sum;  
        temp += i;  
        sum = temp;  
    }  
}
```

`assert(temp - i == sum);`

Another example

```
void list_del(ListHead *data) {  
    ListHead *prev = data->prev;  
    ListHead *next = data->next;  
    assert(prev == NULL || prev->next == data);  
    assert(next == NULL || next->prev == data);  
    if (prev != NULL) prev->next = next;  
    if (next != NULL) next->prev = prev;  
}
```

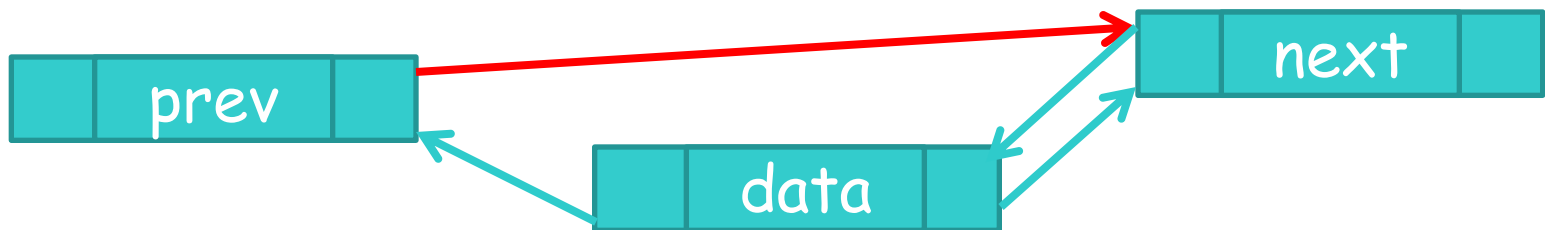
check consistency



Another example

```
void list_del(ListHead *data) {  
    ListHead *prev = data->prev;  
    ListHead *next = data->next;  
    assert(prev == NULL || prev->next == data);  
    assert(next == NULL || next->prev == data);  
    if (prev != NULL) prev->next = next;  
    if (next != NULL) next->prev = prev;  
}
```

check consistency

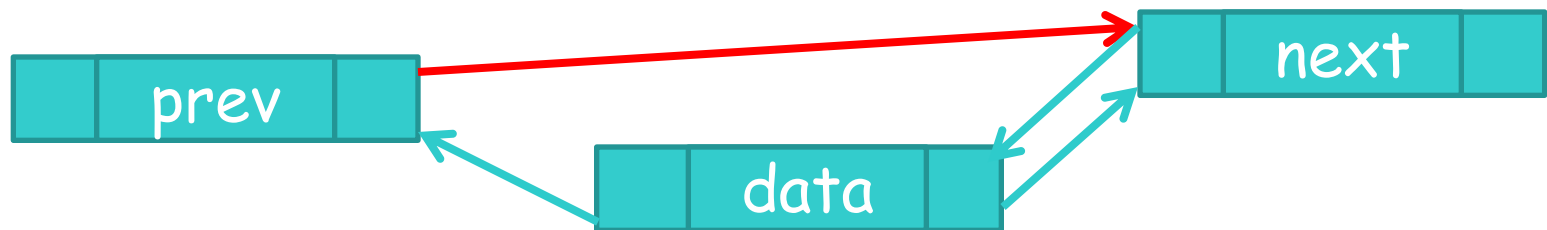


Another example

```
void list_del(ListHead *data) {  
    ListHead *prev = data->prev;  
    ListHead *next = data->next;  
    assert(prev == NULL || prev->next == data);  
    assert(next == NULL || next->prev == data);  
    if (prev != NULL) prev->next = next;  
    if (next != NULL) next->prev = prev;  
}
```

check
consistency

context switch
happens



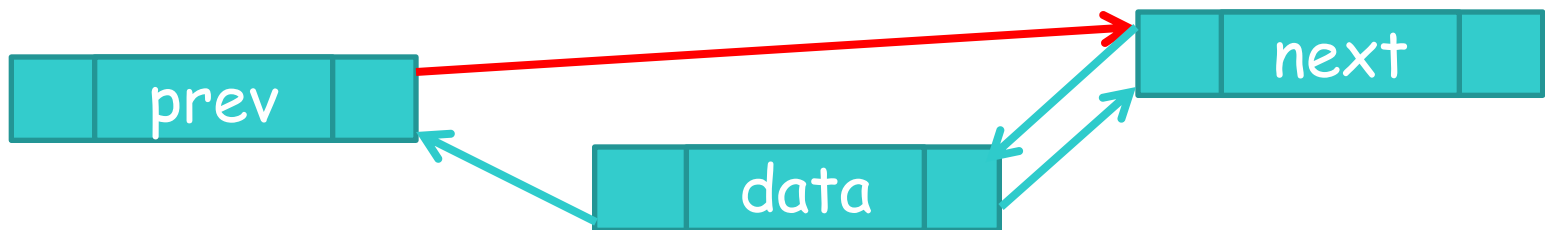
Another example

```
void list_del(ListHead *data) {  
    ListHead *prev = data->prev;  
    ListHead *next = data->next;  
    assert(prev == NULL || prev->next == data);  
    assert(next == NULL || next->prev == data);  
    if (prev != NULL) prev->next = next;  
    if (next != NULL) next->prev = prev;  
}
```



check
consistency

context switch
happens



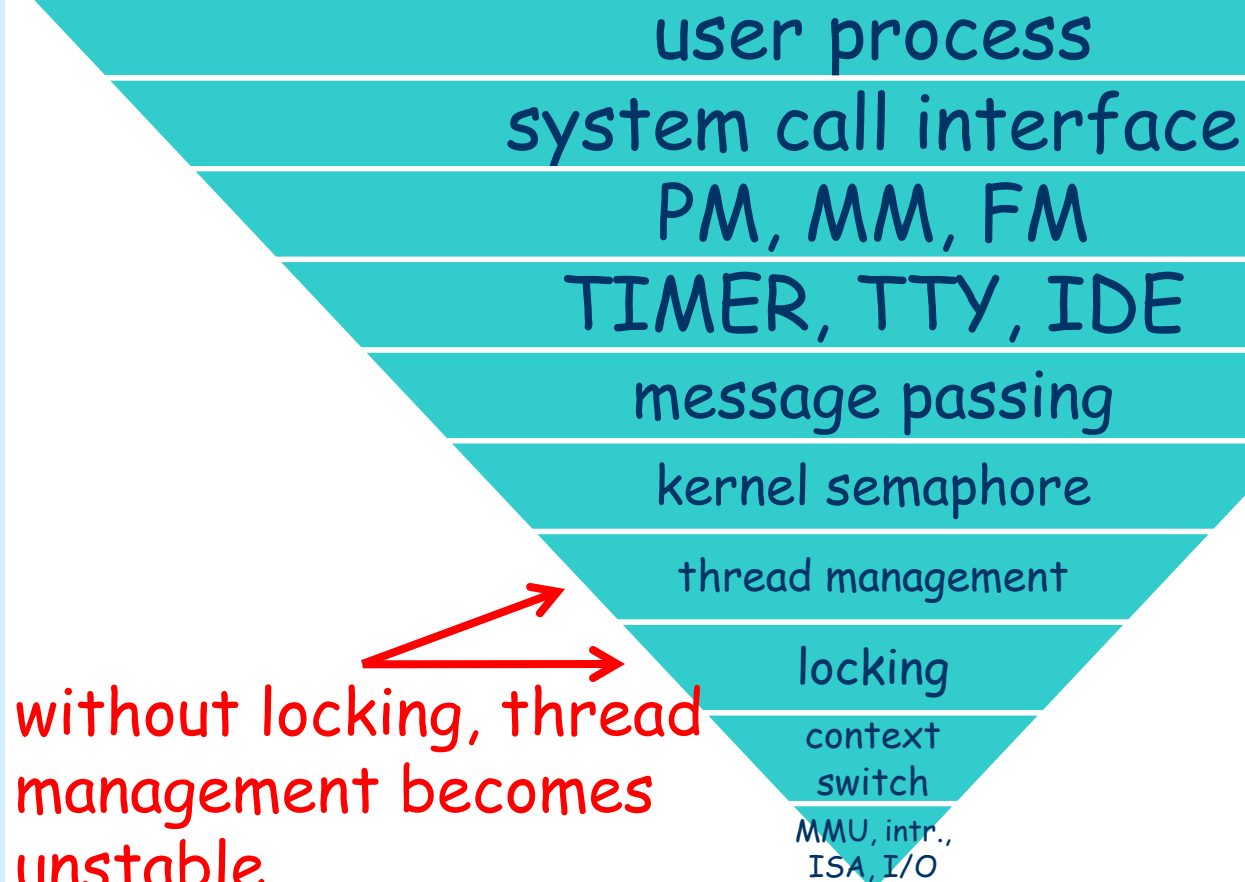


Protect critical regions

- concurrency control
 - goal: guarantee atomicity & consistency
- Without protection, your code will get stuck in mysterious behaviors.
 - nothing is impossible
 - your program will become an awful mess!
 - VERY HARD to backtrack, as well as reproduce

[illegible][illegible]

The architecture of Nanos





linux发行版不支持win8快速启动功能会导致数据丢失

- the “most” largest critical region

Windows 8的新功能Fast Startup能让双系统启动时更快，然而其负面作用是，当用户加载第二款系统（通常是Linux）时本地存储的数据处于风险之中。原因很简单，在双启动配置下，Fast Startup选项并不会完全关闭Windows 8，只是强制它进入休眠模式，从而能在用户退出另一个系统重新启动Windows 8时获得更快的加载速度。

由于是在休眠状态下，Windows 8会继续存储Windows会话信息，包括FAT和NTFS分区文件系统数据。当用户重新切换回Windows 8时，一些文件已经被删去或重写了。

ntfs-3g FUSE文件系统驱动开发人员发现，Linux会试图在Windows分区写入数据，有时会重写存储在memory image中的系统文件，当切换回Windows 8时，Fast Startup只重新加载了系统，却不能重新找回丢失的数据。

为了解决这个问题，ntfs-3g FUSE开发人员不得不开发了新版，让Linux系统将NTFS分区设定为只读。

但是，大多数Linux发行版，包括Ubuntu、Debian和openSUSE等目前都没有解决这个问题，所以强制阻止操作系统重写Fast Startup数据是当前最好的解决办法。



How to discover critical regions?

- (1) the same memory location may access by more than one process
 - usually global variables, or fixed pointers
- (2) one of them may be write operation
- (1) + (2) = data race



Why?

- allow concurrent access
 - shared by two or more processes
- allow update
 - allow write operation
- not atomic
 - remain inconsistent while interrupted
- without concurrency control
 - access critical regions casually



More data races

- process → execution flow
 - process1 v.s. process2
 - process v.s. interrupt
 - process v.s. signal
 - interrupt1 v.s. interrupt2
 - interrupt1 v.s. itself
 - re-enterable & unre-enterable
- more tricky in SMP
 - “real” simultaneousness



Simplification in Nanos

- All hardware interrupts are unre-enterable.
- We do not need to perform extra protections for interrupt procedures.
 - Do NOT enable interrupt during the interrupt procedures.
 - Or your system may get stuck in mysterious behaviors.



Protect critical regions (cont.)

- Do not let two processes (execution flows) enter the same critical region.
- based on sleeping
 - semaphore
- based on waiting
 - locking

- 
-
- Semaphore

Semaphore

- busy-then-sleep





Semaphore (cont.)

- maintain the number of resource available
 - but not the specific resource
- $\text{token} > 0$
 - token units of resource are available
- $\text{token} == 0$
 - no resource remains
 - processes must queue for the resource



P-V operation

- manage the semaphore
- P operation
 - request for one unit of resource
 - block if no resource is available
- V operation
 - release one unit of resource
 - wake up a blocking process queuing for the same resource, if any



Usage

- counting semaphore
- binary semaphore
 - mutex lock
 - only one process can get the resource
 - when resource = critical region, the atomcity can be guaranteed
 - initialize with token = 1

Example - summing

mutex

1



temp

???

sum

0

temp

???

```
void sum_odd(){  
    int i = 1, temp;  
    → for (;i < 100; i += 2){  
        P(mutex);  
        temp = sum;  
        temp += i;  
        sum = temp;  
        V(mutex);  
    }  
}
```

```
void sum_even(){  
    int i = 2, temp;  
    for (;i < 100; i += 2){  
        P(mutex);  
        temp = sum;  
        temp += i;  
        sum = temp;  
        V(mutex);  
    }  
}
```

Example - summing

mutex	0
-------	---

temp	???
------	-----

sum	0
-----	---

temp	???
------	-----

```
void sum_odd(){  
    int i = 1, temp;  
    for (; i < 100; i += 1){  
        P(mutex);  
        temp = sum;  
        temp += i;  
        sum = temp;  
        V(mutex);  
    }  
}
```

get a
token



```
void sum_even(){  
    int i = 2, temp;  
    for (; i < 100; i += 2){  
        P(mutex);  
        temp = sum;  
        temp += i;  
        sum = temp;  
        V(mutex);  
    }  
}
```

Example - summing



mutex	0
-------	---

temp	0
------	---

sum	0
-----	---

temp	???
------	-----

```
void sum_odd(){  
    int i = 1, temp;  
    for (; i < 100; i += 2){  
        P(mutex);  
        temp = sum;  
        temp += i;  
        sum = temp;  
        V(mutex);  
    }  
}
```



```
void sum_even(){  
    int i = 2, temp;  
    for (; i < 100; i += 2){  
        P(mutex);  
        temp = sum;  
        temp += i;  
        sum = temp;  
        V(mutex);  
    }  
}
```

Example - summing

mutex	0
-------	---


temp	1
------	---

sum	0
-----	---

temp	???
------	-----

```
void sum_odd(){  
    int i = 1, temp;  
    for (; i < 100; i += 2){  
        P(mutex);  
        temp = sum;  
        temp += i;  
        sum = temp;  
        V(mutex);  
    }  
}
```

→



```
void sum_even(){  
    int i = 2, temp;  
    for (; i < 100; i += 2){  
        P(mutex);  
        temp = sum;  
        temp += i;  
        sum = temp;  
        V(mutex);  
    }  
}
```

Example - summing

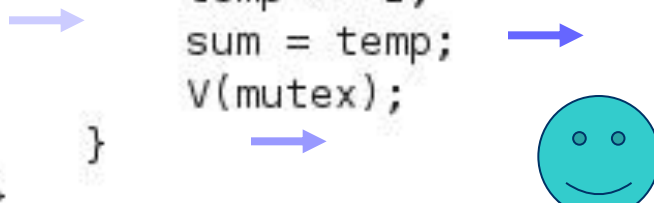
mutex	0
-------	---

temp	1
------	---

sum	0
-----	---

temp	???
------	-----

```
void sum_odd(){  
    int i = 1, temp;  
    for (; i < 100; i += 2){  
        P(mutex);  
        temp = sum;  
        temp += i;  
        sum = temp;  
        V(mutex);  
    }  
}
```



```
void sum_even(){  
    int i = 2, temp;  
    for (; i < 100; i += 2){  
        P(mutex);  
        temp = sum;  
        temp += i;  
        sum = temp;  
        V(mutex);  
    }  
}
```


Example - summing



mutex	0
-------	---

temp	1
------	---

sum	0
-----	---

temp	???
------	-----

```
void sum_odd(){  
    int i = 1, temp;  
    for (;i < 100; i += 2){  
        P(mutex);  
        temp = sum;  
        temp += i;  
        sum = temp;  
        V(mutex);  
    }  
}
```



```
void sum_even(){  
    int i = 2, temp;  
    → for (;i < 100; i += 2){  
        P(mutex);  
        temp = sum;  
        temp += i;  
        sum = temp;  
        V(mutex);  
    }  
}
```

Example - summing

mutex	0
-------	---

temp	1
------	---

sum	0
-----	---

temp	???
------	-----

```
void sum_odd(){  
    int i = 1, temp;  
    for (;i < 100; i += 2){  
        P(mutex);  
        temp = sum;  
        temp += i;  
        sum = temp;  
        V(mutex);  
    }  
}
```



no token is
available

```
void sum_even(){  
    int i = 2, temp;  
    for (;i < 100; i += 2){  
        P(mutex);  
        temp = sum;  
        temp += i;  
        sum = temp;  
        V(mutex);  
    }  
}
```

Example - summing

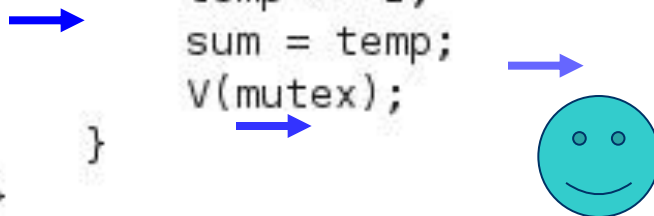
mutex	0
-------	---

temp	1
------	---

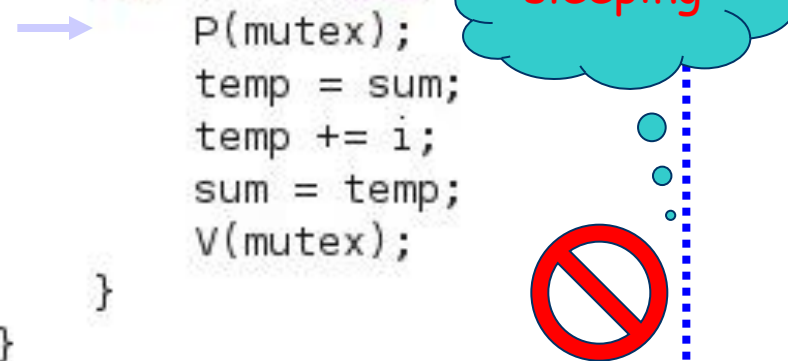
sum	0
-----	---

temp	???
------	-----

```
void sum_odd(){  
    int i = 1, temp;  
    for (; i < 100; i += 2){  
        P(mutex);  
        temp = sum;  
        temp += i;  
        sum = temp;  
        V(mutex);  
    }  
}
```



```
void sum_even(){  
    int i = 2, temp;  
    for (; i < 100; i += 2){  
        P(mutex);  
        temp = sum;  
        temp += i;  
        sum = temp;  
        V(mutex);  
    }  
}
```



Example - summing



mutex	0
-------	---

temp	1
------	---



sum	1
-----	---

temp	???
------	-----

```
void sum_odd(){  
    int i = 1, temp;  
    for (; i < 100; i += 2){  
        P(mutex);  
        temp = sum;  
        temp += i;  
        sum = temp;  
        V(mutex);  
    }  
}
```



```
void sum_even(){  
    int i = 2, temp;  
    for (; i < 100; i += 2){  
        P(mutex);  
        temp = sum;  
        temp += i;  
        sum = temp;  
        V(mutex);  
    }  
}
```



Example - summing

mutex	0
-------	---

temp	1
------	---

sum	1
-----	---

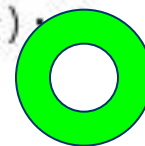
temp	???
------	-----

```
void sum_odd(){  
    int i = 1, temp;  
    → for (;i < 100; i += 2){  
        P(mutex);  
        temp = sum;  
        temp += i;  
        sum = temp;  
        V(mutex);  
    }  
}
```

release
the token

```
void sum_even(){  
    int i = 2, temp;  
    → for (;i < 100; i += 2){  
        P(mutex);  
        temp = sum;  
        temp += i;  
        sum = temp;  
        V(mutex);  
    }  
}
```

woken up



Example - summing

mutex	0
-------	---


temp	1
------	---

sum	1
-----	---

temp	???
------	-----

```
void sum_odd(){  
    int i = 1, temp;  
    → for (;i < 100; i += 2){  
        P(mutex);  
        temp = sum;  
        temp += i;  
        sum = temp;  
        V(mutex);  
    }  
}
```

```
void sum_even(){  
    int i = 2, temp;  
    → for (;i < 100; i += 2){  
        P(mutex);  
        temp = sum;  
        temp += i;  
        sum = temp;  
        V(mutex);  
    }  
}
```



Example - summing

mutex	0
-------	---

temp	1
------	---

sum	1
-----	---

temp	???
------	-----

```
void sum_odd(){  
    int i = 1, temp;  
    for (; i < 100; i += 2){  
        P(mutex);  
        temp = sum;  
        temp += i;  
        sum = temp;  
        V(mutex);  
    }  
}
```

```
void sum_even(){  
    int i = 2, temp;  
    for (; i < 100; i += 2){  
        P(mutex);  
        temp = sum;  
        temp += i;  
        sum = temp;  
        V(mutex);  
    }  
}
```



Example - summing

mutex	0
-------	---


temp	1
------	---

sum	1
-----	---

temp	???
------	-----

```
void sum_odd(){  
    int i = 1, temp;  
    → for (;i < 100; i += 2){  
        P(mutex);  
        temp = sum;  
        temp += i;  
        sum = temp;  
        V(mutex);  
    }  
}
```

```
void sum_even(){  
    int i = 2, temp;  
    → for (;i < 100; i += 2){  
        P(mutex);  
        temp = sum;  
        temp += i;  
        sum = temp;  
        V(mutex);  
    }  
}
```



Example - summing

mutex	0
-------	---


temp	1
------	---

sum	1
-----	---

temp	1
------	---

```
void sum_odd(){  
    int i = 1, temp;  
    → for (;i < 100; i += 2){  
        P(mutex);  
        temp = sum;  
        temp += i;  
        sum = temp;  
        V(mutex);  
    }  
}
```

```
void sum_even(){  
    int i = 2, temp;  
    for (;i < 100; i += 2){  
        P(mutex);  
        temp = sum;  
        temp += i;  
        sum = temp;  
        V(mutex);  
    }  
}
```



Example - summing

mutex	0
-------	---


temp	1
------	---

sum	1
-----	---

temp	3
------	---

```
void sum_odd(){  
    int i = 1, temp;  
    → for (;i < 100; i += 2){  
        P(mutex);  
        temp = sum;  
        temp += i;  
        sum = temp;  
        V(mutex);  
    }  
}
```

```
void sum_even(){  
    int i = 2, temp;  
    for (;i < 100; i += 2){  
        P(mutex);  
        temp = sum;  
        temp += i;  
        sum = temp;  
        V(mutex);  
    }  
}
```



Example - summing

mutex	0
-------	---

temp	1
------	---

sum	3
-----	---

temp	3
------	---



```
void sum_odd(){  
    int i = 1, temp;  
    → for (;i < 100; i += 2){  
        P(mutex);  
        temp = sum;  
        temp += i;  
        sum = temp;  
        V(mutex);  
    }  
}
```

```
void sum_even(){  
    int i = 2, temp;  
    for (;i < 100; i += 2){  
        P(mutex);  
        temp = sum;  
        temp += i;  
        sum = temp;  
        V(mutex);  
    }  
}
```





Implementation

- data structure
 - token
 - queue
- operation
 - P() operation
 - if blocked, put the current process in the queue
 - V() operation
 - if some processes are blocking, wake up one of them
- It is easy to implement semaphore with sleep() and wakeup().



Recursion?

- Different processes may access the same token by P-V operations concurrently.
 - P-V operations are also critical regions!
 - Atomicity & consistency must be guaranteed.
- How to solve this "recursion"?
 - locking!



Summary about semaphore

- easy to understand
 - get: if no resource, then sleep
 - release: wake up one waiting process
- easy to use
 - $P()$, $V()$
- easy to implement
 - counter + thread management
- but the correctness of semaphore depends on the correctness of locking



-
- Locking



Disable interrupt

- the easiest way to implement locking
- goal: disable context switch
 - context switch = interrupt driven stack switch
 - no context switch = no concurrent access
 - guarantee the atomicity of the execution in critical regions



Shortcoming

- not suitable for long critical regions
 - reduce concurrency of the whole system
 - interrupts can not obtain response in time
- malicious user programs
 - do not enable interrupt any longer
- not work for SMP
 - every CPU has its own interrupt controller
 - local APIC
 - other CPUs can still enter the same region
 - atomicity & consistency are no guaranteed



Conclusions

- Disabling interrupt works well when
 - protecting short critical regions
 - in kernel space
 - in a uniprocessor system
- Any problem?

Trap 1 - nested locking

```
void critical_region1() {  
    lock();  
    // ...  
    V(sem);  
    // ...  
    unlock();  
}
```


```
void V(Sem *s) {  
    lock();  
    // ...  
    unlock();  
}
```

not safe
any longer



Trap 2 - sleep during locking

```
void P() {  
    lock();  
    // ...  
    if(counter == 0)  
        sleep();  
    // ...  
    unlock();  
}
```



the consistency of locking
should not be violated by
other processes

Trap 3 - use locking in interrupts

```
void do_timer() {  
    // ...  
    critical_region1();  
    // ...  
}
```

```
void critical_region1() {  
    lock();  
    // ...  
    V(sem);  
    // ...  
    unlock();  
}
```

← should
not sti()



Why it is hard to implement locking?

- We can not predict where we will use locking in the future.
 - make the implementation as flexible as possible
 - so does other designs
- Refactoring your code is another “solution”.
 - but the cost is much higher
 - e.g. extra attentions should be payed, instead of using nested locking directly



How to verify your solution?

- try to find a common rule for locking
 - this may become a unified solution
- try to find consistency conditions
 - make them the conditions of assertions
- no assertion fail = no bad things happen
 - only bad things which may be captured by assertions you have written



How to verify your solution?

- try best to expose inconsistent states
 - use a higher timer interrupt frequency
 - more interrupts = more probability to trigger interrupt-related errors
- run test cases for a long time
 - one night is enough
 - How to detect mysterious reboot when you are not in front of your computer?



Have fun!

- Implementing locking is more “interesting” than just listening.
- This is also a training for software testing.
 - find assertion conditions
- Locking in SMP is more tricky.
 - We may mention this topic next week.