

Message passing

- Message passing
- SMP/multi-core system
- Message passing in SMP

- 
-
- Message passing

Message passing

- in the real world





Abstraction

- message
- sender
- receiver
- send
- receive



Message

- a structure
 - sender
 - receiver
 - type
 - content
- How to interpret the type and content is up to the sender and receiver.



Sender & Receiver

- the entities who manipulate the message
 - in Nanos, they are all threads/processes
- sender:
 - fill the message
 - send it
- receiver:
 - receive the message
 - do something according to the content of the message



Send & Receive

- `send(dest, m):`
 - deliver the message `m` to `dest`
 - every process has a "mail box"
 - put the message into receiver's "mail box"
 - always succeed, return immediately
- `receive(src, m):`
 - pick a message from `src` from the "mail box"
 - `src` can be anyone
 - if no message wanted is available, wait until one comes



The power

- communication
- notification
- synchronization
- multi-value returning
- decoupling

Communication

- send some data to other
 - the natural function

0:13:36
明天 会不会小测？
【提示：此用户正在使用Q+ Web：<http://web.qq.com/>】

0:13:49
小测是什么？

0:14:01
小册？！

0:14:22
测你妹哦

0:14:41
侧你妹啊

0:14:47
到底什么情况？

0:15:06
这种行为！

0:15:15
令人发指！！！！

0:15:29
意思说明天要早起？

0:15:52
到底要不要测？



Notification

- send a message for occurrence of some events
 - this is a special case of communication
 - usually to tell the receiver to handle the events

20:50:42

转自赵老师：

陶主任关于“导师指导的学分”的答复：导师指导的学分，可以在4年级两个学期内给出，但是需要导师明确研究内容，需要提交研究成果，成果不一定是发表论文，可以是软件系统、研究报告，但必须要有成果，有导师的签字认可。学分给多少，由导师根据学生投入的时间和取得的成果建议，由教学副主任签字认可。这个学分可以不要，不是必修。另外，毕业设计的成果和这个学分的成果不得雷同，不能用同一件事情拿两份学分。

大家可以现在开始和自己的导师联系，后面慢慢做项目或者其他事情

Synchronization





Synchronization (cont.)

- implement the relationship of “happen after”
 - if A should happen after B, then block the process until B finishes
 - receive() can exactly do the same thing!
- running

```
void runner(int id) {  
    relax_himself();  
    receive(referee, "ready!");  
    ready();  
    receive(referee, "go!");  
    run();  
}
```



Communication & Synchronization

- communication + synchronization = network
- client
 - send requests (communication)
 - wait for reply (synchronization)
- server
 - wait for request (synchronization)
 - process request
 - send reply (communication)



Multi-value returning

- When replying, a message can carry more than one return value.
 - more flexible than the function calls



Decoupling

- Processes communicate only by message passing.
- Other information is private.
 - invisible for other processes
 - fewer critical regions!
- This is exactly the core idea of micro-kernel.



Implementation

- Every “mail box” is a critical area!
 - Different processes may send messages to a certain process simultaneously.
- Receiver cannot receive message when the “mail box” is empty!
- a problem with mutex and synchronization



Implementation (cont.)

- In Nanos, send() is asynchronous.
 - send() always succeeds without being blocked
 - because we will use send in the interrupt handlers
 - What does this mean?
- low-level API available
 - semaphore
 - locking
- How to implement message passing with them?

- 
-
- SMP/multi-core system



Recap – critical region

- In a uni-processor system, only process-level critical region exists.
 - physical limit
 - only one execution flow can run at any point of time
- the root cause:
 - interrupts
- an effective solution:
 - disable interrupt



More data races

- process → execution flow
 - process1 v.s. process2
 - process v.s. interrupt
 - process v.s. signal
 - interrupt1 v.s. interrupt2
 - interrupt1 v.s. itself
 - re-enterable & unre-enterable
- more tricky in SMP
 - "real" simultaneousness

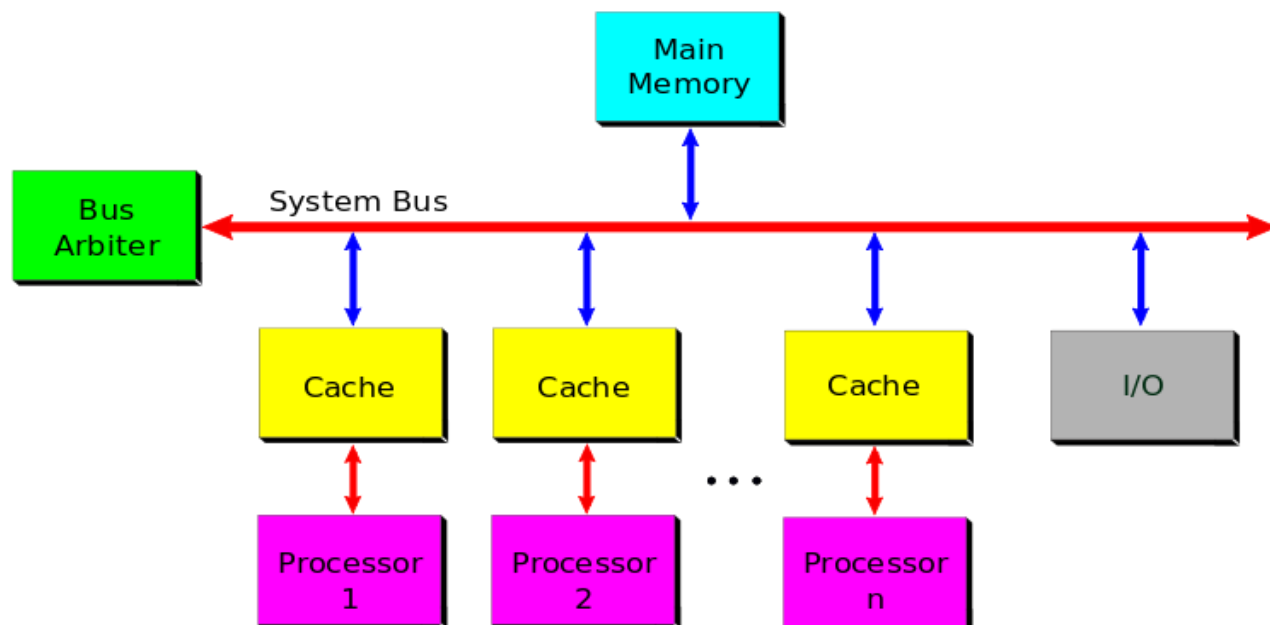


But in SMP...

- Everything becomes tricky.
- the root cause
 - parallelism
 - not only concurrency
- But first you should know how SMP works.

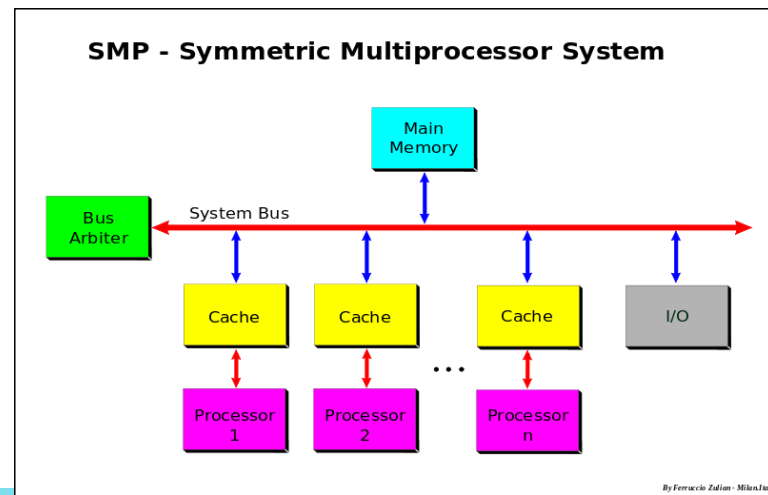
SMP

SMP - Symmetric Multiprocessor System



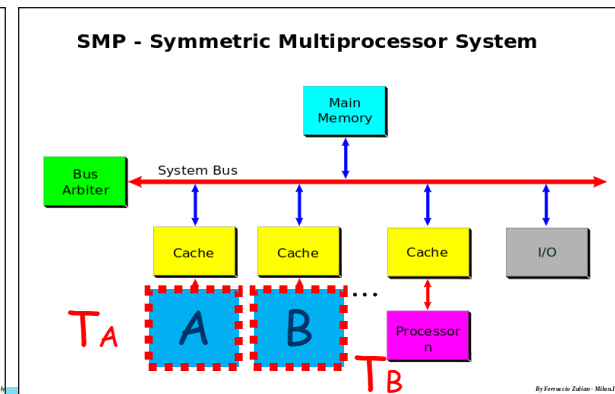
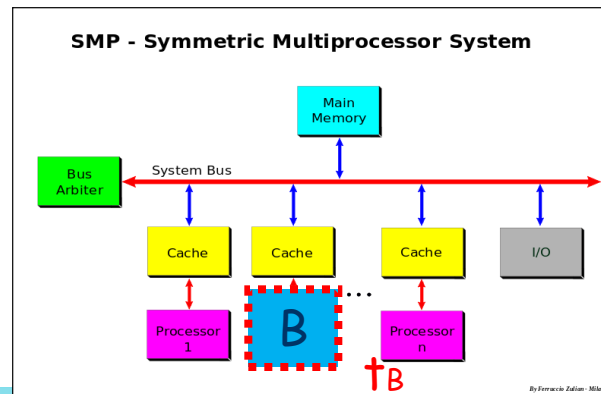
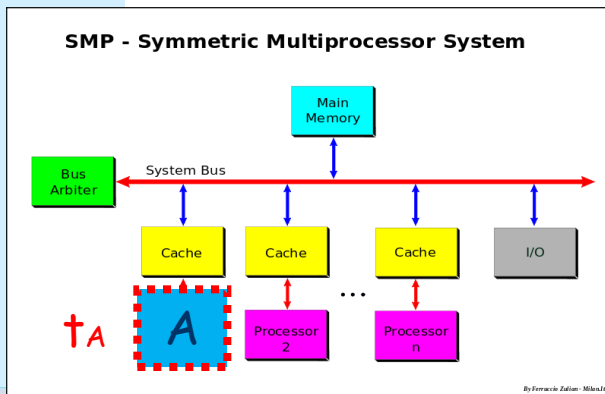
SMP

- Processors connect with each other via system bus.
- Each processor has its own L1 cache.
- Each processor can perform computation individually.



How much can we benefit from SMP?

- For two task A & B without data dependency
 - in theory: $t_A = T_A, t_B = T_B$
 - in fact: $t_A < T_A, t_B < T_B$
 - or even $t_A + t_B < \min(T_A, T_B)$





Amdahl's Law

http://en.wikipedia.org/wiki/Amdahl's_law


- For one task wants to benefit from SMP by multi-threading:
 - B = the fraction of the task can only be serialized
 - $T(n)$ = the time to finish the task with n threads

$$T(n) = T(1) \left(B + \frac{1}{n} (1 - B) \right)$$

- example: $T(1) = 20s$, $B = 5\%$
 - then $T(n) > T(\text{infinity}) = 1s \rightarrow 20x$ faster
- $T(n)$ is the lower bound in theory
 - even worse in practise

Almost NO linear benefit

http://en.wikipedia.org/wiki/Amdahl's_law





In theory

- conclusion from Amdahl's Law
 - How much a task can be parallelized depends the task itself.
- parallelized without any serialization
 - vector addition
- parallelized with certain serialization
 - vector summation
- hard to parallelize at all
 - Hanoi tower



In practise

- What makes us far from the lower bound?
 - data race
 - data dependency
 - the cost from multi-threading
 - the cost from SMP/multi-core system
- even sometime perform worse than the serialization version on uni-core system
- better parallel algorithm
 - how to compute π ?



1. Cost from multi-threading

- Thread creation/destruction has its own cost.

- more than you do in OS lab

```
#define N 10000000  
#define NR_THREAD N  
int a[N], b[N], c[N];  
void add(int tid) { c[tid] = a[tid] + b[tid]; }
```

- Amdahl's Law + cost from multi-threading
 - small scale parallelization (2-8 threads)



2. Data race

- This is exactly the “critical region” problem.
- goal: provide exclusive accessing
- Disabling interrupt does not work in SMP/multi-core system.
- solution: atomic instructions
- atomic inst. → locking → high-level mutex



Atomic instructions

- provide exclusive accessing at ISA level
- RMW (read-modify-write)
 - Test&Set
 - Compare&Swap
 - load-link / store-condition
- take producer-consumer problem as an example

Test&Set

Computer Architecture, Princeton University

- blocking mutex

```
P:    Test&Set (mutex), Rtemp
      if (Rtemp != 0) goto P
      Load Rhead, (head)
spin: Load Rtail, (tail)
      if Rhead == Rtail goto spin
      Load R, (Rhead)
      Rhead = Rhead + 1
      Store Rhead, (head)
V:    Store 0, (mutex)
      process(R)
```

```
Test&Set(addr, reg):
    R[reg] = M[addr];
    if (R[reg] == 0)
        M[addr] = 1;
```

*Critical
Section*

Compare&Swap

Computer Architecture, Princeton University

- non-blocking mutex
- “ABA” problem

```
try: Load Rhead, (head)
spin: Load Rtail, (tail)
      if Rhead == Rtail goto spin
      Load R, (Rhead)
      Rnewhead = Rhead + 1
      Compare&Swap(head), Rhead, Rnewhead
      if (status == fail) goto try
      process(R)
```

```
Compare&Swap (addr, rt, rs):
    if (R[rt] == M[addr])
        M[addr] = R[rs];
        R[rs] = R[rt];
        status = success;
    else
        status = fail;
```




Implement atomic instructions

- lock the system bus
 - when the system bus is locked, no other memory accessing is allowed
- This may slow down the performance in SMP/multi-core system.
- Any better idea?

Load-link / Store-condition

Computer Architecture, Princeton University

try:
spin:

```
Load-link Rhead, (head)
Load Rtail, (tail)
if Rhead == Rtail goto spin
Load R, (Rhead)
Rhead = Rhead + 1
Store-conditional Rhead, (head)
if (status == fail) goto try
process(R)
```

Store-condition(addr, reg):

if(<flag, m> == <1, addr>)

cancel other processors' reservation on addr;

M[addr] = R[reg];

status = success;

else

status = fail;

Load-link(reg, addr):

<flag, m> = <1, addr>;

R[reg] = M[addr];




Load-link / Store-condition

- non-blocking mutex
- avoid "ABA" problem
- modification can be arbitrary
 - however, more complex modification -> higher probability to fail for store-condition
- see "MIPS32 instruction set" for more details


3. Data dependency

- Thread A needs the result from thread B to perform computation.
- Synchronization is needed.



Synchronization

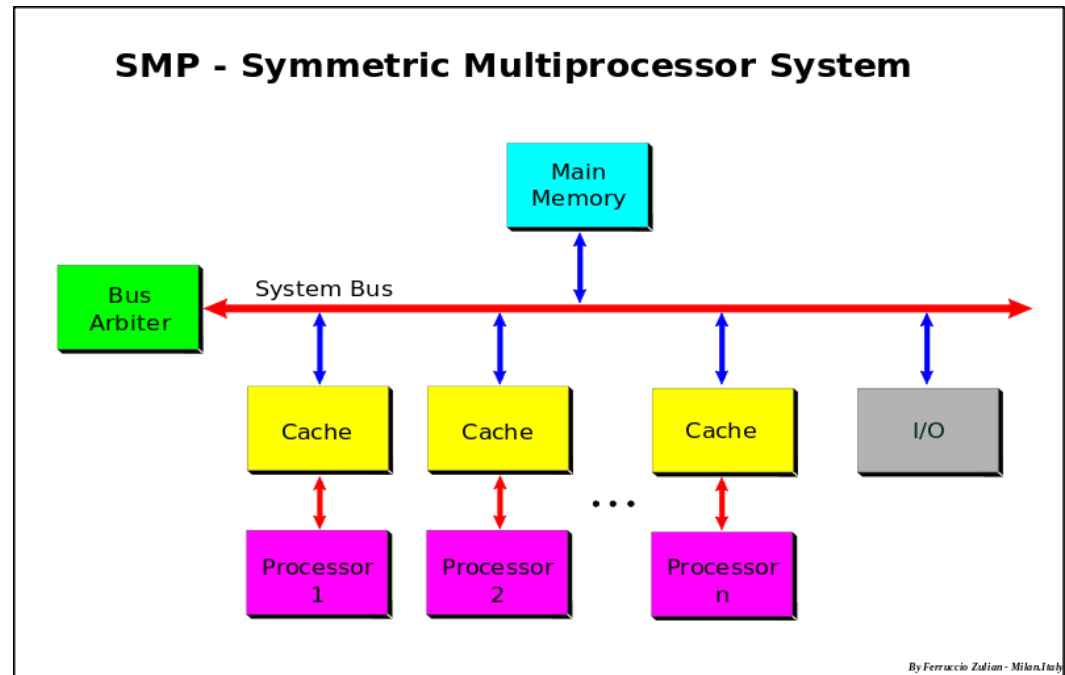
- wait until other threads are ready



- over-synchronization

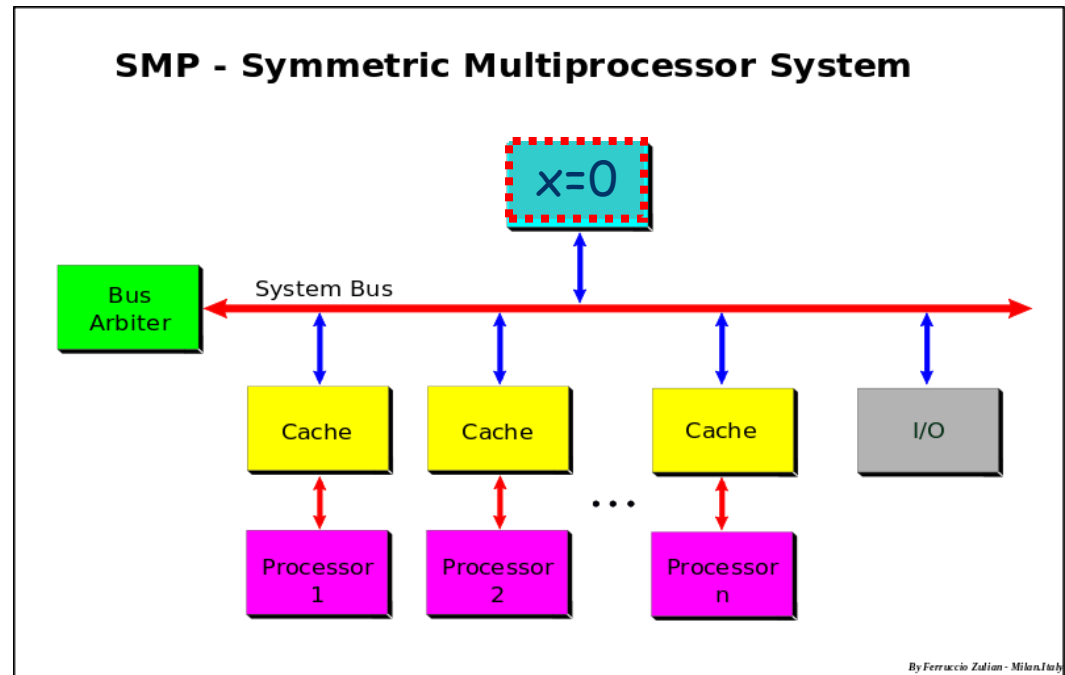
4. Cost from SMP/multi-core system

- Any problem?



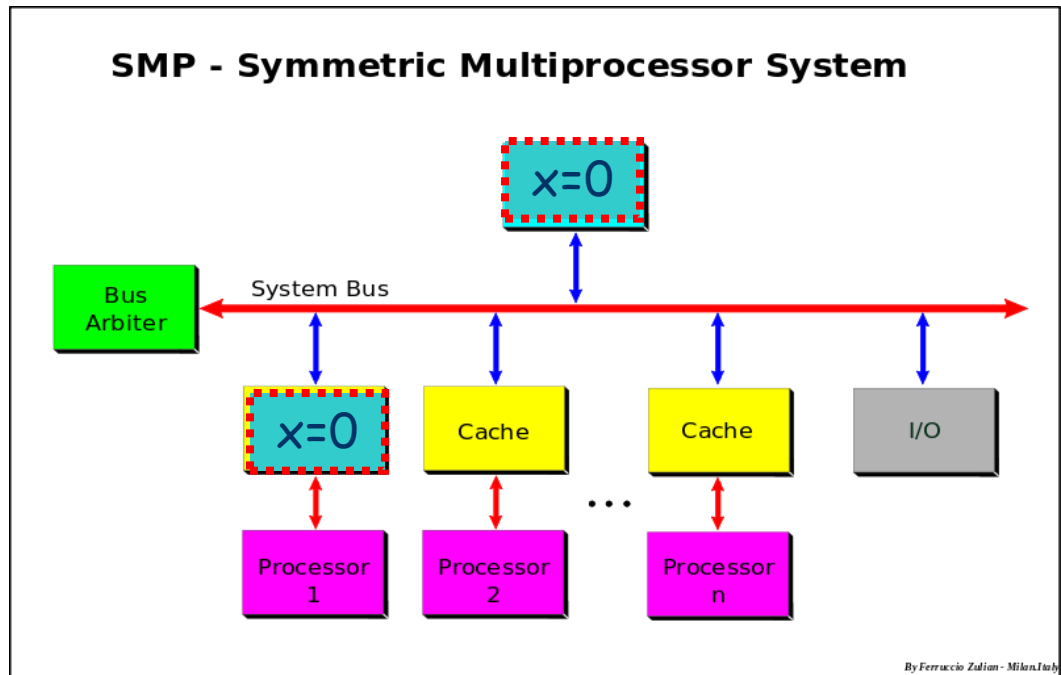
4. Cost from SMP/multi-core system

- Any problem?
 - initially $x = 0$ in memory



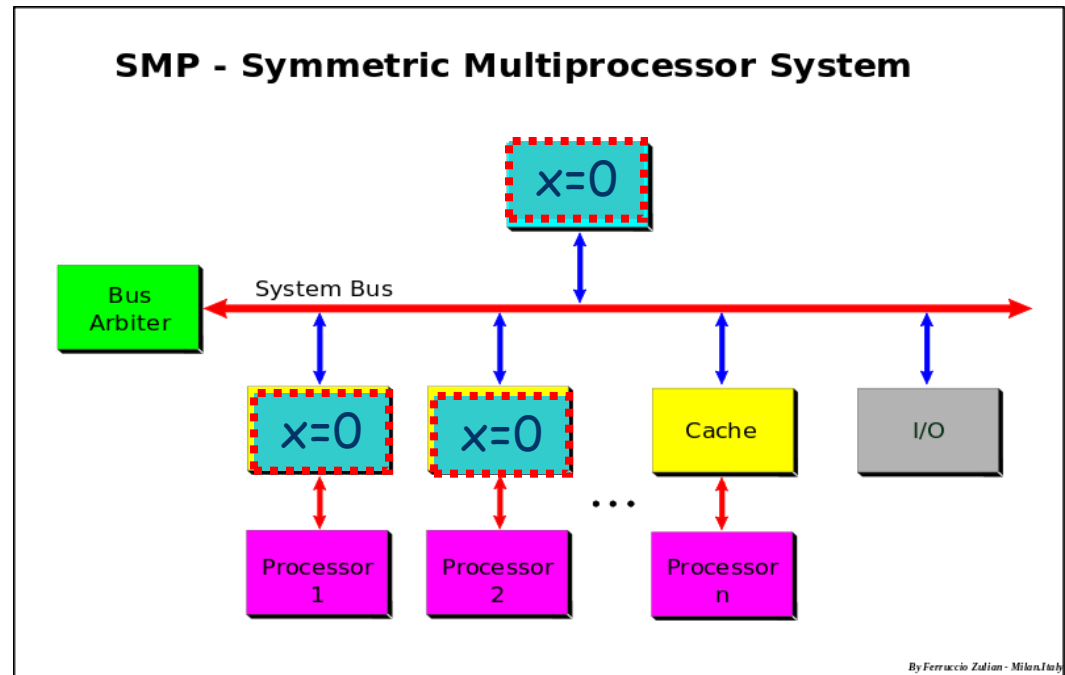
4. Cost from SMP/multi-core system

- Any problem?
 - initially $x = 0$ in memory
 - P1: read x



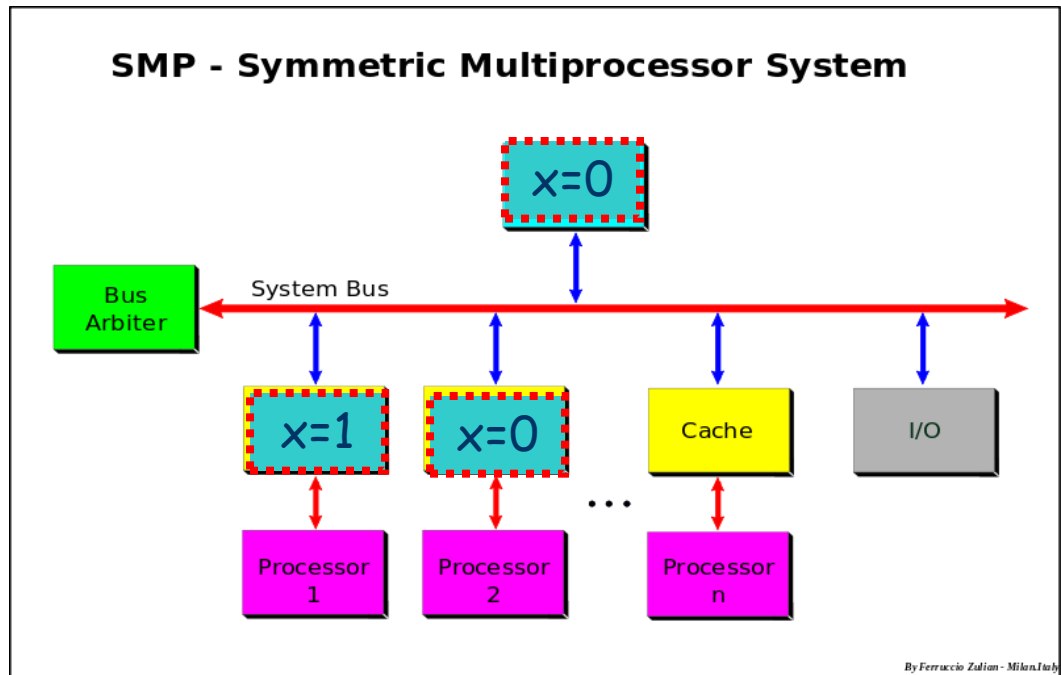
4. Cost from SMP/multi-core system

- Any problem?
 - initially $x = 0$ in memory
 - P1: read x
 - P2: read x



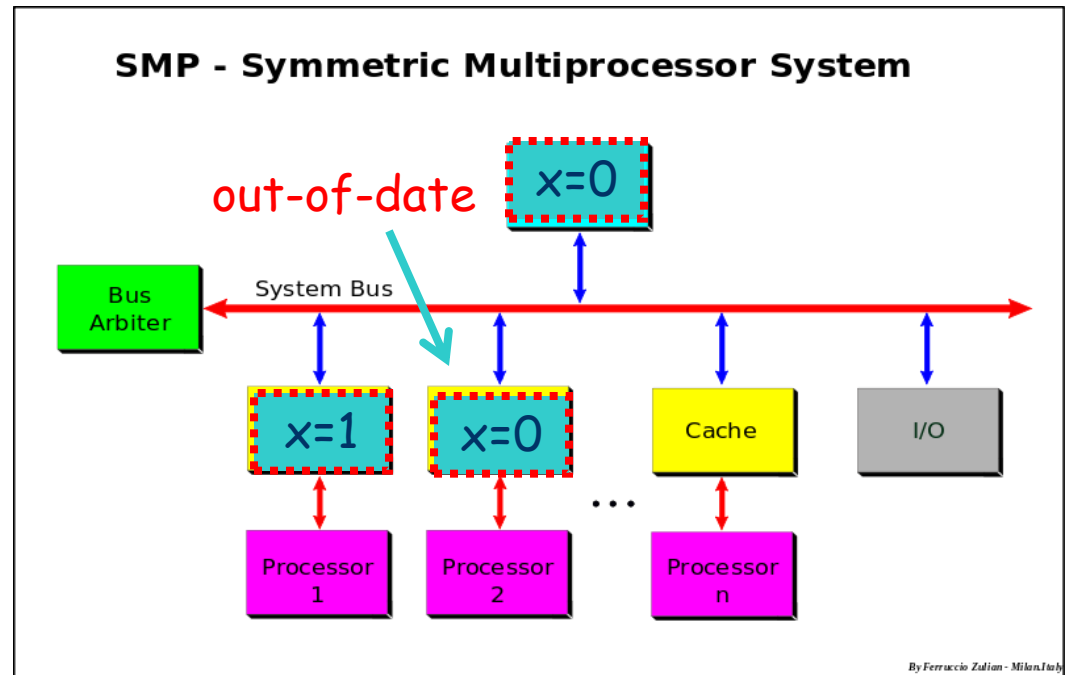
4. Cost from SMP/multi-core system

- Any problem?
 - initially $x = 0$ in memory
 - P1: read x
 - P2: read x
 - P1: write $x = 1$



4. Cost from SMP/multi-core system

- Any problem?
 - initially $x = 0$ in memory
 - P1: read x
 - P2: read x
 - P1: write $x = 1$
 - P2: read x



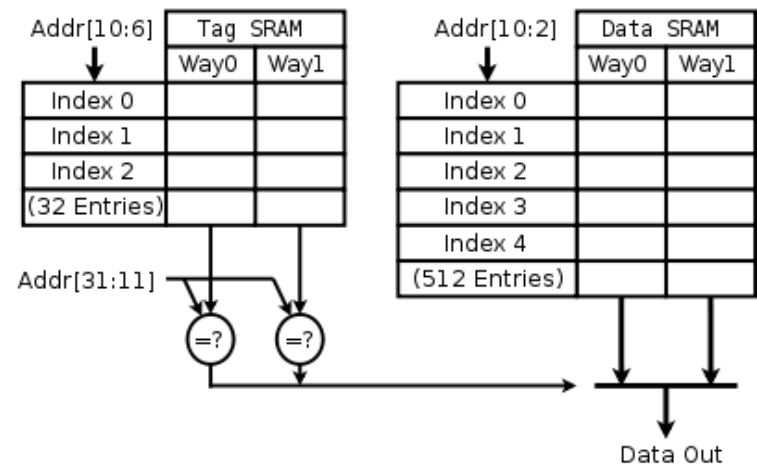


Cache coherence

- Cache of one processor is invisible for other processors.
 - data may be out-of-date
 - where the valid locates is not known
- Mechanism is needed for different caches to communicate.
 - at least notify others about updating
 - snooping (broadcast)

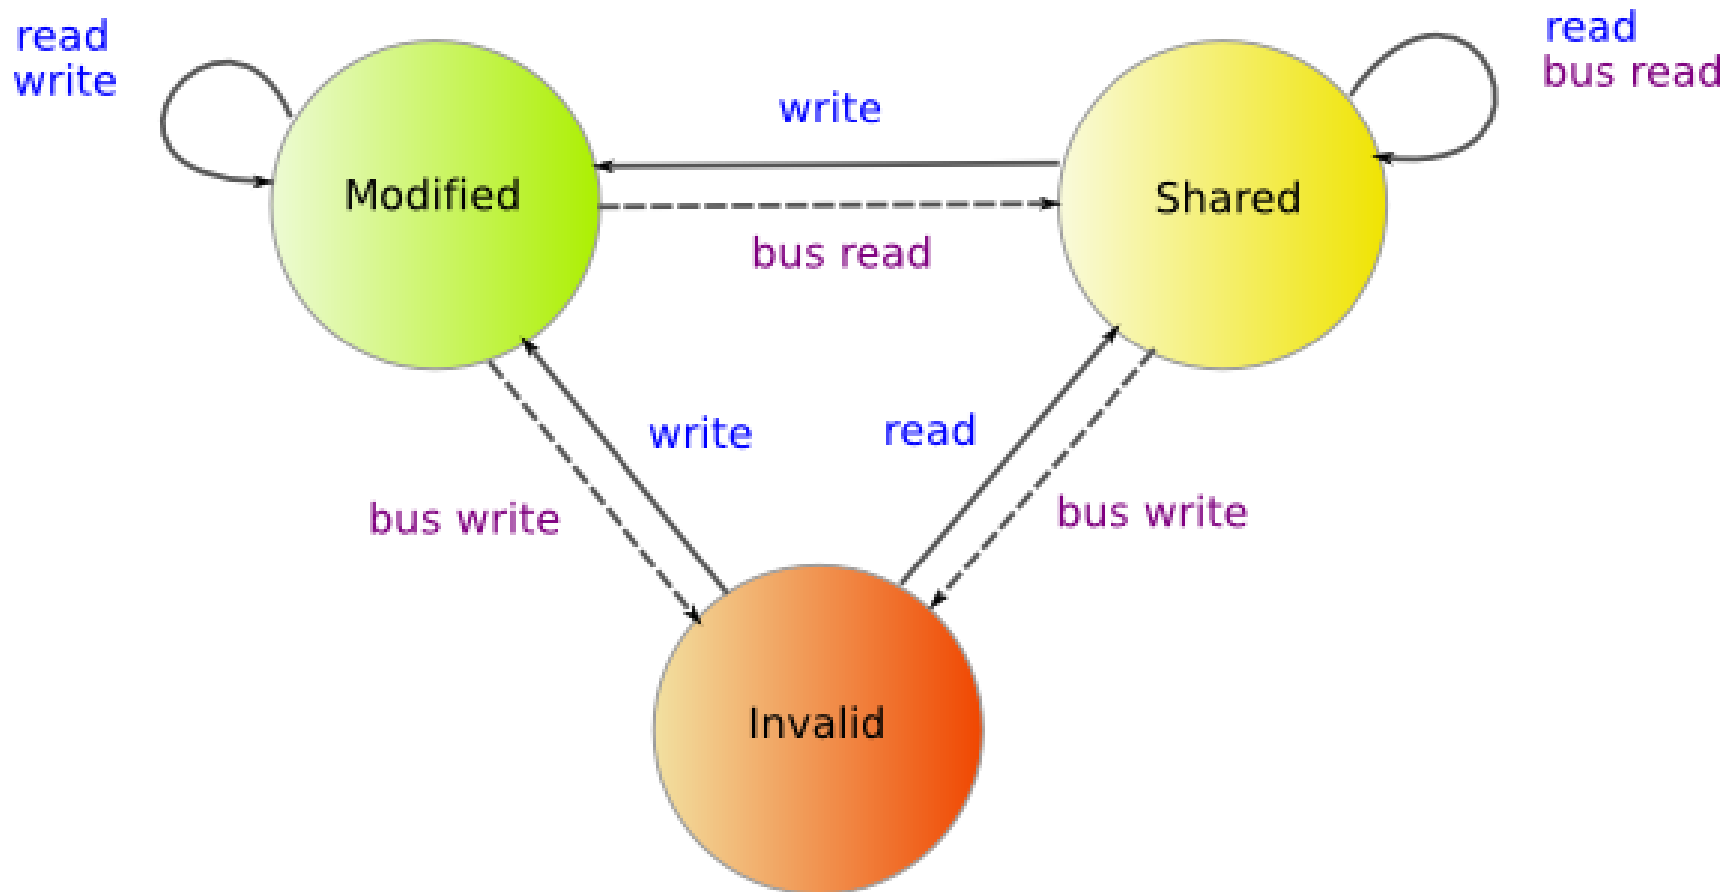
Cache coherence protocol - MSI

- For each cache line, maintain its state bits
 - M(modified): the data is dirty, needs writing back in the future
 - S(shared): the data is clean, and is shared by other caches
 - I(invalid): the data is invalid



4KB, 2-way set-associative 64B line cache read path

Cache coherence protocol - MSI





Performance

- To guarantee cache coherence, threads in SMP/multi-core system may perform badly.
 - Share data misses often dominate cache behavior even though they may only be 10% to 40% of the data accesses.
- Shared memory has little locality.
 - Why?



Memory ordering

- let $x = y = 0$ initially

Processor 1	Processor 2
<code>movl \$1, (x)</code>	<code>movl \$1, (y)</code>
<code>movl (y), %eax</code>	<code>movl (x), %ebx</code>

- What will the result be?
- Is "`%eax = %ebx = 0`" possible?




Memory ordering

Intel® 64 and IA-32 Architectures Software Developer's Manual

- Reads are not reordered with other reads.
- Writes are not reordered with older reads.
- Writes to memory are not reordered with other writes, with the following exceptions:
 - writes executed with the CLFLUSH instruction;
 - streaming stores (writes) executed with the non-temporal move instructions (MOVNTI, MOVNTQ, MOVNTDQ, MOVNTPS, and MOVNTPD); and
 - string operations (see Section 8.2.4.1).
- Reads may be reordered with older writes to different locations but not with older writes to the same location.
- Reads or writes cannot be reordered with I/O instructions, locked instructions, or serializing instructions.
- Reads cannot pass earlier LFENCE and MFENCE instructions.
- Writes cannot pass earlier LFENCE, SFENCE, and MFENCE instructions.
- LFENCE instructions cannot pass earlier reads.
- SFENCE instructions cannot pass earlier writes.
- MFENCE instructions cannot pass earlier reads or writes.

Memory ordering

Processor 1	Processor 2
movl \$1, (x)	movl \$1, (y)
movl (y), %eax	movl (x), %ebx



Processor 1	Processor 2
movl \$1, (x)	movl (x), %ebx
movl (y), %eax	movl \$1, (y)

- How to fix?



Memory ordering

- memory barrier (memory fence)
 - force serialization before and after the barrier instructions
 - lfence, sfence, mfence
- Before any memory operation starts after the barrier instruction, all memory operations before it must finish.
- cost
 - more instruction dependency, less performance



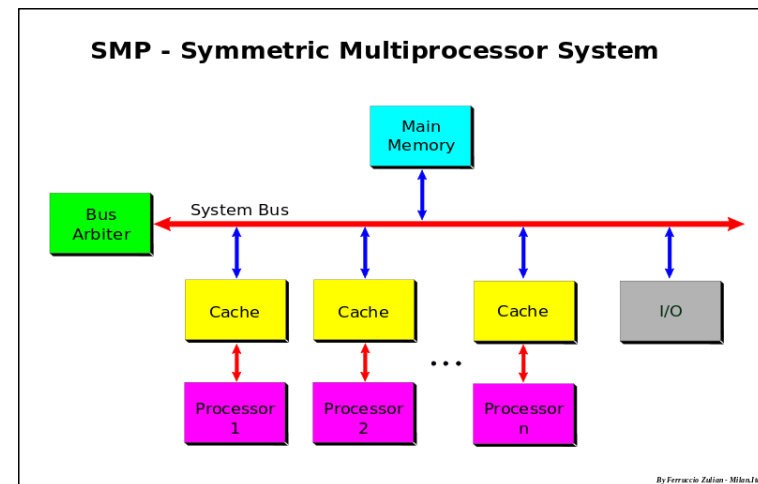
Summary

- SMP/multi-core system sounds attractive.
- But in practise, some performance is lost to guarantee the correctness.
 - atomic instructions, cache coherence, memory barrier...
 - more than you expect!
- In computer science, trade-off is everywhere.
- How we can benefit more from SMP/multi-core system is still in research.

- 
-
- Message passing in SMP

Shared memory

- This is the architecture we use in PC.
 - processors connected by a bus
 - a global memory shared by all processors
 - also connected on a bus
- All communications must go through the bus.
 - more processors → bottleneck






Architecture based on message passing

- Processors connected to switches.
- To communicate, send messages to switches.
- Switches perform forwarding & routing.
- Each processor has its local memory.
 - usually no global memory
 - can use as unified/separate address space(s)
- Yes! This is a network!
 - Interconnection Network
 - NoC (Network on Chip)

Bus IN






Network on Chip

- Now the topic turns to network...
 - How to increase
 - throughput, reliability?
 - How to decrease
 - latency, bottleneck, congestion, design cost, energy cost?
- Trade-off is everywhere.

Topology

[Lecture notes based on slides created by Mary Jane Irwin, and Patterson & Hennessy "Computer Organization and Design"]

Bus IN




- ❑ N processors, 1 switch (●), 1 link (the bus)
- ❑ Only 1 simultaneous transfer at a time
 - NB = link (bus) bandwidth * 1
 - BB = link (bus) bandwidth * 1

Topology

[Lecture notes based on slides created by Mary Jane Irwin, and Patterson & Hennessy "Computer Organization and Design"]

Ring IN




- ❑ N processors, N switches, 2 links/switch, N links
- ❑ N simultaneous transfers
 - $NB = \text{link bandwidth} * N$
 - $BB = \text{link bandwidth} * 2$
- ❑ If a link is as fast as a bus, the ring is only twice as fast as a bus in the worst case, but is N times faster in the best case

Topology

[Lecture notes based on slides created by Mary Jane Irwin, and Patterson & Hennessy "Computer Organization and Design"]

Fully Connected IN




- ❑ N processors, N switches, N-1 links/switch, $(N*(N-1))/2$ links
- ❑ N simultaneous transfers
 - $NB = \text{link bandwidth} * (N*(N-1))/2$
 - $BB = \text{link bandwidth} * (N/2)^2$

Topology

[Lecture notes based on slides created by Mary Jane Irwin, and Patterson & Hennessy "Computer Organization and Design"]

Crossbar (Xbar) Connected IN




- ❑ N processors, N^2 switches (unidirectional), 2 links/switch, N^2 links
- ❑ N simultaneous transfers
 - $NB = \text{link bandwidth} * N$
 - $BB = \text{link bandwidth} * N/2$


Topology

[Lecture notes based on slides created by Mary Jane Irwin, and Patterson & Hennessy "Computer Organization and Design"]

Hypercube (Binary N-cube) Connected IN



2-cube




3-cube

- ❑ N processors, N switches, $\log N$ links/switch, $(N \log N)/2$ links
- ❑ N simultaneous transfers
 - $NB = \text{link bandwidth} * (N \log N)/2$
 - $BB = \text{link bandwidth} * N/2$

Topology

[Lecture notes based on slides created by Mary Jane Irwin, and Patterson & Hennessy "Computer Organization and Design"]

2D and 3D Mesh/Torus Connected IN




- ❑ N processors, N switches, 2, 3, 4 (2D torus) or 6 (3D torus) links/switch, $4N/2$ links or $6N/2$ links
- ❑ N simultaneous transfers
 - $NB = \text{link bandwidth} * 4N$ or $\text{link bandwidth} * 6N$
 - $BB = \text{link bandwidth} * 2 N^{1/2}$ or $\text{link bandwidth} * 2 N^{2/3}$

Topology

[Lecture notes based on slides created by Mary Jane Irwin, and Patterson & Hennessy "Computer Organization and Design"]

Fat Tree




- ❑ N processors, $\log(N-1) \cdot \log N$ switches, 2 up + 4 down = 6 links/switch, $N \cdot \log N$ links
- ❑ N simultaneous transfers
 - $NB = \text{link bandwidth} \cdot N \log N$
 - $BB = \text{link bandwidth} \cdot 4$

Topology

[Lecture notes based on slides created by Mary Jane Irwin, and Patterson & Hennessy "Computer Organization and Design"]

Fat Tree

- Trees are good structures. People in CS use them all the time. Suppose we wanted to make a tree network.



- Any time A wants to send to C, it ties up the upper links, so that B can't send to D.
 - The bisection bandwidth on a tree is horrible - 1 link, at all times
- The solution is to 'thicken' the upper links.
 - More links as the tree gets thicker increases the bisection
- Rather than design a bunch of N-port switches, use pairs

Topology

[Lecture notes based on slides created by Mary Jane Irwin, and Patterson & Hennessy "Computer Organization and Design"]

SGI NUMalink Fat Tree




Figure 2. 512-Processor Dual "Fat-Tree" Interconnect Topology

www.embedded-computing.com/articles/woodacre



Cache coherence

- For unified memory address space, we should still maintain cache coherence.
- Protocol based on snooping does not work.
 - Broadcast has a very high cost.
- directory-base protocols
 - each entry of the directory presents the state of a memory block
 - who has the copy, whether it is dirty...
 - ask the directory control for information



Shared memory v.s. message passing

- shared memory
 - memory is shared
 - communicate with load & store
 - has issue of critical region
 - synchronization with explicit instruction
- message passing
 - memory is private
 - communicate with send & receive
 - no issue of critical region
 - synchronization is implicit