

# Lab1 feedback

- Lab1 feedback
- Hints for implementing locking

- 
- 
- Lab1 feedback
    - 必答题

# Submission

- much better than Lab0



- however

- no project in stage 3
  - 121220046, 121220133
- non-standard character
  - 121220151, 121220158, 121220319, 121242031, 121250198
- 20% penalty



# Shell commands

1. **shell命令** 完成Lab1的内容之后，你整个工程中的.c, .h和.S文件总共有多少行代码？你是使用什么命令得到这个结果的？和Lab1的框架代码相比，你在Lab1中编写了多少行代码？（Hint：使用git checkout可以回到“过去”）你可以把这条命令写入Makefile中，随着实验进度的推进，你可以很方便地统计工程的代码行数，例如敲入make count就会自动运行统计代码行数的命令。再来个难一点的，除去空行之外，你整个工程中的.c, .h和.S文件总共有多少行代码？

```
find . -name '*.chS' | xargs wc
```

```
find . -name '*.chS' | xargs grep -v '^$' | wc
```



# Compiling & linking

2. **编译与链接** 你应该在框架代码中看到`include/x86/io.h`中看到一些由 `static inline` 开头定义的函数. 分别尝试去掉 `static`, `inline` 或者去掉两者, 然后进行编译, 你会看到发生错误. 请解释为什么会发生这些错误? 你有办法证明你的想法吗?

- without "static"
  - multiple definition
- without "inline"
  - unused-function
- without "static" and "inline"
  - multiple definition



# Why?

- What does "static" mean?

```
static int x;  
static void fun() {  
    static int y;  
    y ++;  
}
```

- Can you use "x" in another source file?
- What is "static" from the view of machine?

# static & symbol table

```
// main.c
static void fun() {
    static int yyyyy;
    yyyyy++;
}
```

```
// fun.c
void fun() {
}
```

35:	080483f0	0 FUNC	LOCAL	DEFAULT	14	frame_dummy
36:	080495e4	0 OBJECT	LOCAL	DEFAULT	19	__frame_dummy_init_array_
37:	00000000	0 FILE	LOCAL	DEFAULT	ABS	main.c
38:	0804841c	26 FUNC	LOCAL	DEFAULT	14	fun
39:	08049708	4 OBJECT	LOCAL	DEFAULT	26	yyyyy.1816
40:	00000000	0 FILE	LOCAL	DEFAULT	ABS	fun.c
41:	00000000	0 FILE	LOCAL	DEFAULT	ABS	crtstuff.c
42:	080485e0	0 OBJECT	LOCAL	DEFAULT	18	__FRAME_END__
43:	080495ec	0 OBJECT	LOCAL	DEFAULT	21	__JCR_END__
44:	080495e8	0 NOTYPE	LOCAL	DEFAULT	19	__init_array_end
45:	080495f0	0 OBJECT	LOCAL	DEFAULT	22	__DYNAMIC
46:	080495e4	0 NOTYPE	LOCAL	DEFAULT	19	__init_array_start
47:	080496e4	0 OBJECT	LOCAL	DEFAULT	24	__GLOBAL_OFFSET_TABLE__
48:	08048480	5 FUNC	GLOBAL	DEFAULT	14	__libc_csu_fini
49:	080484ea	0 FUNC	GLOBAL	HIDDEN	14	__i686.get_pc_thunk.bx
50:	00000000	0 NOTYPE	WEAK	DEFAULT	UND	__ITM_deregisterTMCloneTab
51:	080496fc	0 NOTYPE	WEAK	DEFAULT	25	data_start
52:	00000000	0 FUNC	GLOBAL	DEFAULT	UND	printf@@GLIBC_2.0
53:	08049704	0 NOTYPE	GLOBAL	DEFAULT	ABS	__edata
54:	080484f0	0 FUNC	GLOBAL	DEFAULT	16	__fini
55:	0804846c	13 FUNC	GLOBAL	DEFAULT	14	fun
56:	080496fc	0 NOTYPE	GLOBAL	DEFAULT	25	__data_start
57:	00000000	0 NOTYPE	WEAK	DEFAULT	UND	__gmon_start__
58:	08049700	0 OBJECT	GLOBAL	HIDDEN	25	__dso_handle
59:	0804850c	4 OBJECT	GLOBAL	DEFAULT	16	IO stdin used



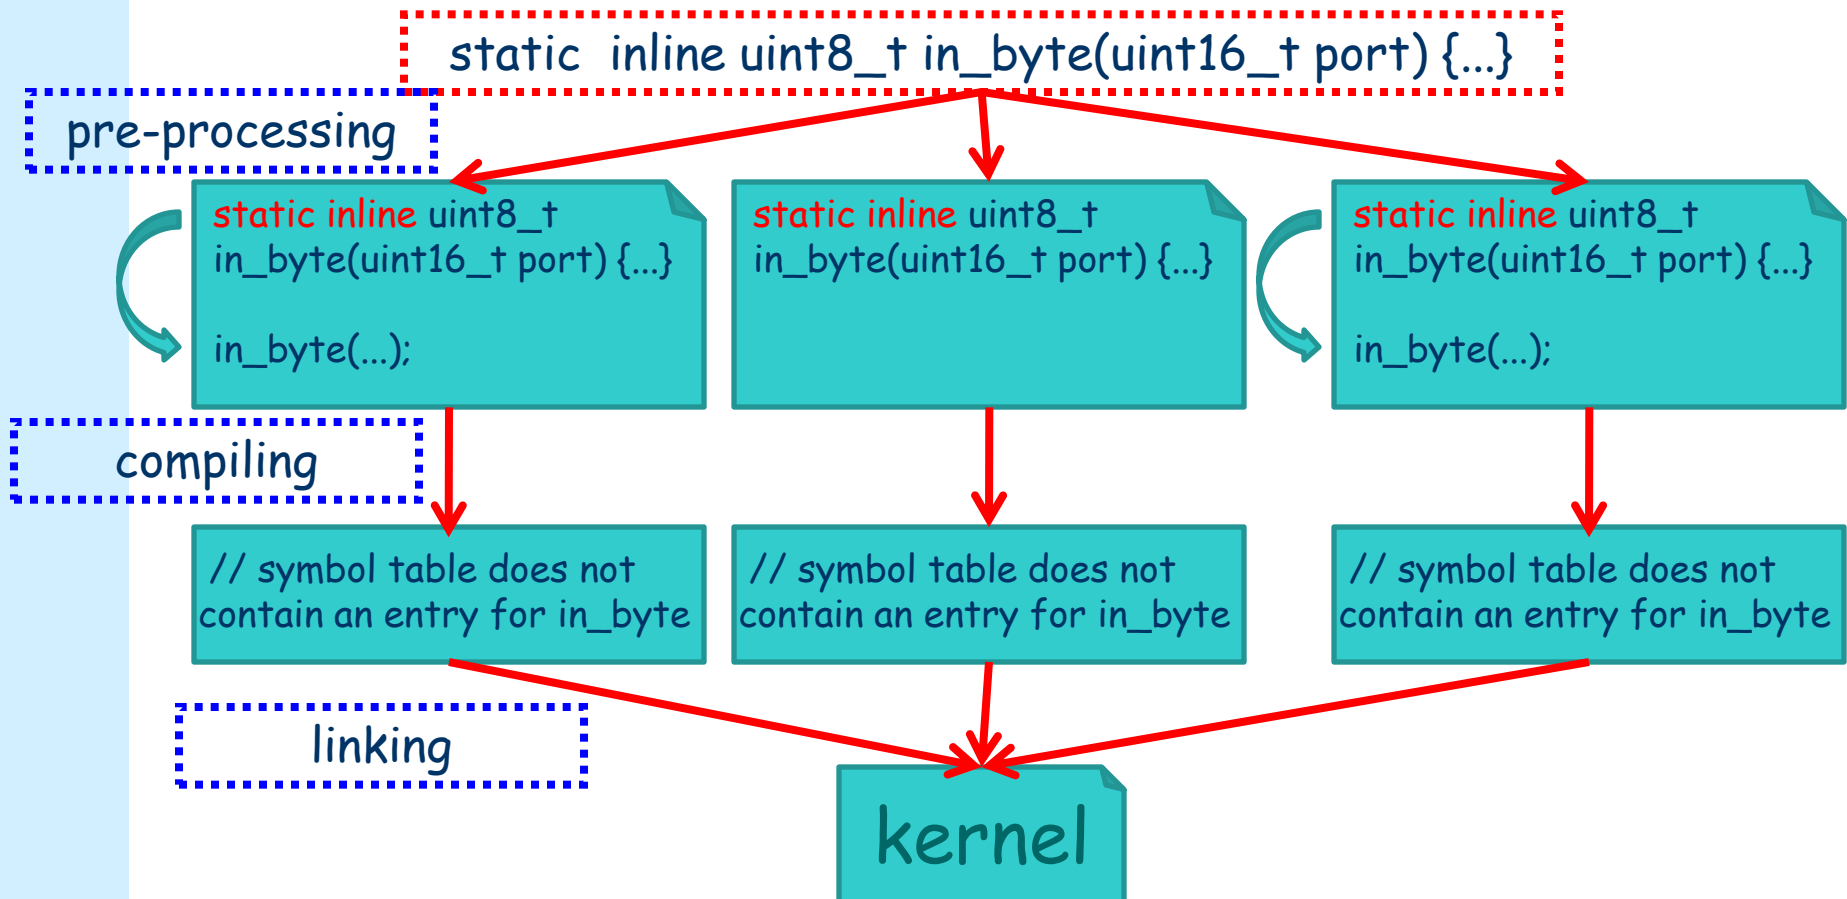
# inline

---

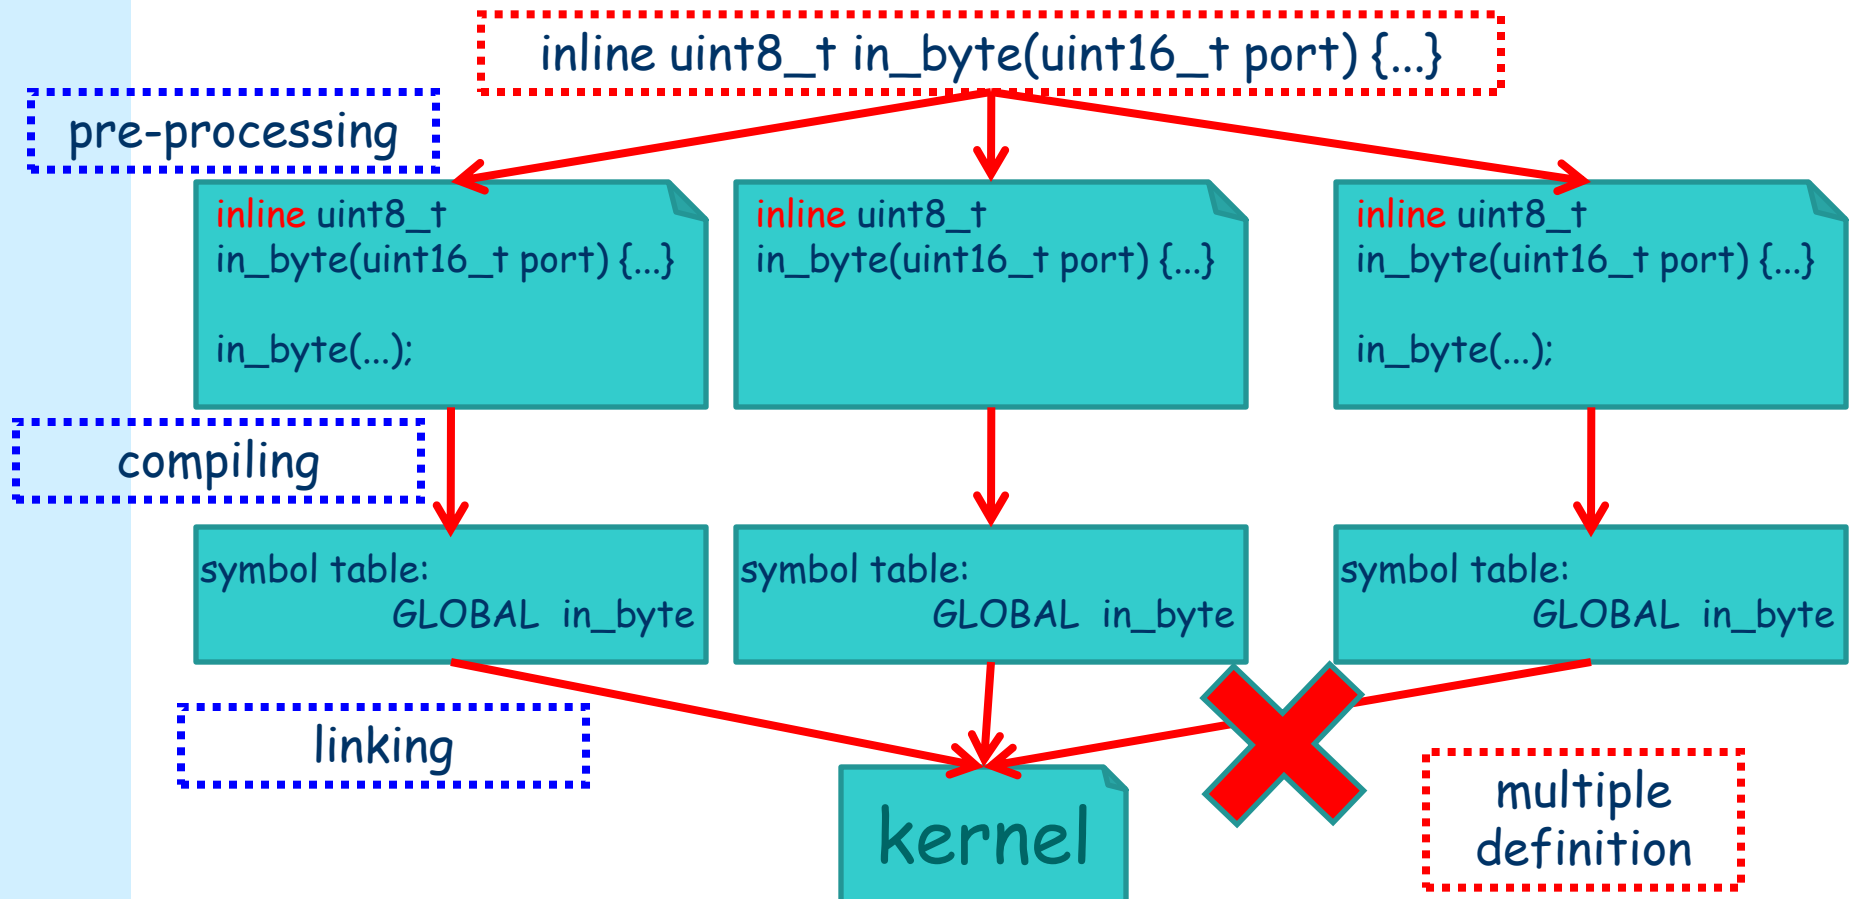
- Unroll the code of a function at calling points.
- After inlining, the body of the function will not exist.
- exceptions:
  - recursive function
  - global function
  - function pointer



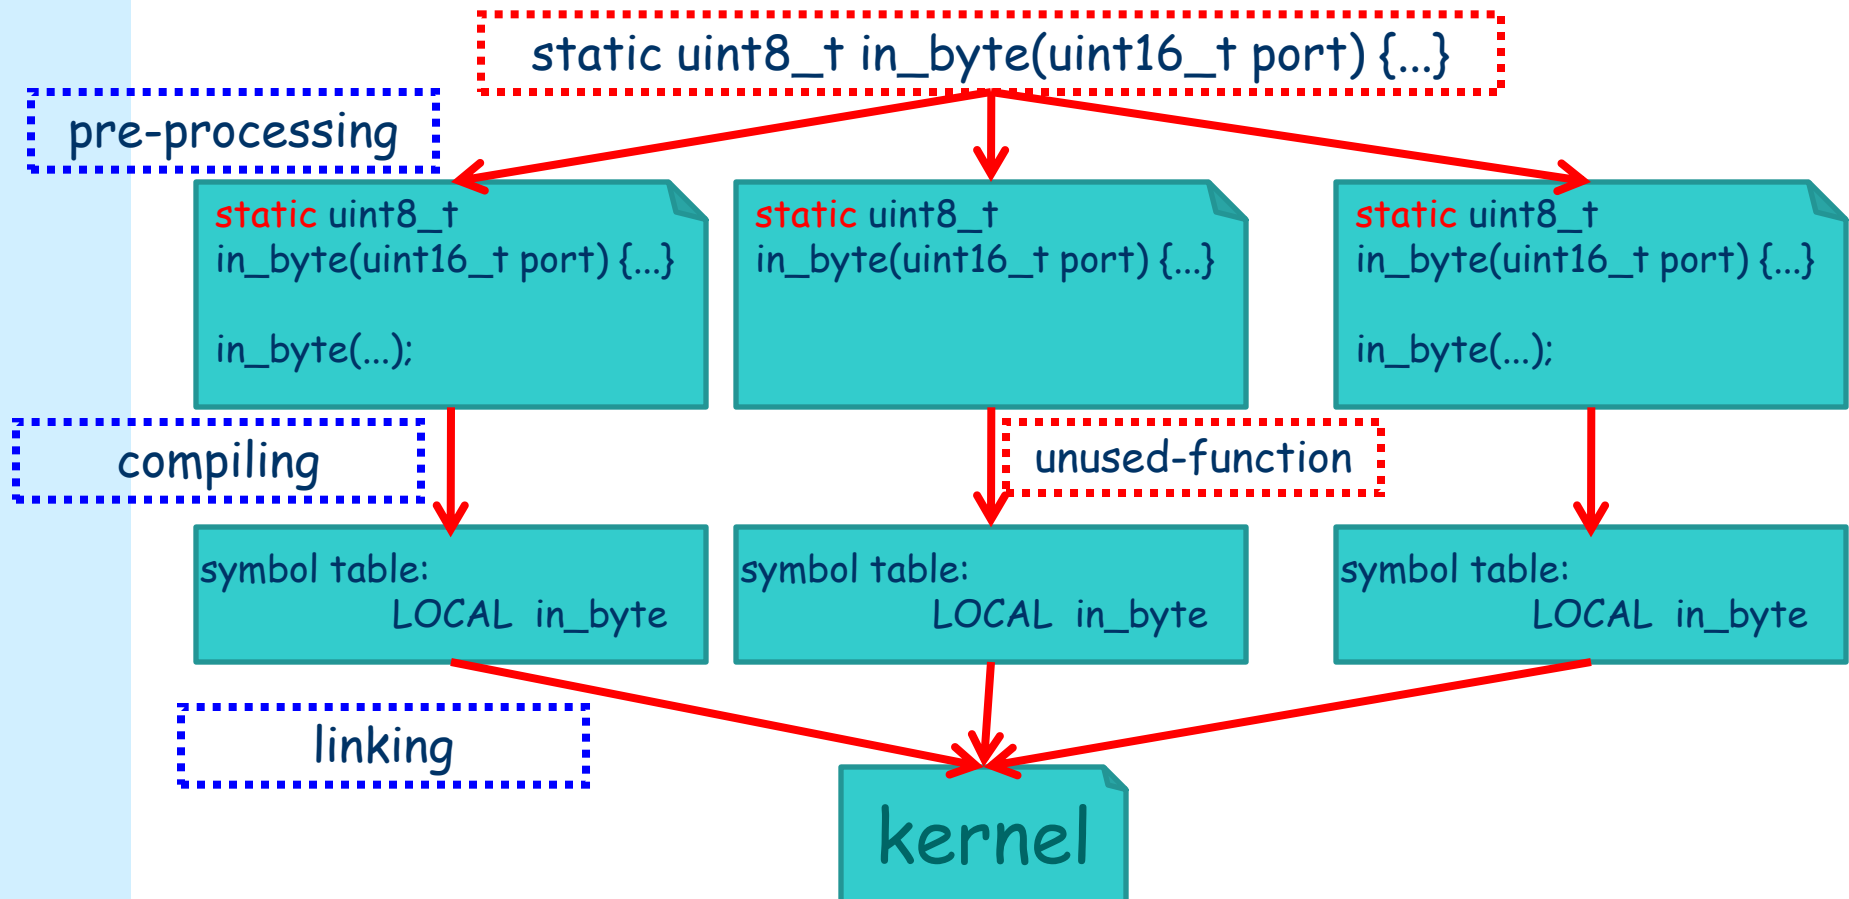
# With static & inline



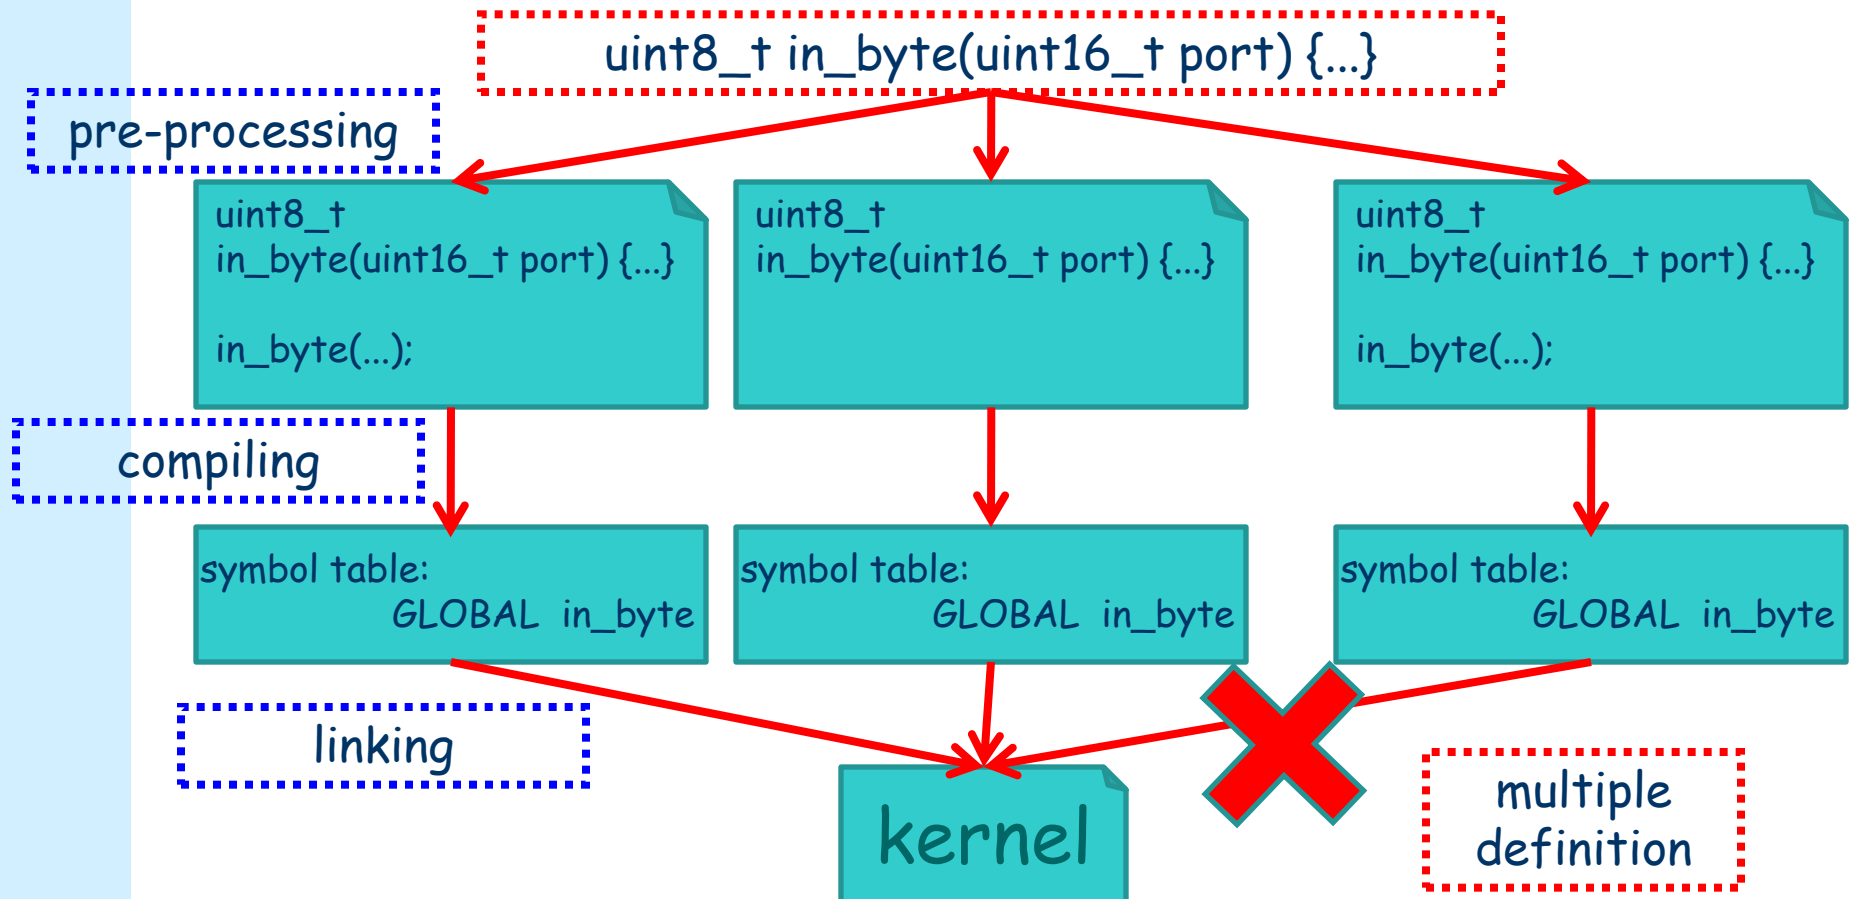
# Without static



# Without inline



# Without static & inline





# Unused-function in gcc

- `man gcc`

## **`-Wunused-function`**

Warn whenever a static function is declared but not defined or a non-inline static function is unused. This warning is enabled by **`-Wall`**.

- gcc has many options
  - only search for those you are interested in

# Compiling & linking

3. **编译与链接** 在include/common.h中添加一行

```
volatile static int dummy;
```

然后编译. 请问编译结果含有多少个dummy变量的实体? 你是如何得到这个结果的? 为什么会产生这样的结果?(Hint: 使用readelf命令. 回答完本题后可以删除添加的代码.)

- **see symbol table**

```
readelf -s kernel | grep -c dummy
```



# Why?

- “static” makes it local for each instance of “dummy”.
  - like the previous question
- “volatile” here is to prevent gcc for optimizing out “dummy”s.
- What about
  - `volatile static int dummy = 0;`
  - `volatile int dummy;`
  - `volatile int dummy = 0;`

4. 使用`man` gcc中的`-MD`选项有什么作用? `-Wall`和`-Werror`有什么作用? 为什么要使用`-Wall`和`-Werror`?

**-MD** **-MD** is equivalent to **-M -MF file**, except that **-E** is not implied. The driver determines file based on whether an **-o** option is given. If it is, the driver uses its argument but with a suffix of .d, otherwise it takes the name of the input file, removes any directory components and suffix, and applies a .d suffix.

If **-MD** is used in conjunction with **-E**, any **-o** switch is understood to specify the dependency output file, but if used without **-E**, each **-o** is understood to specify a target object file.

Since **-E** is not implied, **-MD** can be used to generate a dependency output file as a side-effect of the compilation process.

- Learn to use "man", learn to use everything.





# Makefile

5. 了解Makefile 在Makefile中有一行

```
-include $(OBJS:.o=.d)
```

请解释这行代码的功能.

- include the contents of all \*.d files
  - Where do \*.d files come from?
  - What do \*.d files contain?
  - Why do we include them?



# Makefile

- Where do \*.d files come from?

```
CFLAGS = -m32 -static -ggdb -MD -Wall -I./include -O2 \  
        -fno-builtin -fno-stack-protector -fno-omit-frame-pointer
```

- What do \*.d files contain?

```
src/kernel/main.o: src/kernel/main.c include/common.h include/types.h \  
include/const.h include/assert.h include/x86/x86.h include/x86/cpu.h \  
include/x86/memory.h include/x86/io.h include/common.h include/x86/x86.h \  
include/memory.h
```

- Why do we include them?
  - What happen if they are not included?



# Makefile

6. 了解Makefile 请描述你在终端敲入make后，make程序如何组织.c，.h和.S文件，最终生成disk.img。（这个问题包括两个方面：Makefile的工作方式和编译链接的过程。Hint：make过程中会用到 implicit rules.）
- fresh make & non-fresh make
  - plenty of details
    - variables, functions, include, implicit rules...
  - GNU Make Manual

fresh  
make

```
run: disk.img  
$(QEMU) -serial stdio disk.img
```

fresh  
make

```
run: disk.img  
$(QEMU) -serial stdio disk.img
```



```
disk.img: kernel  
@cd boot; make  
cat boot/bootblock kernel > disk.img
```

fresh  
make

```
run: disk.img  
$(QEMU) -serial stdio disk.img
```

① ↓

```
disk.img: kernel  
@cd boot; make  
cat boot/bootblock kernel > disk.img
```

② ↓

```
kernel: $(OBJS)  
$(LD) $(LDFLAGS) -e os_init -Ttext 0xC0100000 -o kernel $(OBJS)
```

# fresh make

```
run: disk.img  
$(QEMU) -serial stdio disk.img
```

1 ↓

```
disk.img: kernel  
@cd boot; make  
cat boot/bootblock kernel > disk.img
```

2 ↓

```
kernel: $(OBJS)  
$(LD) $(LDFLAGS) -e os_init -Ttext 0xC0100000 -o kernel $(OBJS)
```

3 ↓

```
main.o: main.c  
$(CC) $(CPPFLAGS)  
$(CFLAGS) -c -o main.o main.c
```

# fresh make

```
run: disk.img  
$(QEMU) -serial stdio disk.img
```

1 ↓

```
disk.img: kernel  
@cd boot; make  
cat boot/bootblock kernel > disk.img
```

2 ↓

```
kernel: $(OBJS)  
$(LD) $(LDFLAGS) -e os_init -Ttext 0xC0100000 -o kernel $(OBJS)
```

3 ↓

```
main.o: main.c  
$(CC) $(CPPFLAGS)  
$(CFLAGS) -c -o main.o main.c
```

4 ↓



main.c



# fresh make

```
run: disk.img  
$(QEMU) -serial stdio disk.img
```

① ↓

```
disk.img: kernel  
@cd boot; make  
cat boot/bootblock kernel > disk.img
```


② ↓

```
kernel: $(OBJS)  
$(LD) $(LDFLAGS) -e os_init -Ttext 0xC0100000 -o kernel $(OBJS)
```

③ ↓

```
main.o: main.c  
$(CC) $(CPPFLAGS)  
$(CFLAGS) -c -o main.o main.c
```

④ ↓    ⑤ ↑



main.c

# fresh make

```
run: disk.img  
$(QEMU) -serial stdio disk.img
```

1 ↓

```
disk.img: kernel  
@cd boot; make  
cat boot/bootblock kernel > disk.img
```

2 ↓

```
kernel: $(OBJS)  
$(LD) $(LDFLAGS) -e os_init -Ttext 0xC0100000 -o kernel $(OBJS)
```

3 ↓


```
main.o: main.c  
$(CC) $(CPPFLAGS)  
$(CFLAGS) -c -o main.o main.c
```

4 ↓




main.c

↑ 5



main.o

6 ↓



main.d

# fresh make

```
run: disk.img  
$(QEMU) -serial stdio disk.img
```

① ↓

```
disk.img: kernel  
@cd boot; make  
cat boot/bootblock kernel > disk.img
```

② ↓

```
kernel: $(OBJS)  
$(LD) $(LDFLAGS) -e os_init -Ttext 0xC0100000 -o kernel $(OBJS)
```

③ ↓

⑦ ↑

```
main.o: main.c  
$(CC) $(CPPFLAGS)  
$(CFLAGS) -c -o main.o main.c
```

④ ↓



main.c

⑤ ↑



main.o

⑥ ↓



main.d

# fresh make

```
run: disk.img  
$(QEMU) -serial stdio disk.img
```

1 ↓

```
disk.img: kernel  
@cd boot; make  
cat boot/bootblock kernel > disk.img
```

2 ↓

```
kernel: $(OBJS)  
$(LD) $(LDFLAGS) -e os_init -Ttext 0xC0100000 -o kernel $(OBJS)
```

3 ↓

7 ↑

```
main.o: main.c  
$(CC) $(CPPFLAGS)  
$(CFLAGS) -c -o main.o main.c
```

8 ...

4 ↓



main.c

5 ↑



main.o

6 ↓



main.d

# fresh make

```
run: disk.img  
$(QEMU) -serial stdio disk.img
```

① ↓

```
disk.img: kernel  
@cd boot; make  
cat boot/bootblock kernel > disk.img
```

② ↓

```
kernel: $(OBJS)  
$(LD) $(LDFLAGS) -e os_init -Ttext 0xC0100000 -o kernel $(OBJS)
```

③ ↓

⑦ ↑


```
main.o: main.c  
$(CC) $(CPPFLAGS)  
$(CFLAGS) -c -o main.o main.c
```

⑧ ...

⑨ ↓


```
do_irq.o: do_irq.S  
$(CC) $(ASFLAGS)  
-c -o do_irq.o do_irq.S
```

④ ↓




main.c

⑤ ↑



main.o

⑥ ↓



main.d

# fresh make

```
run: disk.img  
$(QEMU) -serial stdio disk.img
```

① ↓

```
disk.img: kernel  
@cd boot; make  
cat boot/bootblock kernel > disk.img
```

② ↓

```
kernel: $(OBJS)  
$(LD) $(LDFLAGS) -e os_init -Ttext 0xC0100000 -o kernel $(OBJS)
```

③ ↓

⑦ ↑

```
main.o: main.c  
$(CC) $(CPPFLAGS)  
$(CFLAGS) -c -o main.o main.c
```

⑧ ...

⑨ ↓

```
do_irq.o: do_irq.S  
$(CC) $(ASFLAGS)  
-c -o do_irq.o do_irq.S
```

④ ↓

⑤ ↑



main.c



main.o



main.d

⑥ ↓

⑩ ↓



do\_irq.S

# fresh make

```
run: disk.img  
$(QEMU) -serial stdio disk.img
```

① ↓

```
disk.img: kernel  
@cd boot; make  
cat boot/bootblock kernel > disk.img
```

② ↓

```
kernel: $(OBJS)  
$(LD) $(LDFLAGS) -e os_init -Ttext 0xC0100000 -o kernel $(OBJS)
```

③ ↓

⑦ ↑

```
main.o: main.c  
$(CC) $(CPPFLAGS)  
$(CFLAGS) -c -o main.o main.c
```

⑧ ...

⑨ ↓

```
do_irq.o: do_irq.S  
$(CC) $(ASFLAGS)  
-c -o do_irq.o do_irq.S
```

④ ↓

⑤ ↑



main.c



main.o



main.d

⑥ ↓

⑩ ↓

⑪ ↑



do\_irq.S

# fresh make

```
run: disk.img  
$(QEMU) -serial stdio disk.img
```

① ↓

```
disk.img: kernel  
@cd boot; make  
cat boot/bootblock kernel > disk.img
```

② ↓

```
kernel: $(OBJS)  
$(LD) $(LDFLAGS) -e os_init -Ttext 0xC0100000 -o kernel $(OBJS)
```

③ ↓

⑦ ↑

```
main.o: main.c  
$(CC) $(CPPFLAGS)  
$(CFLAGS) -c -o main.o main.c
```

⑧ ...

⑨ ↓

```
do_irq.o: do_irq.S  
$(CC) $(ASFLAGS)  
-c -o do_irq.o do_irq.S
```

④ ↓

⑤ ↑



main.c



main.o



main.d


⑥

⑩ ↓


⑪ ↑



do\_irq.S



do\_irq.o



do\_irq.d

⑫



# fresh make

```
run: disk.img  
$(QEMU) -serial stdio disk.img
```

① ↓

```
disk.img: kernel  
@cd boot; make  
cat boot/bootblock kernel > disk.img
```

② ↓

```
kernel: $(OBJS)  
$(LD) $(LDFLAGS) -e os_init -Ttext 0xC0100000 -o kernel $(OBJS)
```

③ ↓

⑦ ↑

```
main.o: main.c  
$(CC) $(CPPFLAGS)  
$(CFLAGS) -c -o main.o main.c
```

⑧ ...

⑨ ↓

⑬ ↑

```
do_irq.o: do_irq.S  
$(CC) $(ASFLAGS)  
-c -o do_irq.o do_irq.S
```

④ ↓

⑤ ↑



main.c



main.o



main.d

⑥ ↓

⑩ ↓


⑪ ↑



do\_irq.S



do\_irq.o



do\_irq.d

⑫ ↓

# fresh make


```
run: disk.img  
$(QEMU) -serial stdio disk.img
```

1

```
disk.img: kernel  
@cd boot; make  
cat boot/bootblock kernel > disk.img
```

2

```
kernel: $(OBJS)  
$(LD) $(LDFLAGS) -e os_init -Ttext 0xC0100000 -o kernel $(OBJS)
```



3

7

9

13

14

```
main.o: main.c  
$(CC) $(CPPFLAGS)  
$(CFLAGS) -c -o main.o main.c
```

8

```
do_irq.o: do_irq.S  
$(CC) $(ASFLAGS)  
-c -o do_irq.o do_irq.S
```

4


5

6


10

11


12




main.c



main.o




main.d



do\_irq.S

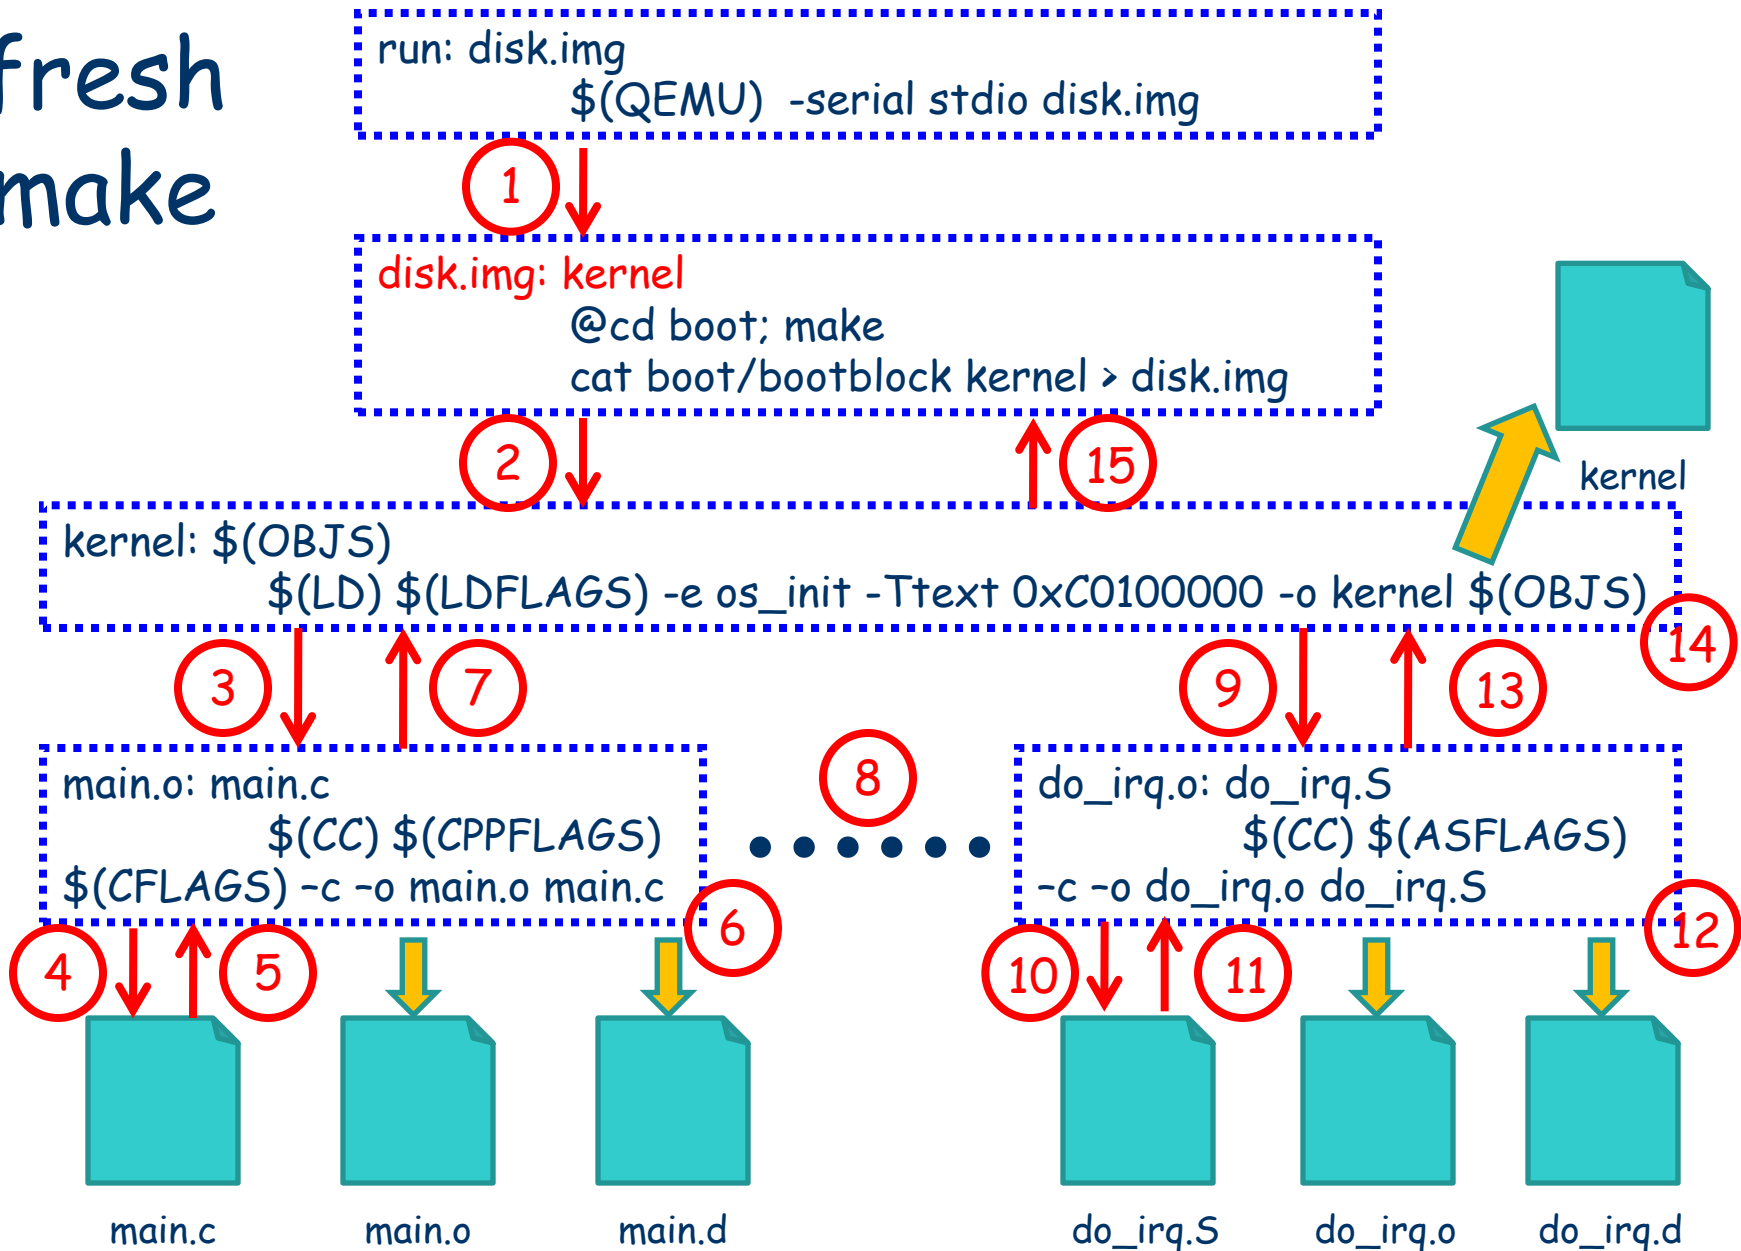


do\_irq.o

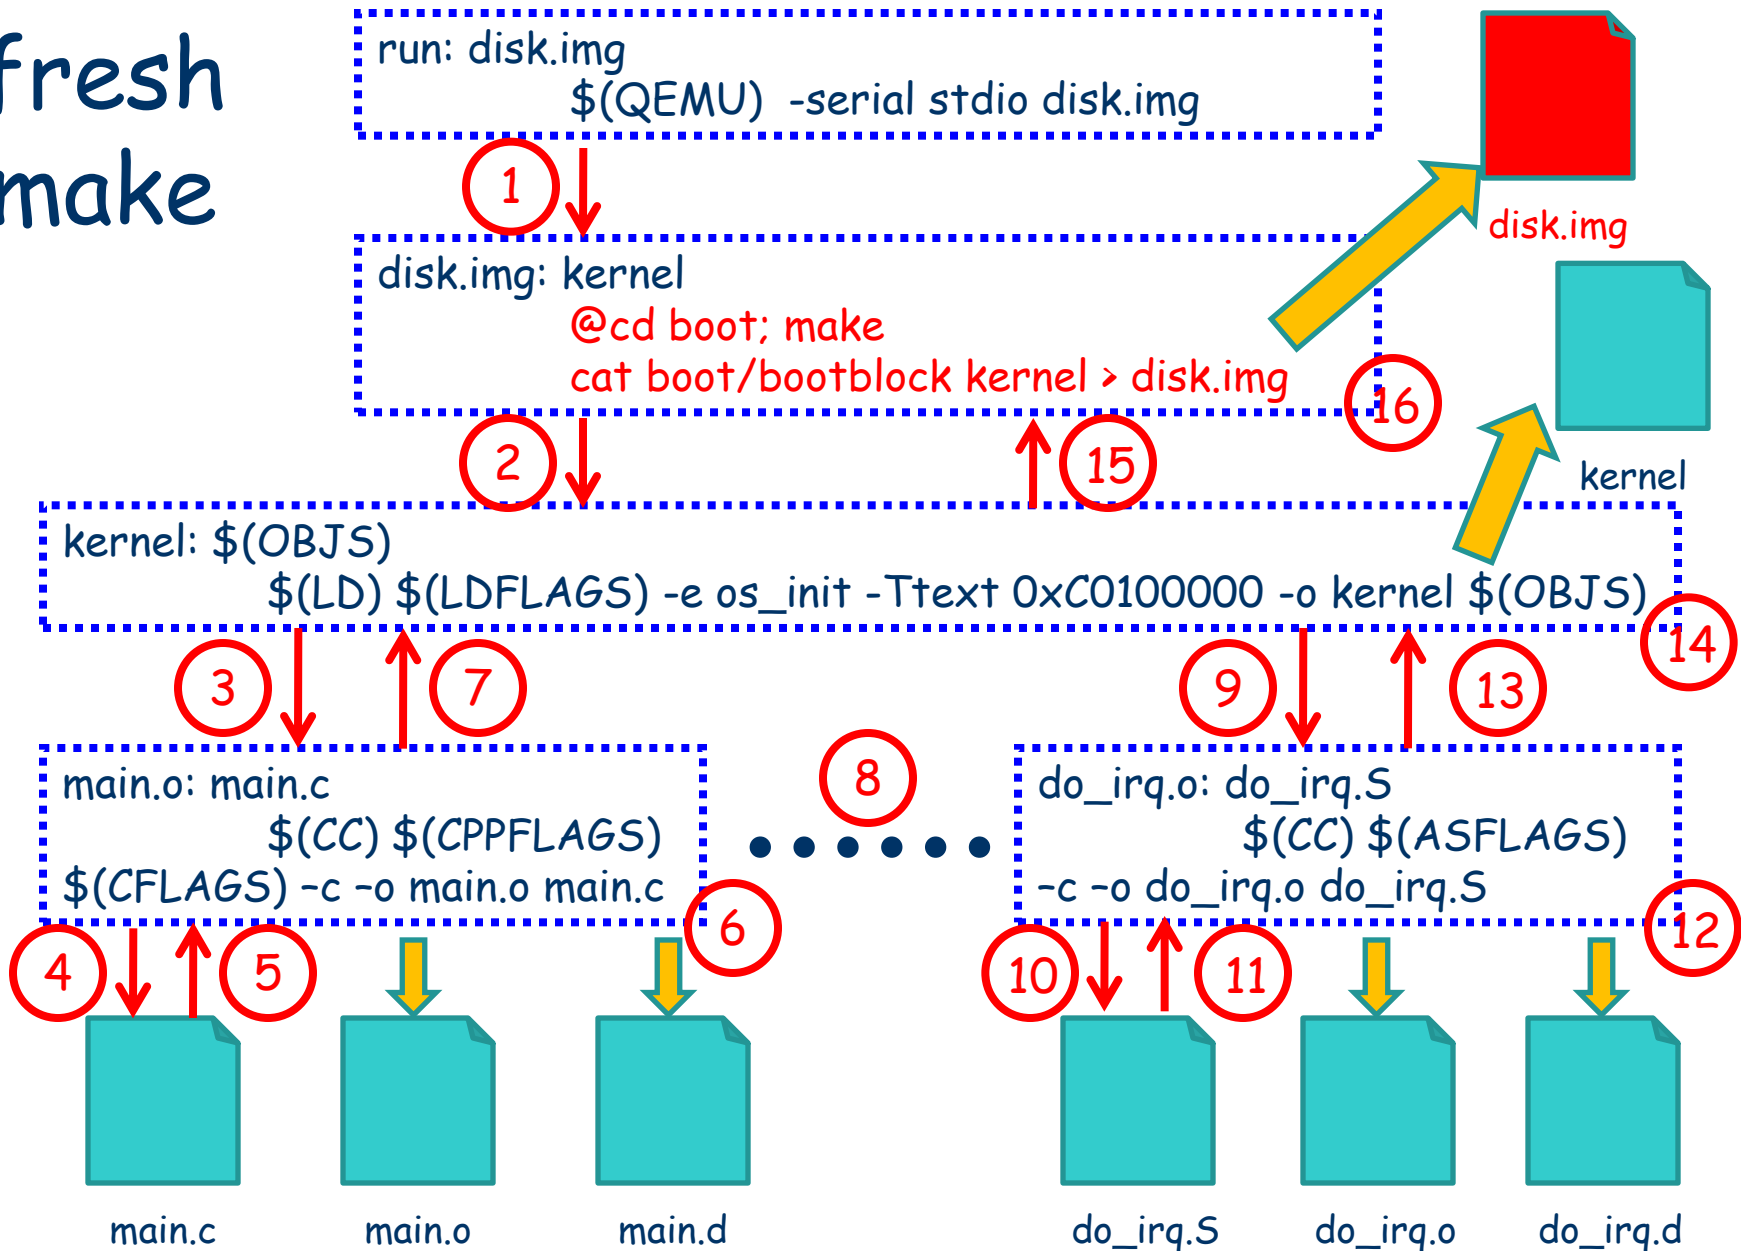


do\_irq.d

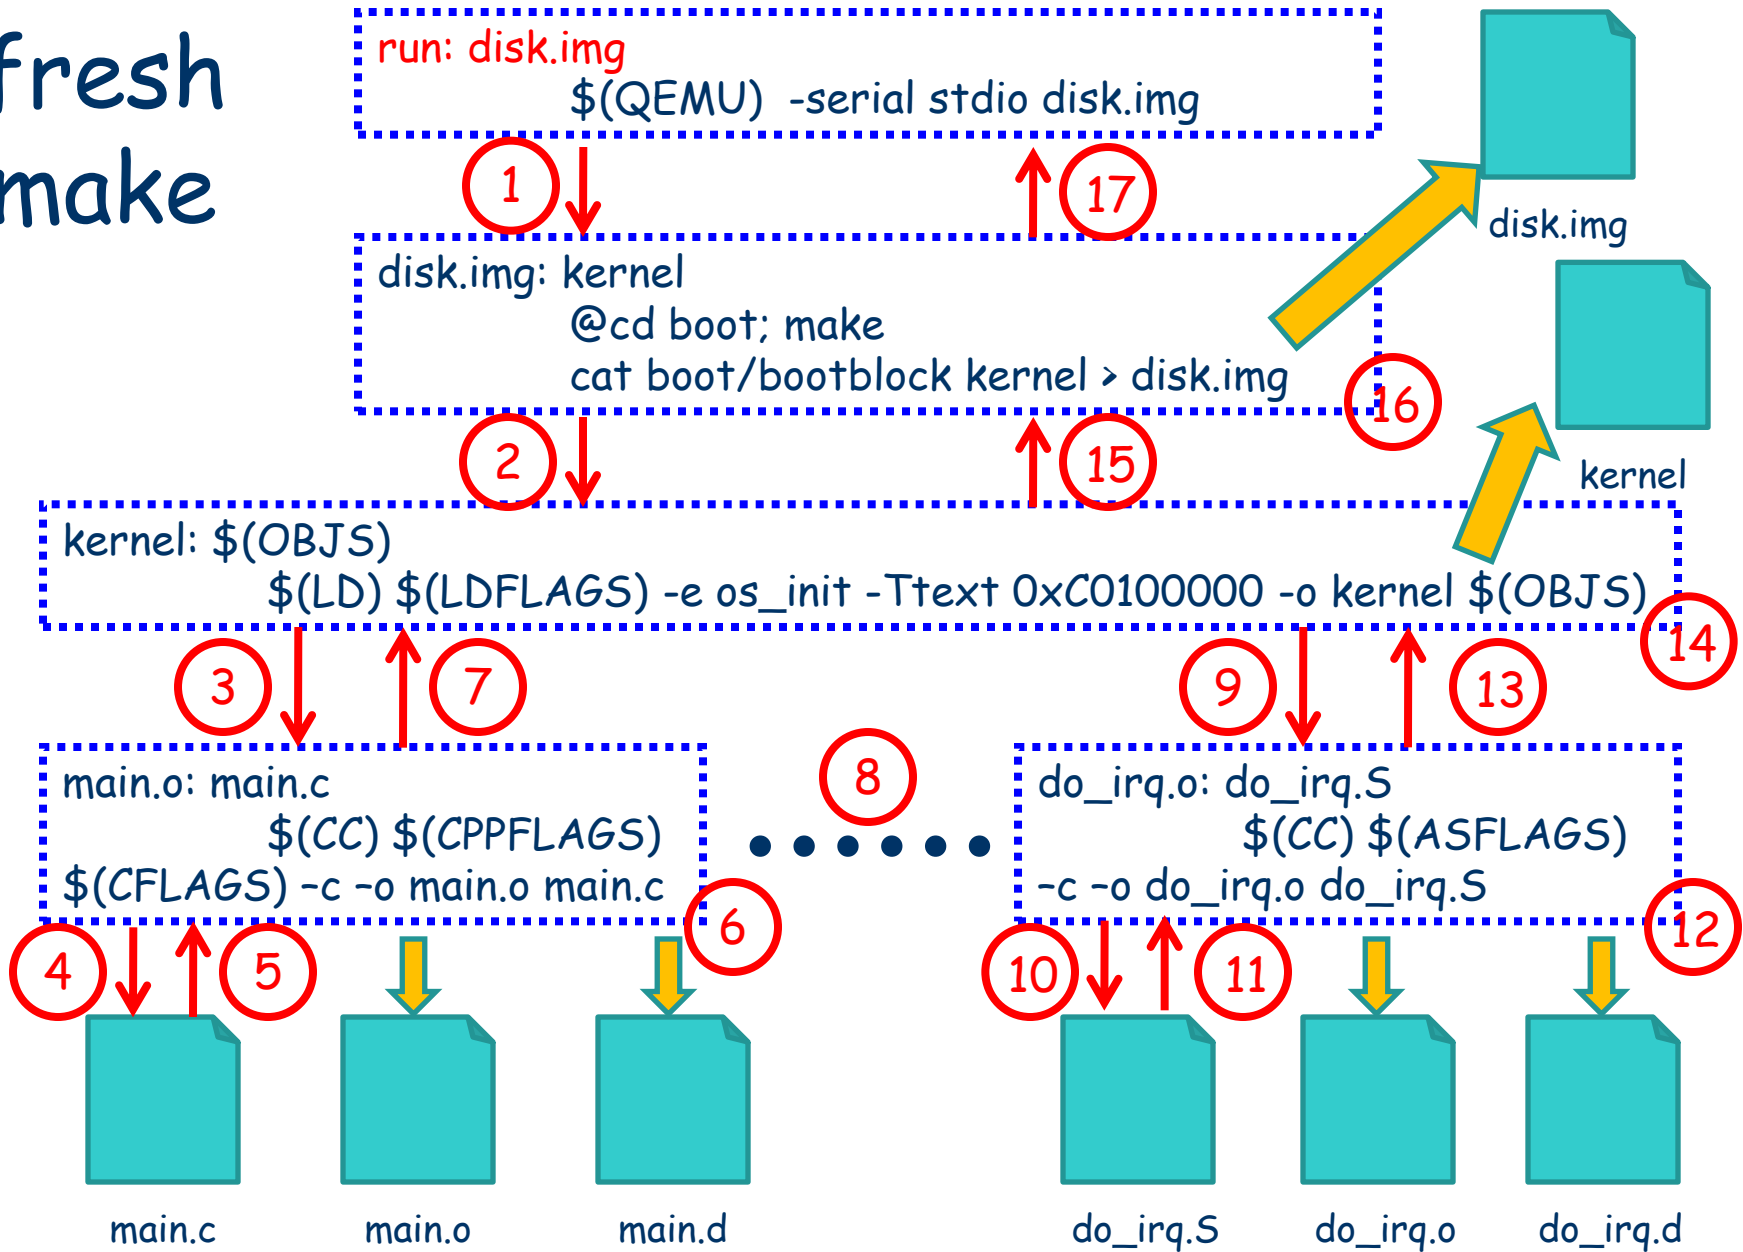
# fresh make



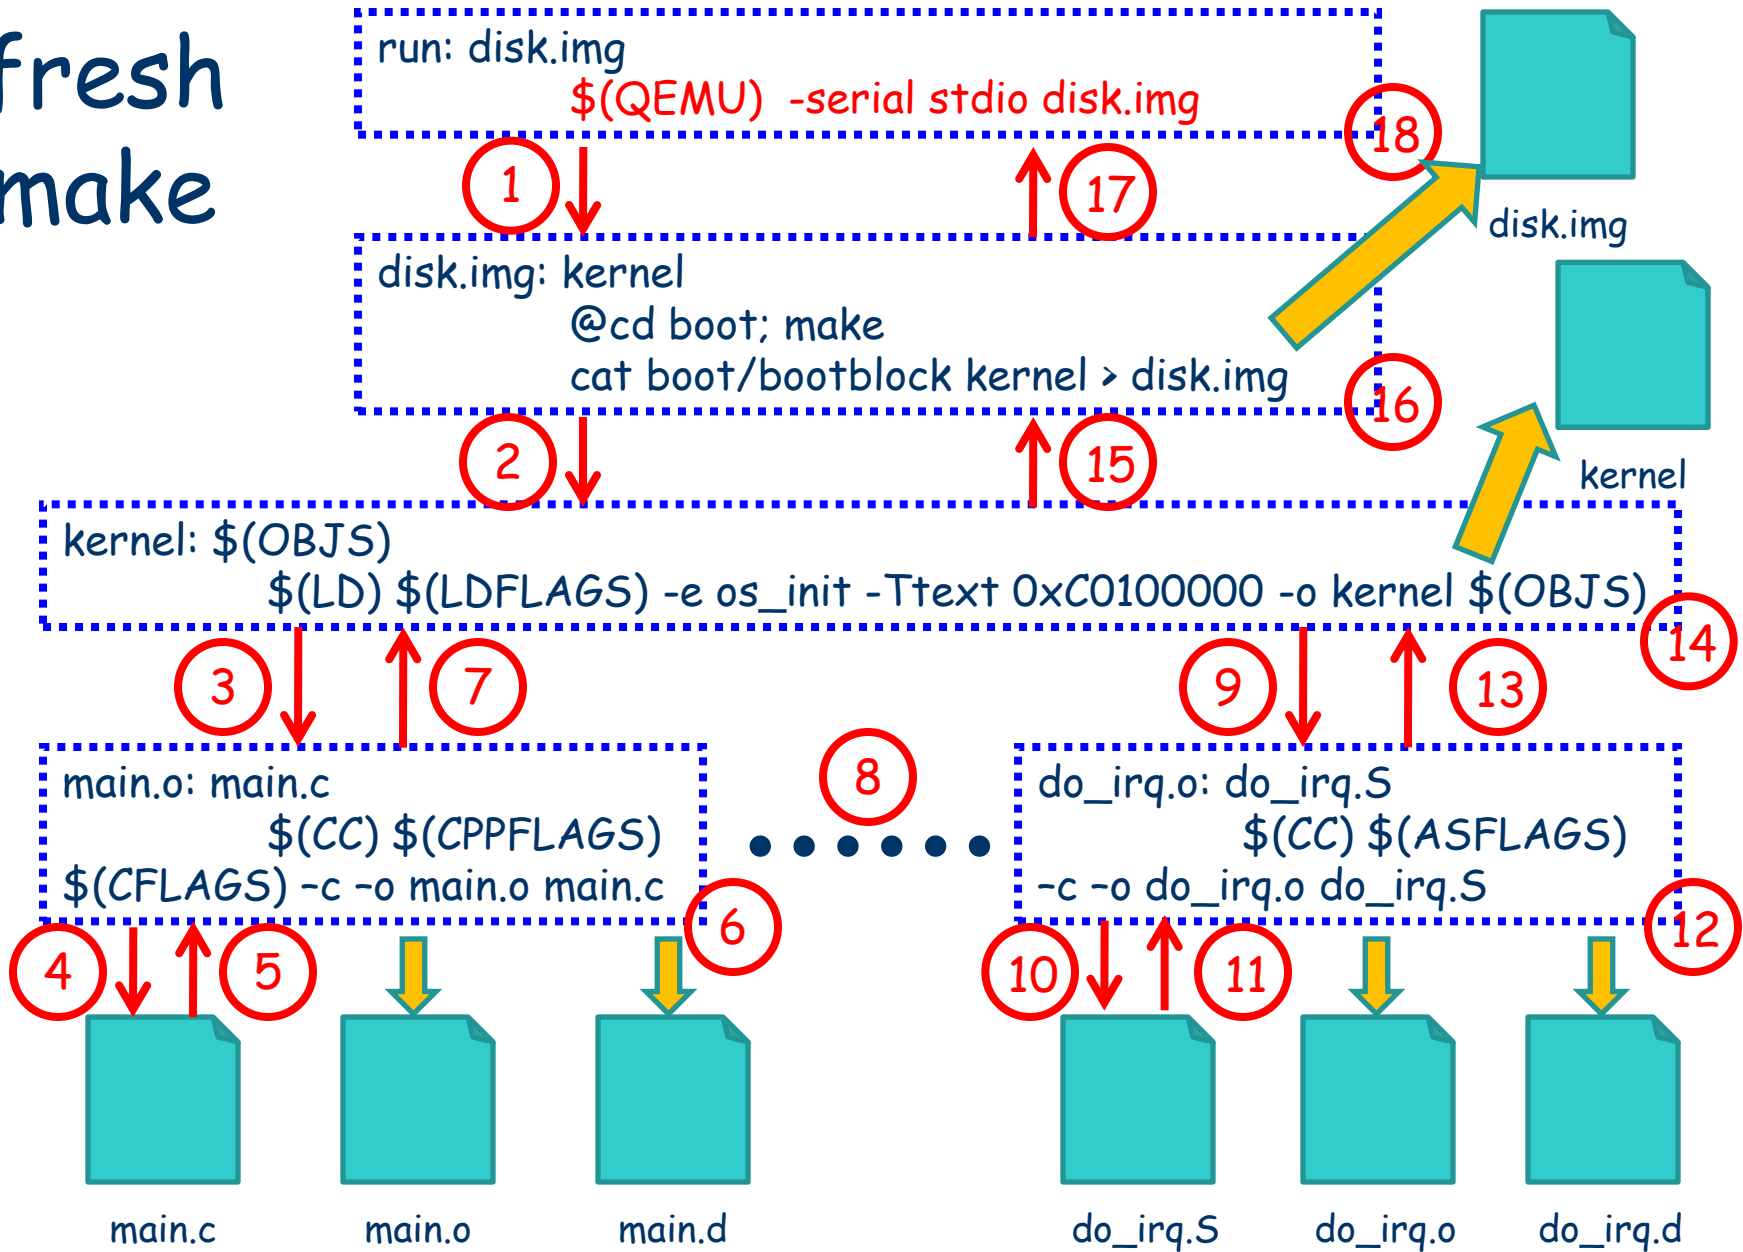
# fresh make



fresh  
make



fresh  
make



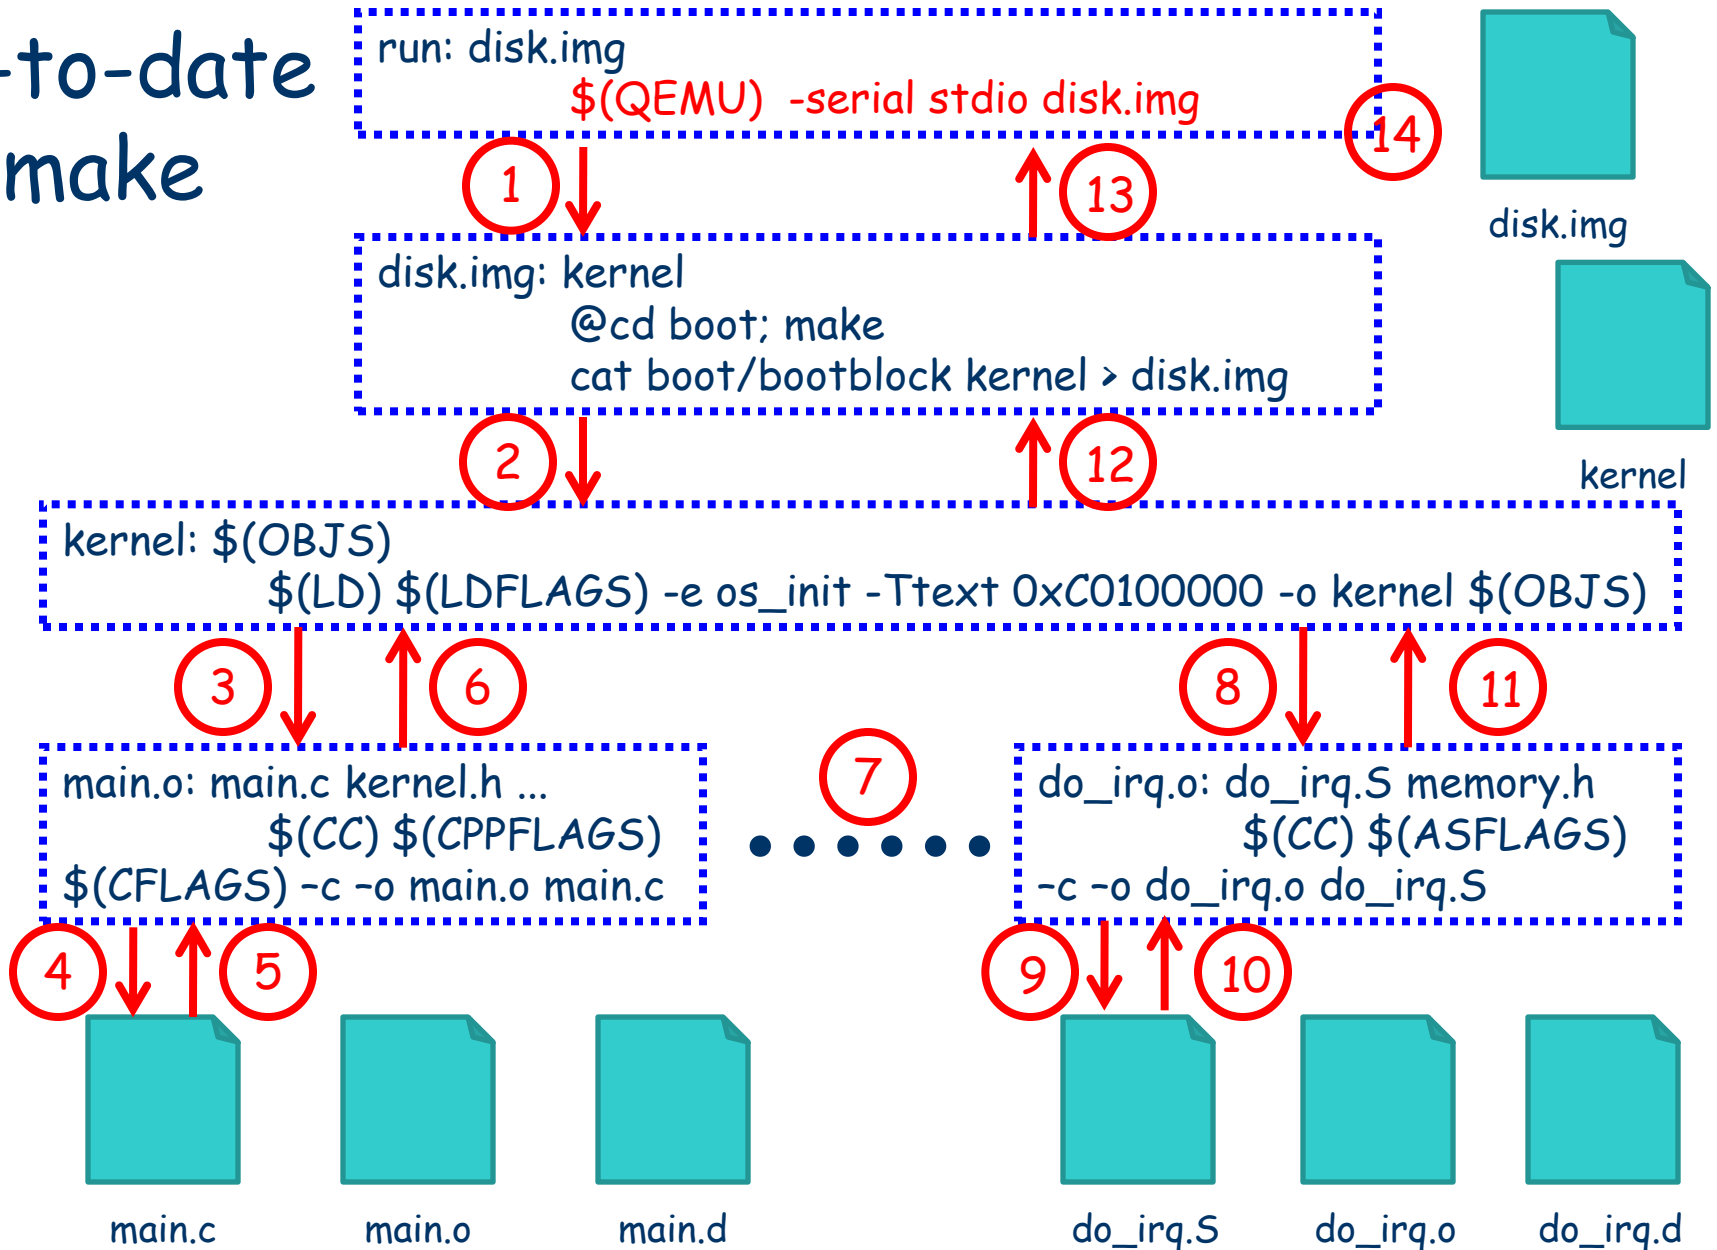


# Makefile

---

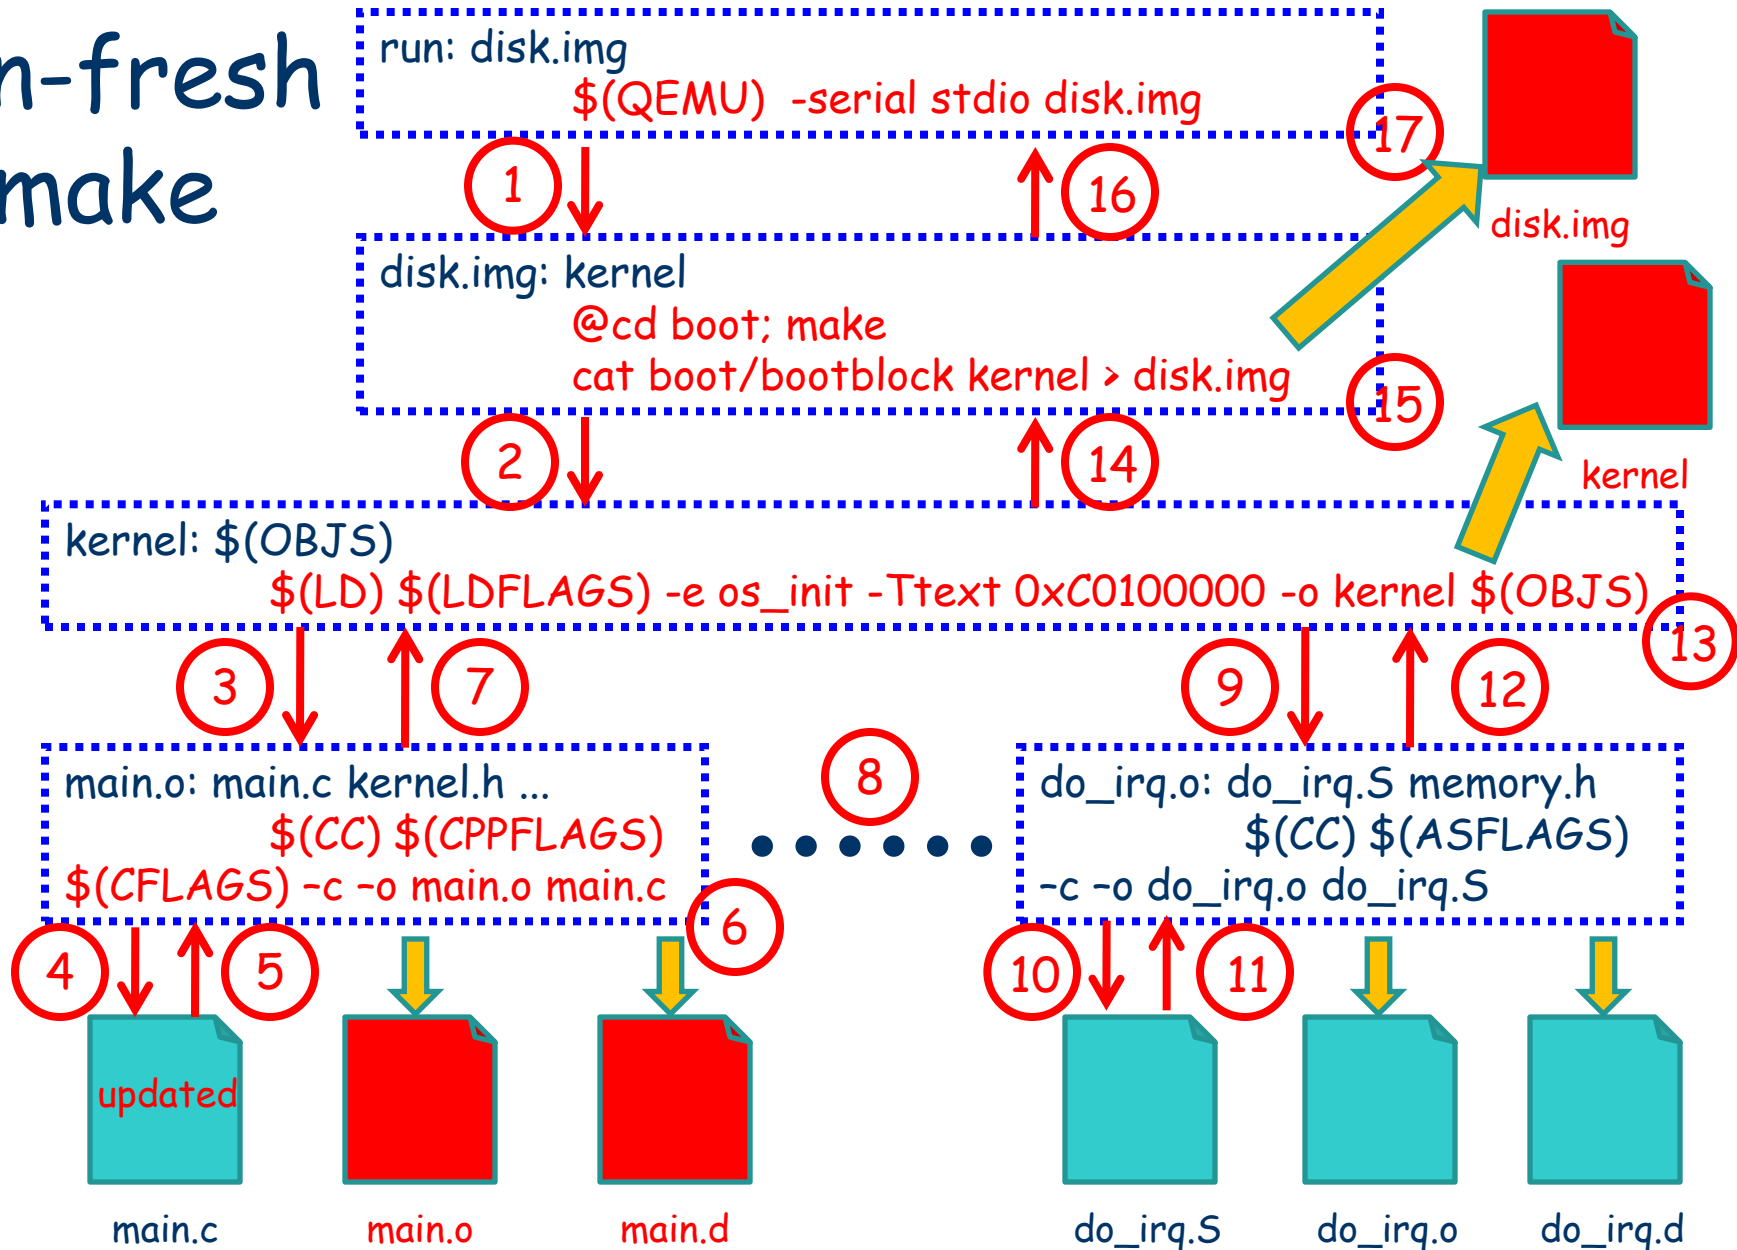
- What happen when
  - no file is updated?
  - a source file is updated?
  - a header file is updated?
  - boot/main.c is updated?
- Why “\$(QEMU) -serial stdio disk.img” always executes?

up-to-date  
make





non-fresh  
make





# volatile

7. 理解volatile 在include/x86/cpu.h中有两个函数write\_gdtr()和write\_idtr(), 请选择其中一个函数来回答下列问题: 函数体中有一行

```
static volatile uint16_t data[3];
```

尝试去掉volatile后重新编译并运行, 你发现了什么问题? 请对比去掉volatile前后编译结果的不同, 并解释为什么去掉volatile之后会导致运行出现问题. (Hint: 使用objdump命令, 你可能还需要查阅i386手册. 如果你使用gcc 4.7.3的版本, 你可能不会观察到运行之后出现的问题, 但你仍然可以对比编译结果的不同.)

- **mysterious reboot**

- may not occur with gcc version later than 4.7.2

# volatile

Terminal

File Edit View Search Terminal Help

```
109 c010012d: e8 1e 09 00 00    call    c0100a50 <memset>
110 c0100132: b8 80 50 12 c0    mov     $0xc0125080,%eax
111 c0100137: 66 a3 b2 50 12 c0 mov     %ax,0xc01250b2
112 c010013d: c1 e8 10          shr     $0x10,%eax
113 c0100140: 66 a3 b4 50 12 c0 mov     %ax,0xc01250b4
114 c0100146: b8 b0 50 12 c0    mov     $0xc01250b0,%eax
115 c010014b: 0f 01 10          lgdtl   (%eax)
116 c010014e: b8 00 50 12 c0    mov     $0xc0125000,%eax
117 c0100153: 89 c2            mov     %eax,%edx
118 c0100155: 66 a3 aa 50 12 c0 mov     %ax,0xc01250aa
119 c010015b: c1 ea 10          shr     $0x10,%edx
120 c010015e: c1 e8 18          shr     $0x18,%eax
121 c0100161: 88 15 ac 50 12 c0 mov     %dl,0xc01250ac
122 c0100167: 66 c7 05 88 50 12 c0 movw    $0xffff,0xc0125088
123 c010016e: ff ff
124 c0100170: 66 c7 05 8a 50 12 c0 movw    $0x0,0xc012508a
125 c0100177: 00 00
126 c0100179: c6 05 8c 50 12 c0 00 movb    $0x0,0xc012508c
127 c0100180: c6 05 8d 50 12 c0 9a movb    $0x9a,0xc012508d
128 c0100187: c6 05 8e 50 12 c0 cf movb    $0xcf,0xc012508e
129 c010018e: c6 05 8f 50 12 c0 00 movb    $0x0,0xc012508f
130 c0100195: 66 c7 05 90 50 12 c0 movw    $0xffff,0xc0125090
131 c010019c: ff ff
132 c010019e: 66 c7 05 92 50 12 c0 movw    $0x0,0xc0125092
133 c01001a5: 00 00
134 c01001a7: c6 05 94 50 12 c0 00 movb    $0x0,0xc0125094
135 c01001ae: c6 05 95 50 12 c0 92 movb    $0x92,0xc0125095
136 c01001b5: c6 05 96 50 12 c0 cf movb    $0xcf,0xc0125096
137 c01001bc: c6 05 97 50 12 c0 00 movb    $0x0,0xc0125097
138 c01001c3: 66 c7 05 98 50 12 c0 movw    $0xffff,0xc0125098
139 c01001ca: ff ff
140 c01001cc: 66 c7 05 9a 50 12 c0 movw    $0x0,0xc012509a
141 c01001d3: 00 00
142 c01001d5: c6 05 9c 50 12 c0 00 movb    $0x0,0xc012509c
143 c01001dc: c6 05 9d 50 12 c0 fa movb    $0xfa,0xc012509d
144 c01001e3: c6 05 9e 50 12 c0 cf movb    $0xcf,0xc012509e
145 c01001ea: c6 05 9f 50 12 c0 00 movb    $0x0,0xc012509f
146 c01001f1: 66 c7 05 a0 50 12 c0 movw    $0xffff,0xc01250a0
147 c01001f8: ff ff
148 c01001fa: 66 c7 05 a2 50 12 c0 movw    $0x0,0xc01250a2
149 c0100201: 00 00
150 c0100203: c6 05 a4 50 12 c0 00 movb    $0x0,0xc01250a4
151 c010020a: c6 05 a5 50 12 c0 f2 movb    $0xf2,0xc01250a5
152 c0100211: c6 05 a6 50 12 c0 cf movb    $0xcf,0xc01250a6
153 c0100218: c6 05 a7 50 12 c0 00 movb    $0x0,0xc01250a7
154 c010021f: 66 c7 05 b0 50 12 c0 movw    $0x2f,0xc01250b0
155 c0100226: 2f 00
156 c0100228: c7 05 08 50 12 c0 10 movl    $0x10,0xc0125008
157 c010022f: 00 00 00
158 c0100232: 66 c7 05 a8 50 12 c0 movw    $0x63,0xc01250a8
159 c0100239: 63 00
```

without  
"volatile"

```
109 c010012d: e8 1e 09 00 00    call    c0100a50 <memset>
110 c0100132: b8 80 50 12 c0    mov     $0xc0125080,%eax
111 c0100137: 66 c7 05 80 50 12 c0 movw    $0x2f,0xc0103000
112 c010013e: 2f 00
113 c0100140: 66 a3 02 30 10 c0 mov     %ax,0xc0103002
114 c0100146: c1 e8 10          shr     $0x10,%eax
115 c0100149: 66 a3 04 30 10 c0 mov     %ax,0xc0103004
116 c010014f: b8 00 30 10 c0    mov     $0xc0103000,%eax
117 c0100154: 0f 01 10          lgdtl   (%eax)
118 c0100157: b8 00 50 12 c0    mov     $0xc0125000,%eax
119 c010015c: 89 c2            mov     %eax,%edx
120 c010015e: 66 a3 aa 50 12 c0 mov     %ax,0xc01250aa
121 c0100164: c1 ea 10          shr     $0x10,%edx
122 c0100167: c1 e8 18          shr     $0x18,%eax
123 c010016a: 88 15 ac 50 12 c0 mov     %dl,0xc01250ac
124 c0100170: 66 c7 05 88 50 12 c0 movw    $0xffff,0xc0125088
125 c0100177: ff ff
126 c0100179: 66 c7 05 8a 50 12 c0 movw    $0x0,0xc012508a
127 c0100180: 00 00
128 c0100182: c6 05 8c 50 12 c0 00 movb    $0x0,0xc012508c
129 c0100189: c6 05 8d 50 12 c0 9a movb    $0x9a,0xc012508d
130 c0100190: c6 05 8e 50 12 c0 cf movb    $0xcf,0xc012508e
131 c0100197: c6 05 8f 50 12 c0 00 movb    $0x0,0xc012508f
132 c010019e: 66 c7 05 90 50 12 c0 movw    $0xffff,0xc0125090
133 c01001a5: ff ff
134 c01001a7: 66 c7 05 92 50 12 c0 movw    $0x0,0xc0125092
135 c01001ae: 00 00
136 c01001b0: c6 05 94 50 12 c0 00 movb    $0x0,0xc0125094
137 c01001b7: c6 05 95 50 12 c0 92 movb    $0x92,0xc0125095
138 c01001be: c6 05 96 50 12 c0 cf movb    $0xcf,0xc0125096
139 c01001c5: c6 05 97 50 12 c0 00 movb    $0x0,0xc0125097
140 c01001cc: 66 c7 05 98 50 12 c0 movw    $0xffff,0xc0125098
141 c01001d3: ff ff
142 c01001d5: 66 c7 05 9a 50 12 c0 movw    $0x0,0xc012509a
143 c01001dc: 00 00
144 c01001de: c6 05 9c 50 12 c0 00 movb    $0x0,0xc012509c
145 c01001e5: c6 05 9d 50 12 c0 fa movb    $0xfa,0xc012509d
146 c01001ec: c6 05 9e 50 12 c0 cf movb    $0xcf,0xc012509e
147 c01001f3: c6 05 9f 50 12 c0 00 movb    $0x0,0xc012509f
148 c01001fa: 66 c7 05 a0 50 12 c0 movw    $0xffff,0xc01250a0
149 c0100201: ff ff
150 c0100203: 66 c7 05 a2 50 12 c0 movw    $0x0,0xc01250a2
151 c010020a: 00 00
152 c010020c: c6 05 a4 50 12 c0 00 movb    $0x0,0xc01250a4
153 c0100213: c6 05 a5 50 12 c0 f2 movb    $0xf2,0xc01250a5
154 c010021a: c6 05 a6 50 12 c0 cf movb    $0xcf,0xc01250a6
155 c0100221: c6 05 a7 50 12 c0 00 movb    $0x0,0xc01250a7
156 c0100228: c7 05 08 50 12 c0 10 movl    $0x10,0xc0125008
157 c010022f: 00 00 00
158 c0100232: 66 c7 05 a8 50 12 c0 movw    $0x63,0xc01250a8
159 c0100239: 63 00
```

with  
"volatile"

~/project/2012os/os-lab1/code-bad.txt[1]

[text] unix utf-8 Ln 109, Col 33/10430\

~/project/2012os/os-lab1/code.txt[2]

[text] unix utf-8 Ln 109, Col 33/10424\


0 vimdiff\*

\*debian\* 16:24 21-Apr-14

# volatile

- Why mysterious reboot?
- What is data[0] for GDTR/IDTR?
  - P.156 in "i386 manual"

Figure 9-2. Pseudo-Descriptor Format for LIDT and SIDT



# Loading kernel

8. 加载内核 在boot/main.c的bootmain()函数中, 最后计算出entry的值. 这个值是多少? 尝试使用不同的方法获取这个值.

```
1 ELF Header:
2 Magic: 7f 45 4c 46 01 01 01 00 00 00 00 00 00 00 00 00
3 Class: ELF32
4 Data: 2's complement, little endian
5 Version: 1 (current)
6 OS/ABI: UNIX - System V
7 ABI Version: 0
8 Type: EXEC (Executable file)
9 Machine: Intel 80386
10 Version: 0x1
11 Entry point address: 0xc01002c0
12 Start of program headers: 52 (bytes into file)
13 Start of section headers: 133568 (bytes into file)
14 Flags: 0x0
15 Size of this header: 52 (bytes)
16 Size of program headers: 32 (bytes)
17 Number of program headers: 3
18 Size of section headers: 40 (bytes)
19 Number of section headers: 17
20 Section header string table index: 14
```

- gdb
- readelf

- How does entry point generate?



# Loading kernel

9. **加载内核** 在boot/main.c的bootmain()函数中, 有两处代码需要减去KOFFSET的值(分别在计算pa和entry时), 但在Lab0中相应代码并没有减去KOFFSET. 请尝试去掉这两处减去KOFFSET的操作, 然后重新编译并运行, 你发现了什么问题? 请解释为什么在Lab1中需要减去KOFFSET的操作.

```
$(LD) $(LDFLAGS) -e os_init -Ttext 0xc0100000 -o kernel $(OBJS)
```

- kernel thinks it is located at 0xc0100000
  - makes the kernel mapping identical to the one in Linux
    - why?
  - but MBR loads the kernel at 0x100000
- All addresses in the kernel binary are virtual addresses.

# Loading kernel


- What about...

```
for(; ph < eph; ph++) {  
-     pa = (unsigned char*)(ph->paddr - KOFFSET); /* physical address */  
+     pa = (unsigned char*)(ph->paddr); /* physical address */  
    readseg(pa, ph->filesz, ph->off); /* load from disk */  
    for (i = pa + ph->filesz; i < pa + ph->memsz; *i++ = 0);  
}  
/* Here we go! */  
entry = (void(*) (void))(elf->entry - KOFFSET);
```

```
for(; ph < eph; ph++) {  
    pa = (unsigned char*)(ph->paddr - KOFFSET); /* physical address */  
    readseg(pa, ph->filesz, ph->off); /* load from disk */  
    for (i = pa + ph->filesz; i < pa + ph->memsz; *i++ = 0);  
}  
/* Here we go! */  
-entry = (void(*) (void))(elf->entry - KOFFSET);  
+entry = (void(*) (void))(elf->entry);
```



# Paging





# Paging

11. **分页机制** 在src/kernel/main.c的os\_init()函数中有一处注释“Before setting up correct paging, no global variable can be used”. 尝试在main.c中定义一个全局变量

```
volatile int x = 0;
```

然后在调用init\_page()前使用这个全局变量

```
x = 10000;
```

并在串口初始化结束后输出它的值

```
printk("x = %d\n", x);
```

重新编译并运行，你发现了什么问题？请解释为什么在开启分页之前不能使用全局变量，但却可以使用局部变量（在init\_page()函数中使用了局部变量）。细心的你会发现，在开启分页机制之前，init\_page()中仍然使用了一些全局变量，但却没有造成错误，这又是为什么？（回答完本题后可以删除添加的代码。）


# Paging

- "x = 0"
  - kernel thinks it is located at 0xc0100000
  - &x = 0xc0101234
  - local variables are accessed via stack

- call init\_page ?

(128M) 0x08000000

0x00000000



- 
- 
- Lab1 feedback
    - 蓝框题



# Initialize current state

- the initial value of esp in GPR can be arbitrary
  - why?
- P.364 in "i386 manual"

```
ELSE (* OperandSize = 32, instruction = POPAD *)  
    EDI ← Pop();  
    ESI ← Pop();  
    EBP ← Pop();  
    throwaway ← Pop (); (* Skip ESP *)  
    EBX ← Pop();  
    EDX ← Pop();  
    ECX ← Pop();  
    EAX ← Pop();  
FI;
```



# Hardware context switch

## 硬件实现的上下文切换

事实上，上下文切换分为两种，分别由硬件和软件实现。Nanos, Windows和Linux都是使用软件实现的上下文切换。请搜索硬件实现上下文切换的相关信息，思考一下两者之间有什么不同，为何现代操作系统大多数都采用软件实现的方式？

- <http://stackoverflow.com/questions/2711044/why-doesnt-linux-use-the-hardware-context-switch-via-the-tss>



# PCB definition

## 另类的PCB(这个问题有难度)

有一种PCB的定义如下：

```
#define KSTACK_SIZE 4096
union PCB {
    uint8_t kstack[KSTACK_SIZE];
    struct {
        void *tf;
        // other fields
    };
};
```

这样的定义方式有什么好处？在SMP(对称多处理器)的环境下，就必须采取类似这样的方式来定义PCB，你知道为什么吗？



# PCB definition

- Any other way to obtain "current"?
  - %esp & 0xffffffff000
- For SMP
  - suppose we define "PCB\* current[NR\_CPU]"
  - How does one core obtain its "current"?



# Context switch

## 对上下文切换过程的思考

- 在asm\_do\_irq调用irq\_handle之前，有一条指令保存了当前栈顶指针的值。这条指令有什么目的？如果将其去掉，会有什么影响？
- 在你完成堆栈切换后，在src/kernel/irq/do\_irq.S中有一条指令必须去掉，否则会发生错误。思考一下为什么在完成堆栈切换之前需要保留该指令，完成堆栈切换之后却必须将其去掉？该指令本来想干什么？现在其作用是否在哪里实现了？如果你觉得很晕，你可以用纸笔画出堆栈的变化，人工模拟上下文切换过程。



# Context switch

```
41     movw $SELECTOR_KERNEL(SEG_KERNEL_DATA), %ax
42     movw %ax, %ds
43     movw %ax, %es
44
45     pushl %esp
46     call irq_handle
47
48 # YOU NEED TO SWITCH STACK TO current->tf
49 # SO YOU NEED TWO ADD TWO LINES OF INTERRUPT CODE
50 # HINT:
51 #     1. USE movl INSTRUCTION
52 #     2. USE (address) CAN REFERENCE MEMORY LOCATION
53 #     3. YOU MAY FLUSH ANY GENRAL PURPOSE REGISTER AS
54 #     4. REGISTERS ARE REFERENCED BY "%%", SUCH AS %es
55
56 ##### your work #####
57
58 #####
59 #####
60     #addl $4, %esp #when you finish this task, this
61
62     popal
```


trap frame

ESP  
p1.tf

p2.tf


0xffffffff

0x00000000



# Context switch

```
41     movw $SELECTOR_KERNEL(SEG_KERNEL_DATA), %ax
42     movw %ax, %ds
43     movw %ax, %es
44
45     pushl %esp
46     call irq_handle
47
48 # YOU NEED TO SWITCH STACK TO current->tf
49 # SO YOU NEED TWO ADD TWO LINES OF INTERRUPT CODE
50 # HINT:
51 #     1. USE movl INSTRUCTION
52 #     2. USE (address) CAN REFERENCE MEMORY LOCATION
53 #     3. YOU MAY FLUSH ANY GENRAL PURPOSE REGISTER AS
54 #     4. REGISTERS ARE REFERENCED BY "%%", SUCH AS %es
55
56 ##### your work #####
57
58 #####
59 #####
60     #addl $4, %esp #when you finish this task, this
61
62     popal
```



# Context switch

```
41     movw $SELECTOR_KERNEL(SEG_KERNEL_DATA), %ax
42     movw %ax, %ds
43     movw %ax, %es
44
45     pushl %esp
46     call irq_handle
47
48 # YOU NEED TO SWITCH STACK TO current->tf
49 # SO YOU NEED TWO ADD TWO LINES OF INTERRUPT CODE
50 # HINT:
51 #     1. USE movl INSTRUCTION
52 #     2. USE (address) CAN REFERENCE MEMORY LOCATION
53 #     3. YOU MAY FLUSH ANY GENRAL PURPOSE REGISTER AS
54 #     4. REGISTERS ARE REFERENCED BY "%", SUCH AS %es
55
56 ##### your work #####
57
58
59 #####
60     #addl $4, %esp #when you finish this task, this
61
62     popal
```

0xffffffff


trap frame

p1.tf

ESP

p2.tf

0x00000000



# Never-return function

ESP



???
EFLAGS
CS
EIP
1000
EAX
EBX
ECX
EDX
old_ESP
EBP
ESI
EDI

## 不能返回的函数

目前你必须保证线程函数永远不会返回，否则将会发生错误。为什么从线程函数返回就会发生错误？



# Use printk in do\_irq.S

## 在汇编代码中调用printk

你可能需要在asm\_do\_irq中调用printk来帮助你进行调试。想一想怎么在汇编代码中调用它，并尝试付诸实践。

```
.extern printk  
msg:  
.asciz "esp = %x\n"  
  
pushl %esp  
pushl $msg  
call printk  
addl $8, %esp
```



# Stack overflow

```
void stackoverflow(int x)
{
    if(x==0)
        printf("%d ", x);
    if(x>0)
        stackoverflow(x-1);
}
void keep_stackoverflow()
{
    while(1){stackoverflow(16384*1000);}
}
```

# Stack overflow

```
143 c01001f0 <stackoverflow>:
144 c01001f0: 55                push    %ebp
145 c01001f1: 89 e5             mov     %esp,%ebp
146 c01001f3: 83 ec 18          sub     $0x18,%esp
147 c01001f6: 8b 45 08           mov     0x8(%ebp),%eax
148 c01001f9: 83 f8 00           cmp     $0x0,%eax
149 c01001fc: 74 07             je      c0100205 <stackoverflow+0x15>
150 c01001fe: 7e 19             jle     c0100219 <stackoverflow+0x29>
151 c0100200: 83 e8 01          sub     $0x1,%eax
152 c0100203: 75 fb             jne     c0100200 <stackoverflow+0x10>
153 c0100205: c7 44 24 04 00 00 00 movl    $0x0,0x4(%esp)
154 c010020c: 00
155 c010020d: c7 04 24 2c 2b 10 c0 movl    $0xc0102b2c,(%esp)
156 c0100214: e8 d7 08 00 00    call    c0100af0 <printk>
157 c0100219: c9               leave
158 c010021a: c3               ret
159 c010021b: 90               nop
160 c010021c: 8d 74 26 00      lea     0x0(%esi,%eiz,1),%esi
```

# Threads with parameters

ESP →

1234
???
EFLAGS
CS
EIP
1000
EAX
EBX
ECX
EDX
old_ESP
EBP
ESI
EDI

## 带有参数的线程(有些难度)

我们知道在Linux下可以编写从外部读入参数的程序，只需要把main函数的参数声明改为`int main(int argc, char *argv[])`即可。在创建线程的时候如何实现类似的功能？

先来个简单一点的吧，让`create_kthread`多接受一个整型参数：

`create_kthread(void *fun, int arg)`。然后只需要编写一个测试函数：





- 
- Hints for implementing locking



# Locking

---

- upgraded version of cli() & sti() to solve
  - nested locking
  - sleep during locking
  - locking in interrupt
- Have you triggered assertion fail/mysterious reboot when testing your implementation?



# Where does the problem come from?

- 
- Interrupt is not enabled when it should.
  - Interrupt is enabled when it should not.
  - How to find these bugs?



# Assertion

- check the status of IF bit in EFLAGS

```
#define INTR assert(read_eflags() & IF_MASK)  
#define NOINTR assert(~read_eflags() & IF_MASK)
```

- insert them in your code
- consistency
  - NOINTR when in critical region or during interrupt
  - INTR otherwise

# Trap 1 - nested locking

```
void V(Sem *s) {  
    INTR;  
    lock(); NOINTR;  
    // ...  
    wakeup(p); NOINTR;  
    // ...  
    unlock();  
    INTR;  
}
```


```
void wakeup(PCB *p) {  
    lock(); NOINTR;  
    // ...  
    unlock();  
}
```

not safe  
any longer



## Trap 2 – sleep during locking

```
void P(Sem *s) {  
    INTR;  
    lock(); NOINTR;  
    // ...  
    if(counter == 0) { sleep(); NOINTR; }  
    // ...  
    unlock();  
    INTR;  
}
```



the consistency of locking  
should not be violated by  
other processes

# Trap 3 - use locking in interrupts

```
void timer_handler() {  
    // ...  
    NOINTR;  
    V(sem);  
    NOINTR;  
    // ...  
}
```

```
void V(Sem *s) {  
    lock(); NOINTR;  
    // ...  
    unlock();  
}
```

← should not sti()



# Test case

```
void
test_consumer(void) {
    while (1) {
        P(&mutex); INTR;
        P(&full); INTR;
        // ...
        V(&empty); INTR;
        V(&mutex); INTR;
    }
}
```





# Implementation

- 1. nested locking
- 2. sleep in locking
- 3. locking in interrupt
  
- (1) → locking counter
  - `assert(lock_cnt >= 0);`
- (2) → counter per thread
- (3) → store IF before the first locking, restore it when leaving the most outside critical region



# Lab2 is out!

---

- the second stage
  - implement message passing
  - create 4 kernel thread to print "abcdabcd..."
    - communicate with message passing
- the third stage
  - add device drivers
    - test your message passing
  - to be continue...
- Have fun!