

System call

- `exec()`
- `fork()`
- `exit()`
- other system calls



- `exec()`



In the previous task

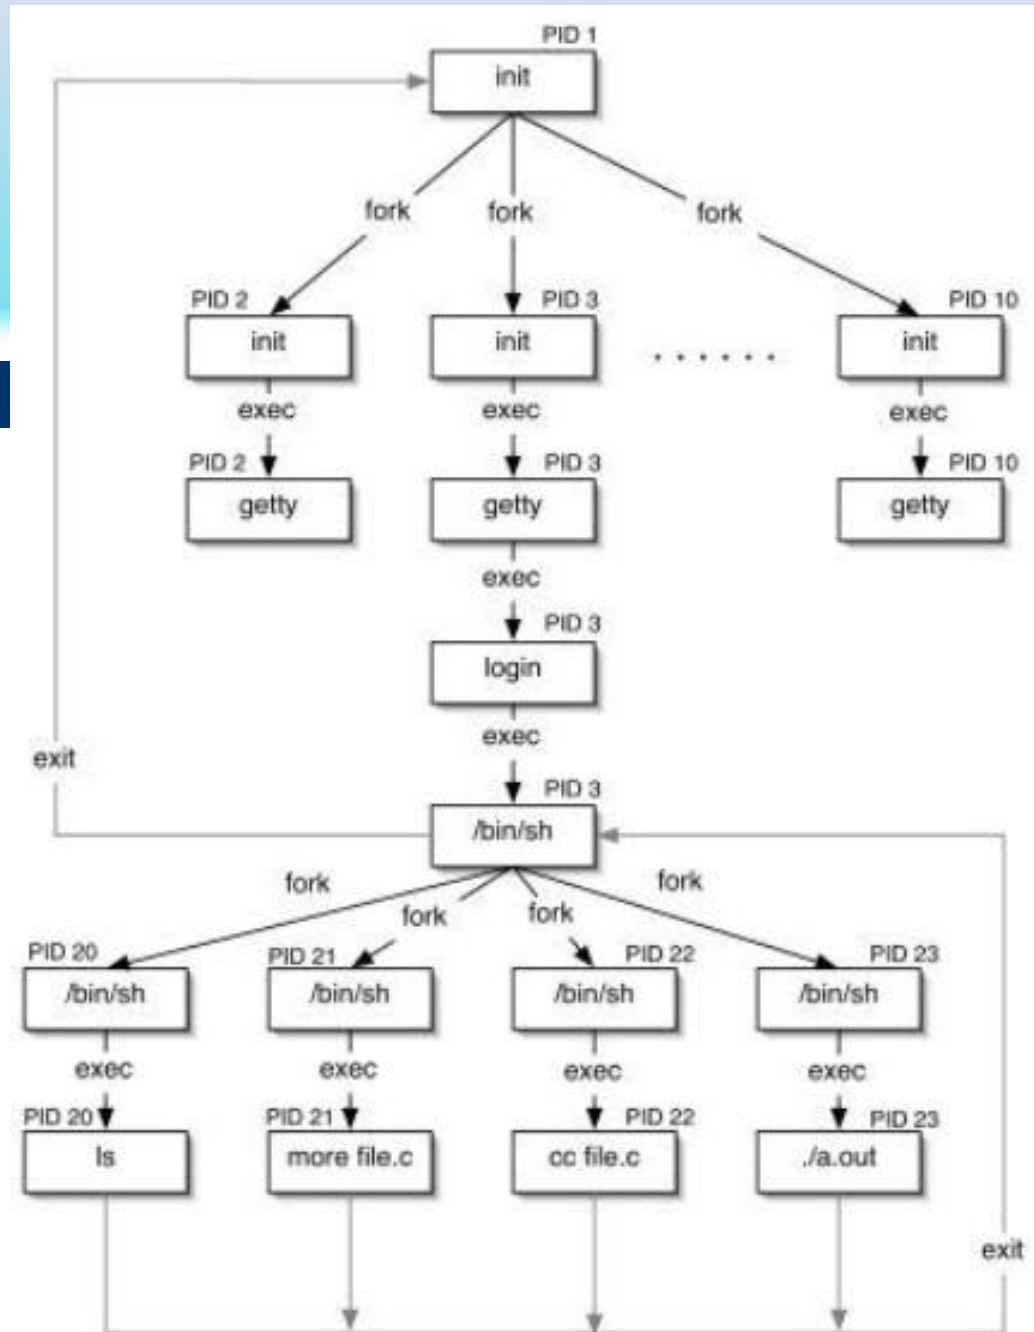
- You have loaded the first user process by PM in kernel.
- But there is only one user process in the OS.
- OS should provide user processes system calls for process management to let them
 - create a child process
 - execute other programs
 - exit normally



Load other programs

- Loading a new program is a legal demand for user process.
- How to execute other programs in user space?
- `exec()` - execute a program
 - replace itself with another program

`fork() + exec()`
= everything!





exec()

- It seems complicated
 - reclaim ALMOST all resource
 - re-allocate the address space
 - load the new program
- mainly handled by PM
 - communicate with MM and FM



Reclaim resource

- PCB, semaphore, message, address space, file descriptor table...
- File descriptor table should not be reclaimed.
 - we will explain it in lab4
- PID does not need to change.

```
1 #include <stdio.h>
2 #include <unistd.h>
3
4 int main(){
5     printf("my pid = %d\n", getpid());
6     execl("./test", NULL); // execute itself again
7     return 0;
8 }
```



The following work

- re-allocate address space
 - just the same as creating address space for the first user program "0"
- load the new program
 - just the same as loading the first user program "0"
- re-initialize PCB
- put the process into ready queue

Arguments

- Executing with arguments is allowed.

```
1 #include <stdio.h>
2
3 int main(int argc, char *argv[]){
4     int i;
5     for( i = 0; i < argc; i ++){
6         printf("argv[%d] = %s\n", i, argv[i]);
7     }
8     return 0;
9 }
```

```
[525][0: ~/test]$ ./test -abc  ( ^ (∞) ^ )
argv[0] = ./test
argv[1] = -abc
argv[2] = 
argv[3] = ( ^ (∞) ^ )
[526][0: ~/test]$
```




exec() with arguments

- can be "arbitrary" numbers of arguments
 - there is a limit
- How to pass them as system call arguments?
 - see "man exec"
- Nanos simplification: encode multiple arguments into a single string.
 - `exec(3, "abc 234 third_arg")`
 - `int main(char *args)`
 - `CFLAGS += -Wno-main`

Pass arguments to user program

- They are arguments of main().
- Where are they located?
 - stack!
- pay attention to pointers
- For multiple arguments, how to implement argc & argv?



Arguments of main() in gdb

Starting program: /home/user/test/test -abc o (J o

Breakpoint 1, main (argc=4, argv=0xbffff4c4) at test.c:5

```
5          for( i = 0; i < argc; i ++ ) {
```

(gdb) x/100c 0xbffff643


```
0xbffff643:  47 '/' 104 'h' 111 'o' 109 'm' 101 'e' 47 '/' 117 'u' 115 's'
0xbffff64b: 101 'e' 114 'r' 47 '/' 116 't' 101 'e' 115 's' 116 't' 47 '/'
0xbffff653: 116 't' 101 'e' 115 's' 116 't' 0 '\000' 45 '-' 97 'a' 9
8 'b'
0xbffff65b: 99 'c' 0 '\000' -27 '\345' -101 '\233' -89 '\24
7' 0 '\000' 111 'o' -17 '\357'
0xbffff663: -68 '\274' -120 '\210' -30 '\342' -107 '\225' -
81 '\257' -30 '\342' -106 '\226' -95 '\241'
0xbffff66b: -30 '\342' -107 '\225' -80 '\260' -17 '\357' -
68 '\274' -119 '\211' 111 'o' 0 '\000'
0xbffff673: 79 'O' 82 'R' 66 'B' 73 'I' 84 'T' 95 '_' 83 'S' 79 'O'
0xbffff67b: 67 'C' 75 'K' 69 'E' 84 'T' 68 'D' 73 'I' 82 'R' 61 '='
0xbffff683: 47 '/' 116 't' 109 'm' 112 'p' 47 '/' 111 'o' 114 'r' 98 'b'
0xbffff68b: 105 'i' 116 't' 45 '-' 121 'y' 122 'z' 104 'h' 0 '\000' 8
3 'S'
0xbffff693: 83 'S' 72 'H' 95 '_' 65 'A' 71 'G' 69 'E' 78 'N' 84 'T'
```



Obtain arguments from `exec()`

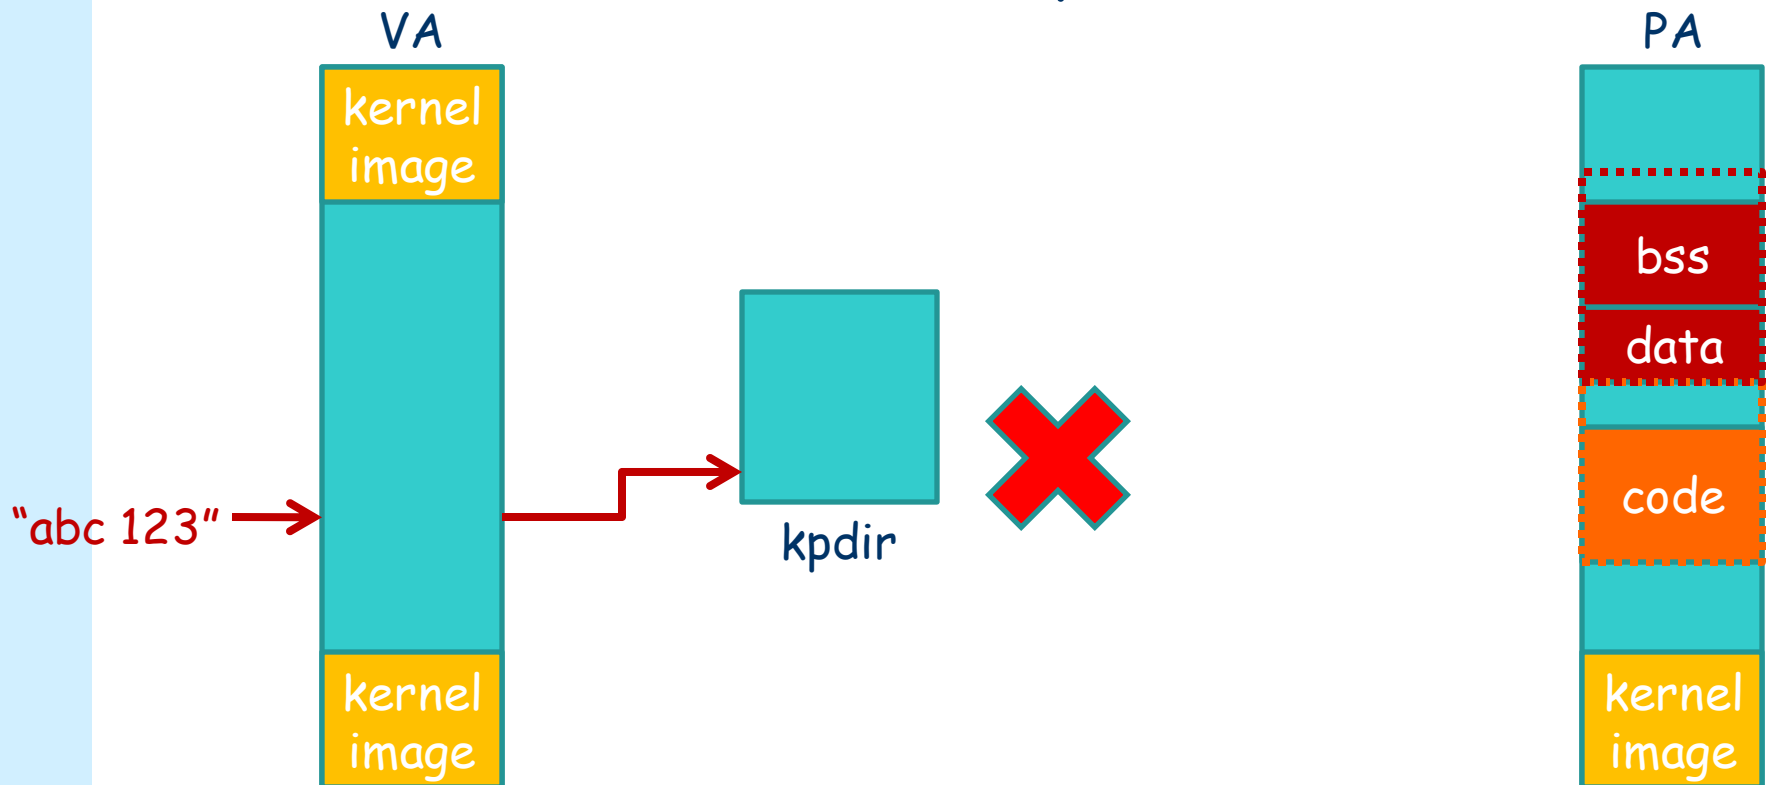
- It seems trivial, but it does not.
- PM can not use the argument string passed from user process directly.
- Why?

User process' view



PM's view

- PM shares the address space with kernel.





Solution

- PM should simulate the process of address translation to get the “physical” address of the argument string.
- Use the “physical” address to access the argument string.
- Why this works?
- How to implement the simulation?



Return value

- When `exec()` succeeds, it never returns.
 - It is replaced by another program successfully.
 - PM does not need to send a reply to the “original” user process.
- What should be done when `exec()` fails?
 - notify the user process by a special return value
 - or simply call `panic` in Nanos



-
- `fork()`



In the previous task

- You have loaded the first user process.
- But `exec()` cannot produce new processes.
- Now it is the time to implement `fork()`!
 - allow "creating" new processes in user space



fork()

- duplicate itself
 - address space, process state, resource...
 - except for PID
- PM, MM, FM should cooperate to handle a fork request
- user process calls fork()
 - trap into kernel
 - send message to PM
 - wait for reply



PM's work

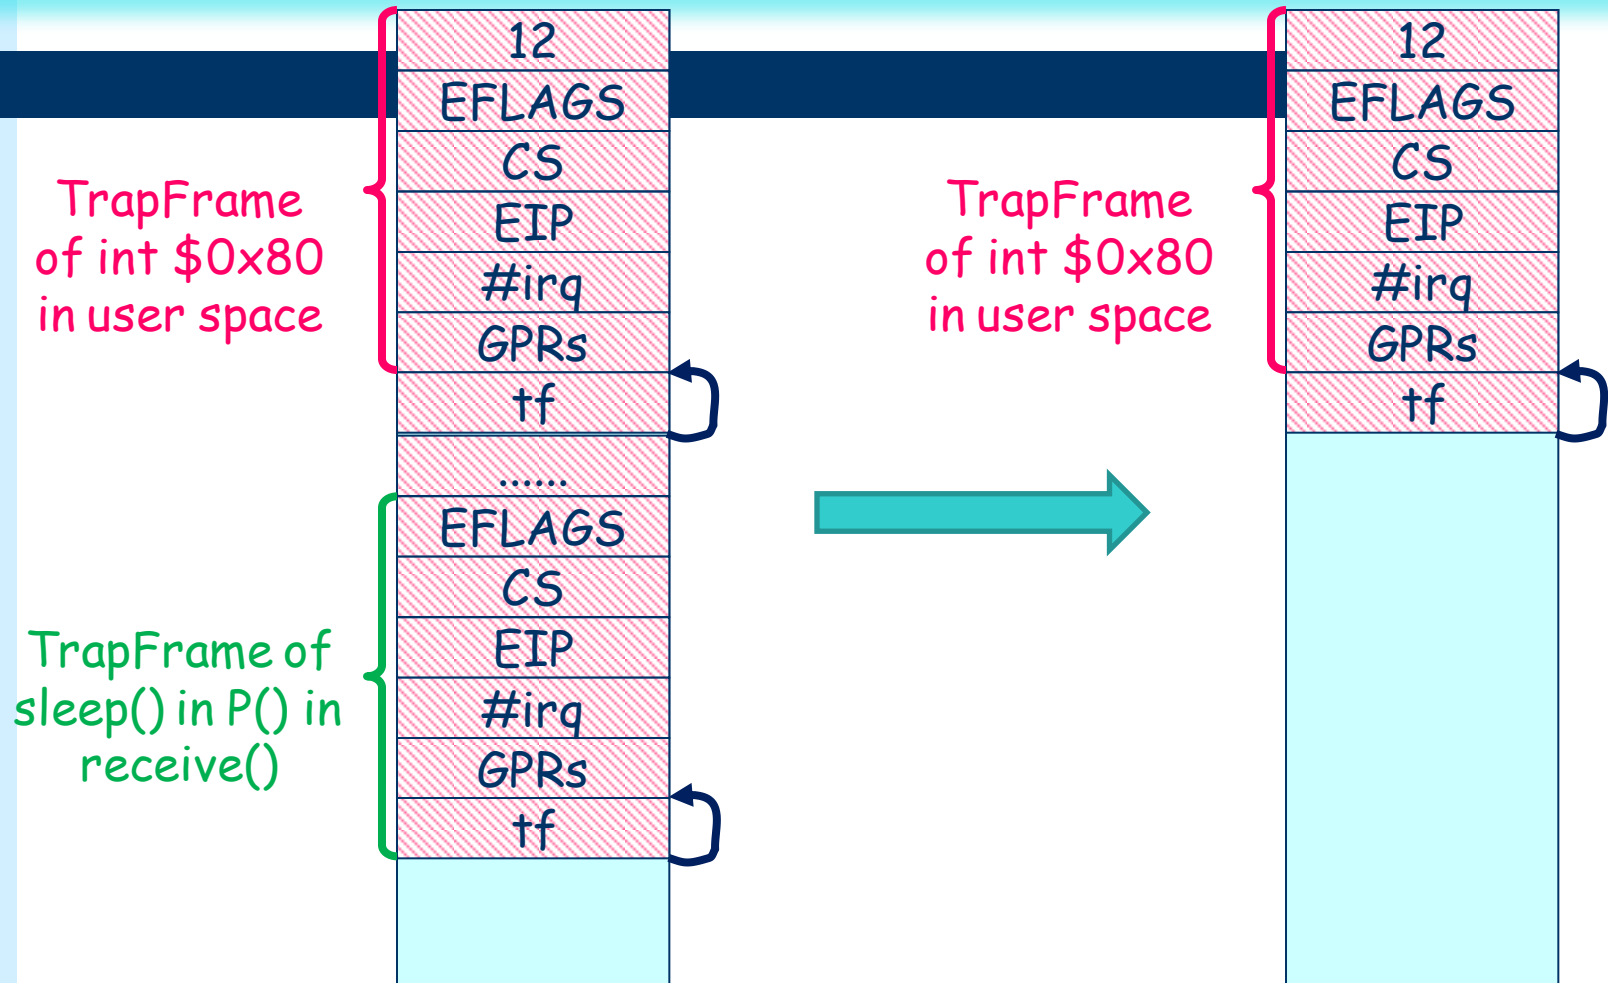
- allocate a free PCB
- clone the process state
 - flags
 - current state
 - pcb->tf
- pay attention to pointer fields !!!



PM's work (cont.)

- Father process is now blocked during a system call.
 - waiting for PM's reply
- To make the child process blocked is tricky.
 - set the same states of message queue and semaphore as the father process's
- A simplification is to let the child process run at right.
 - as if it just receives message reply from PM

Simplification



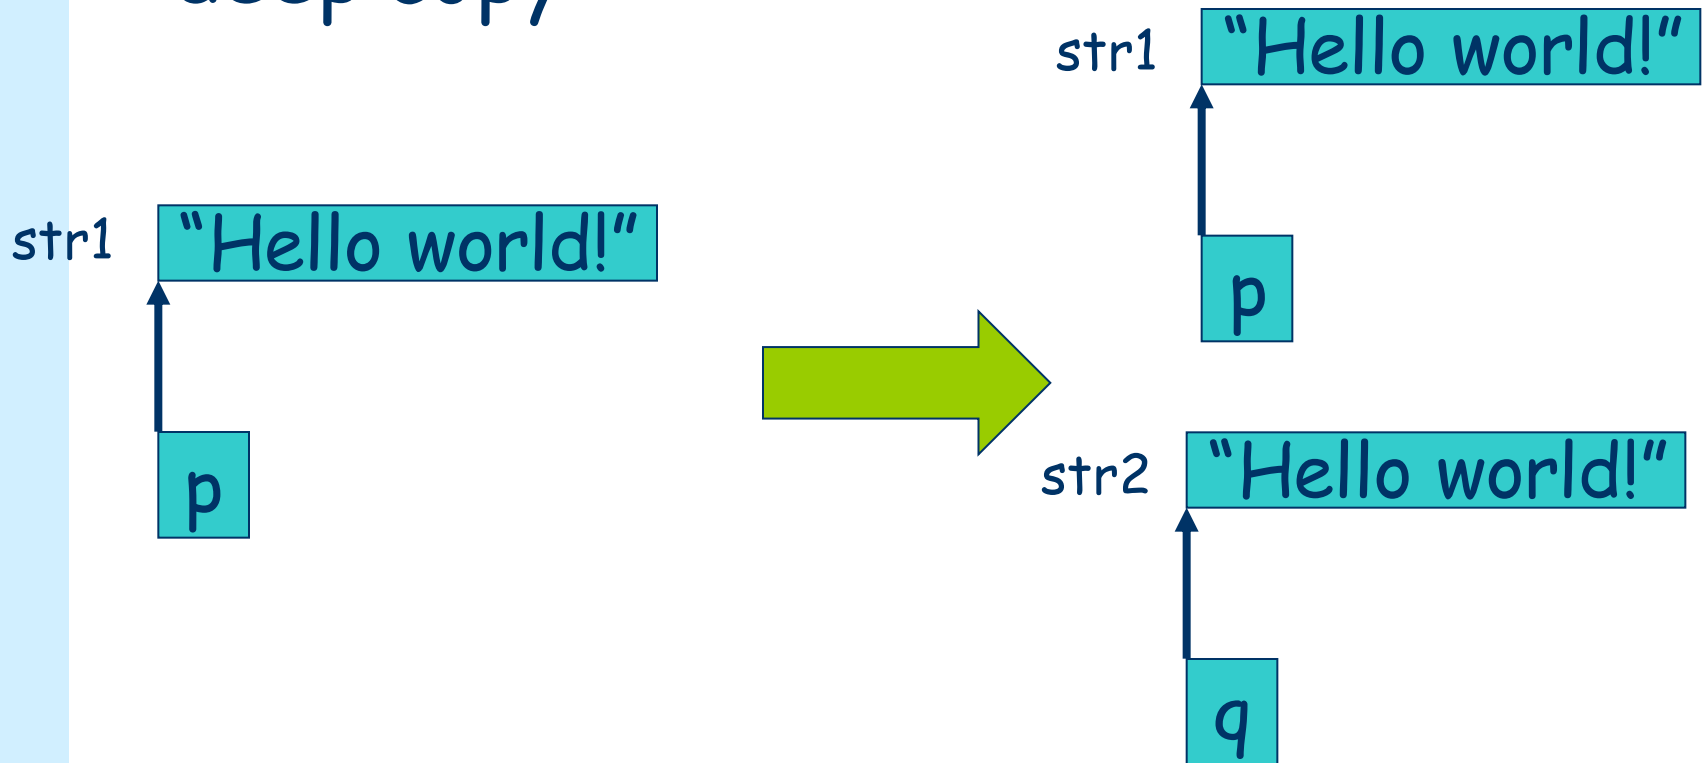


MM's work

- create new address space for child process
 - allocate new page directory and page tables, as well as physical pages
 - map memory above 0xc0000000 to kernel*
- Code is read only, so it can be shared by mapping to the same physical page.
 - This is optional. For simplicity, copy the code, too.
- Data and stack should not be shared.
 - pay attention to kernel stacks


Deep copy & shallow copy

- deep copy




Deep copy & shallow copy

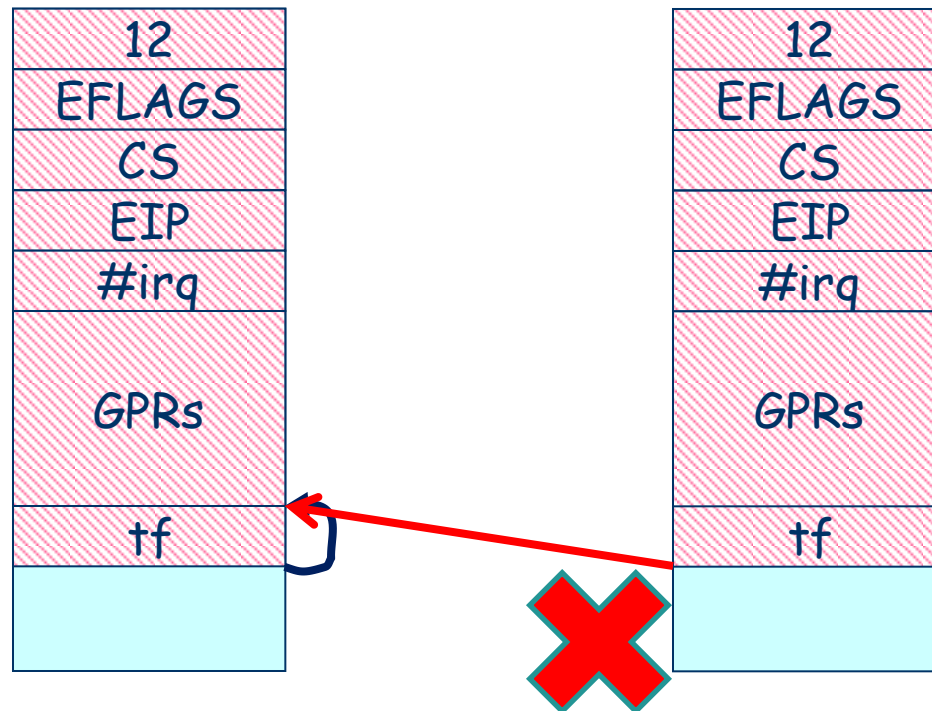
- shallow copy



Problem



Problem





FM's work

- Nothing to do in Lab3.
- It will maintain the file descriptor table in Lab4.



When finished

- put the child process into ready queue
 - do not block the child process for simplicity
- send a reply message to father process
- father and child are running !
- return value
 - `fork()` returns 0 for child, and the PID of the child process for father



-
- `exit()`



exit()

- inform the kernel about process termination
 - kernel should reclaim all resource
- It is straightforward.
 - just reclaim all resource
 - including PCB
 - the process disappears



exit() (cont.)

- The reason why all test threads/processes you created before cannot return:
 - there is not a mechanism for normal process termination
- How to make process exit automatically once returning from main()?



Compiler hack

```
_start() {  
    // initialization  
    main();  
    exit();  
}
```

- make `_start()` the real entry point

Compiler hack (cont.)

```
[502][0: ~/test]$ readelf -e test
```

ELF Header:

```
  Magic:   7f 45 4c 46 01 01 01 00 00 00 00 00 00 00 00 00
  Class:                               ELF32
  Data:                                   2's complement, little endian
  Version:                             1 (current)
  OS/ABI:                              UNIX - System V
  ABI Version:                         0
  Type:                                EXEC (Executable file)
  Machine:                             Intel 80386
  Version:                             0x1
  Entry point address:                 0x8048310
```

```
253 Disassembly of section .text:
```

```
254
```

```
255 08048310 <_start>:
```

```
256 8048310: 31 ed                xor    %ebp,%ebp
```


```
257 8048312: 5e                  pop    %esi
```

Compiler hack (cont.)

```
(gdb) si
0x08048410 in main (argc=134513604, argv=0x1) at test.c:9
9      }
0x08048410 <main+76>:    c3      ret
(gdb)
0xb7e9eca6 in __libc_start_main () from /lib/i686/cmov/libc.so.6
0xb7e9eca6 <__libc_start_main+230>:    89 04 24      mov    %eax, (%esp)
(gdb)
0xb7e9eca9 in __libc_start_main () from /lib/i686/cmov/libc.so.6
0xb7e9eca9 <__libc_start_main+233>:    e8 72 86 01 00 call   0xb7eb7320 <exit>
(gdb) si
0xb7eb7320 in exit () from /lib/i686/cmov/libc.so.6
0xb7eb7320 <exit+0>:    55      push    %ebp
(gdb)
0xb7eb7321 in exit () from /lib/i686/cmov/libc.so.6
0xb7eb7321 <exit+1>:    89 e5    mov    %esp, %ebp
(gdb)
0xb7eb7323 in exit () from /lib/i686/cmov/libc.so.6
0xb7eb7323 <exit+3>:    53      push    %ebx
(gdb)
```

Stack hack

- make the return address of main() points to the exit code






-
- other system calls


getpid()

- get the pid of current process
 - no need to communicate with servers
- 打酱油1号



sleep()

- block itself for several seconds
 - TIMER serves as an alarm
- 打酱油2号



waitpid()

- wait for a process to terminate
 - when a process *A* exits, notify those processes waiting for *A*
- How to implement?

