

Enhance protection and memory management

- User process in ring3
- More enhancement

- 
-
- User process in ring3

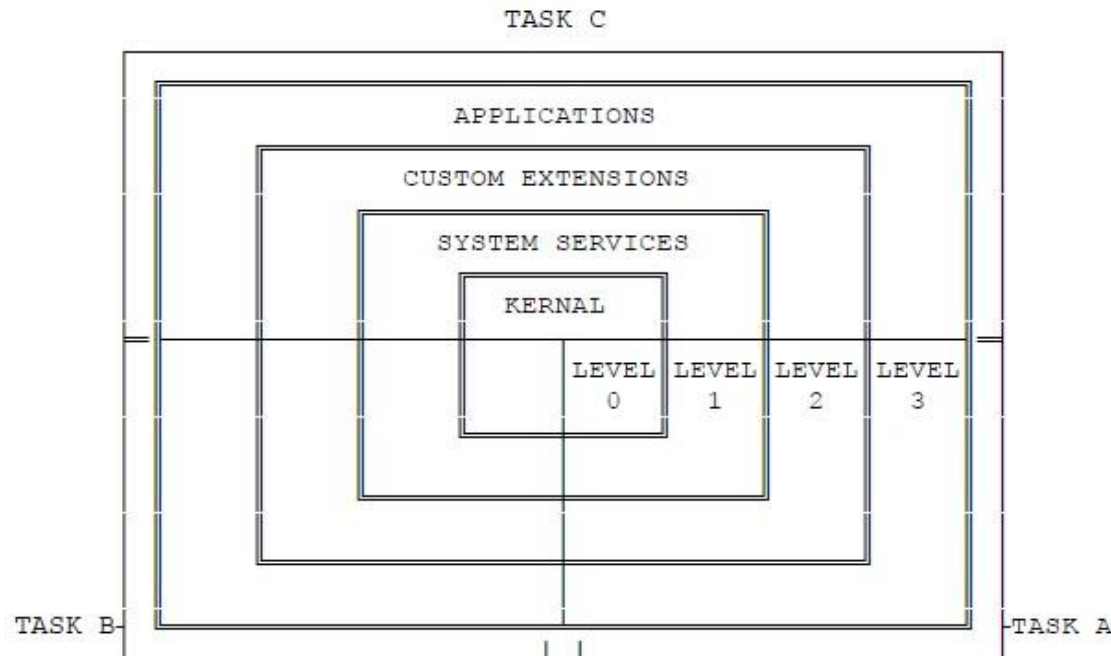


Put user processes in ring3

- When creating the first process
 - use the user code segment & user data segment
 - allocate one page for user stack
 - 0xbffff000 - 0xbfffffff
- But this will cause some problems.
 - We have to fix it.

Levels of privilege

- There are 4 types of the privilege level for a segment in 386 CPU.





A problem

- In real OS, user process has low privilege level.
 - code, data and stack are in ring3
- Interrupt handlers cannot run with user stack according to the hardware mechanism.
 - Why?
 - nonconforming stack: $SS.DPL > CPL$
- What if interrupt handlers are permitted to run with user stack?
 - Can you construct a counter-example to harm the system?

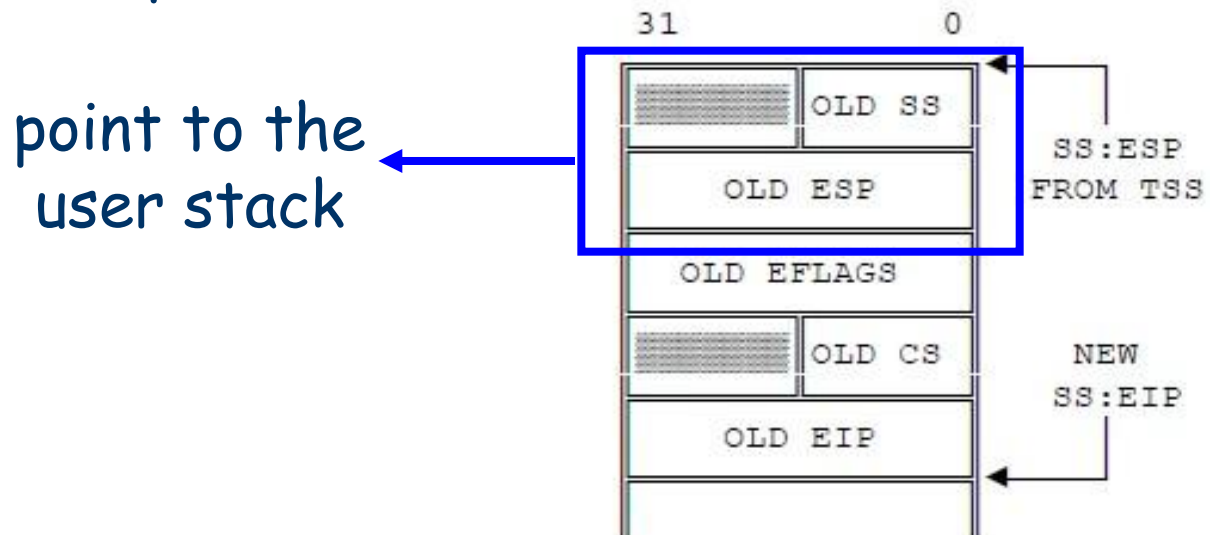


Solution

- Before handling the interrupt, perform stack switching.
 - use another stack instead of the user stack
 - stack switching is performed by hardware
 - kernel guarantees enough space in the new stack to handle the interrupt

New member of current state

- The old stack pointer is stored in the new stack.
 - still can go back to user stack when returning from an interrupt

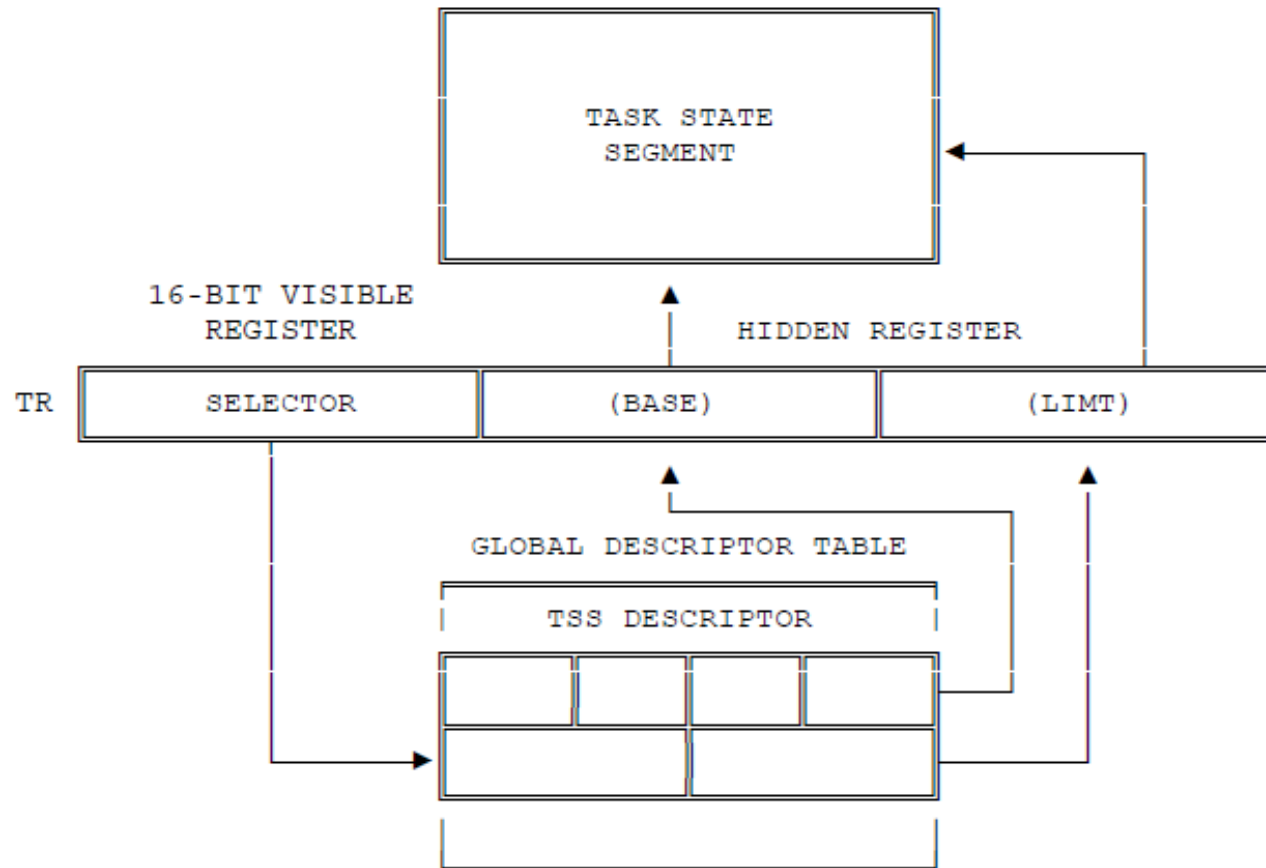


Finding new stack

- TSS
 - task state segment
- 3 stacks for 3 PL
 - ring0, ring1, ring2
 - Why no stack for ring3?
- We only use the stack for ring0.

31	23	15	7	0	
I/O MAP BASE					0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 T
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0					LDT 64
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0					GS 60
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0					FS 56
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0					DS 54
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0					SS 50
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0					CS 46
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0					ES 42
					EDI 44
					ESI 40
					EBP 36
					ESP 32
					EBX 28
					EDX 24
					ECX 20
					EAX 16
					EFLAGS 12
					INSTRUCTION POINTER (EIP) 8
					CR3 (PDPR) 4
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0					SS2 16
					ESP2 12
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0					SS1 10
					ESP1 0C
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0					SS0 8
					ESP0 4
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0					BACK LINK TO PREVIOUS TSS 0

Finding TSS





What should you do?

- set `tf->ss` & `tf->esp` to the user stack when creating the first user process
- set `tss.esp0` to the new process' kernel stack during context switching
- That's all!
- Why does each user process own its kernel stack?
 - What if we set only one kernel stack for all user processes?



Everything else is done by hardware

- stack switching
 - finding TSS
 - loading kernel stack
 - store old stack pointer in the kernel stack
 - store current state
- returning from interrupt to user space
 - When executing iret, switch back to the user stack according to `tf->ss` & `tf->esp`.



-
- Other enhancement



Catch "Segmentation fault"

- When a user process tries to perform illegal operation, kill it, then print "Segmentation fault" to tty, instead of triggering system panic.
- First, enhance `exit()` and `waitpid()`.
 - `void exit(int exit_code);`
 - `int waitpid(pid_t pid);`
 - return the exit code of the process "pid"



Catch “Segmentation fault” (cont.)

- One implementation is as following
 - Distinguish between illegal operations from user processes and bugs from kernel.
 - How?
 - After identifying illegal operations, issue an `exit()` system with a special exit code.
 - shell waits for child process to exit, it checks `chlid process' exit code`.
 - If it is the special exit code, print “Segmentation fault” on the screen.



Set lower privilege level for user pages

- restricting addressable domain
 - clear U/S bit in PDE & PTE for kernel space
 - set U/S bit in PDE & PTE for user space
- type checking
 - clear R/W bit in PDE & PTE for pages for user code
 - set R/W bit in PDE & PTE for pages for others
- For combining two-level protection, see i386 manual for details.



Copy-on-write

- Copy-on-write may dramatically decrease the cost of `fork()`.
- When handling `fork()`, maps child process' code and data to the father's pages, instead of cloning the whole address space.
 - User stack should still be cloned.
- Label these pages as read-only in the two parties' PTE.



Copy-on-write (cont.)

- Once any of the two parties writes to a labeled page, a page fault exception is triggered.
 - Distinguish this accessing from an illegal operation.
 - How to distinguish?
 - Clone this page for the child process, and set the two pages writable.
 - Pages for code are still read-only.



Copy-on-write (cont.)

- When a process exits, the shared pages should not be reclaimed right away.
 - Another process still needs to access these pages.
 - A page is reclaimed when no other process will access it.
- How to maintain the process-page relationship?



malloc and free

- heap management
 - allocate and free memory dynamically
- Where is heap?
 - next to the static data
 - grow up (to the direction of high address)
- malloc() and free() are not system calls.
 - OS will not manage user heaps directly.



malloc and free (cont.)

- `void* sbrk(int inc);`
 - system call to change the size of data segment
 - increase the size of data segment by "inc" bytes
 - when $inc < 0$, the size of data segment is decreased
 - return the top of data segment
- `malloc(nr_byte)`
 - first check whether there is a continuous space of size "nr_byte" in heap
 - if yes, return the start address of this memory space
 - if not, call `sbrk()` for more heap space, then perform allocation



malloc and free (cont.)

- call `sbrk(0)` when `malloc()` is called for the first time
 - obtain the top of static data / the start of heap
- When handling `fork()`, pages for user heap should be cloned, too.
 - can apply copy-on-write
- When handling `exit()`, pages for user heap should be reclaimed, too.
 - avoid memory leak